

SQL vs NoSQL





Somkiat Puisungnoen

Somkiat Puisungnoen

Update Info 1 View Activity Log 10+ ...

Timeline About Friends 3,138 Photos More

When did you work at Opendream? X

... 22 Pending Items

Intro

Software Craftsmanship

Software Practitioner at สยามชัมนาณกิจ พ.ศ. 2556

Agile Practitioner and Technical at SPRINT3r

Post Photo/Video Live Video Life Event

What's on your mind?

Public Post

Somkiat Puisungnoen 15 mins · Bangkok · ⚙️

Java and Bigdata



somkiat.cc

Page Messages Notifications 3 Insights Publishing Tools Settings Help ▾

somkiat.cc
@somkiat.cc

Home Posts Videos Photos

Like Liked Following Share ...

+ Add a Button



Agenda

- Database model
- RDBMS
- SQL vs. NoSQL
- Workshop
 - Redis
 - PostgreSQL
- SQL tuning with PostgreSQL



Database ?



Database = RDBMS ?



Relation DataBase Management System





Postgre^{SQL}

Informix®



RDBMS

Create a protective wrapper around data

Give data a pace to live & control modification

Secure data by enforce access permission

Manage system performance for read/write data



Relational = Database model



Table

customer_id	fname	lname	last_purchase_date	status_level
1234	Alice	Smith	23-Jan-2016	silver
2345	Bob	Washington	14-July-2018	bronze
2352	Charlie	Johnson	20-Dec-2018	gold
2678	Chendong	Mao	19-Dec-2017	gold
3422	Javier	Valdez	23-Dec-2018	gold
3577	Avery	Winston	5-Jan-2019	silver
4240	Charles	Davis	15-Jan-2017	bronze
6235	Vihaan	Patel	12-Nov-2018	gold
2352	Katherine	Coltrane	04-Apr-2016	silver



Database Models

Rank			DBMS	Database Model
Dec 2020	Nov 2020	Dec 2019		
1.	1.	1.	Oracle 	Relational, Multi-model 
2.	2.	2.	MySQL 	Relational, Multi-model 
3.	3.	3.	Microsoft SQL Server 	Relational, Multi-model 
4.	4.	4.	PostgreSQL 	Relational, Multi-model 
5.	5.	5.	MongoDB 	Document, Multi-model 
6.	6.	6.	IBM Db2 	Relational, Multi-model 
7.	7.	↑ 8.	Redis 	Key-value, Multi-model 
8.	8.	↓ 7.	Elasticsearch 	Search engine, Multi-model 
9.	9.	↑ 11.	SQLite 	Relational
10.	10.	10.	Cassandra 	Wide column
11.	11.	↓ 9.	Microsoft Access	Relational
12.	12.	↑ 13.	MariaDB 	Relational, Multi-model 
13.	13.	↓ 12.	Splunk	Search engine
14.	14.	↑ 15.	Teradata 	Relational, Multi-model 
15.	15.	↓ 14.	Hive	Relational

<https://db-engines.com/en/ranking>



Relational

Rank			DBMS	Database Model
Dec 2020	Nov 2020	Dec 2019		
1.	1.	1.	Oracle +	Relational, Multi-model i
2.	2.	2.	MySQL +	Relational, Multi-model i
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model i
4.	4.	4.	PostgreSQL +	Relational, Multi-model i
5.	5.	5.	MongoDB +	Document, Multi-model i
6.	6.	6.	IBM Db2 +	Relational, Multi-model i
7.	7.	↑ 8.	Redis +	Key-value, Multi-model i
8.	8.	↓ 7.	Elasticsearch +	Search engine, Multi-model i
9.	9.	↑ 11.	SQLite +	Relational
10.	10.	10.	Cassandra +	Wide column
11.	11.	↓ 9.	Microsoft Access	Relational
12.	12.	↑ 13.	MariaDB +	Relational, Multi-model i
13.	13.	↓ 12.	Splunk	Search engine
14.	14.	↑ 15.	Teradata +	Relational, Multi-model i
15.	15.	↓ 14.	Hive	Relational

<https://db-engines.com/en/ranking>



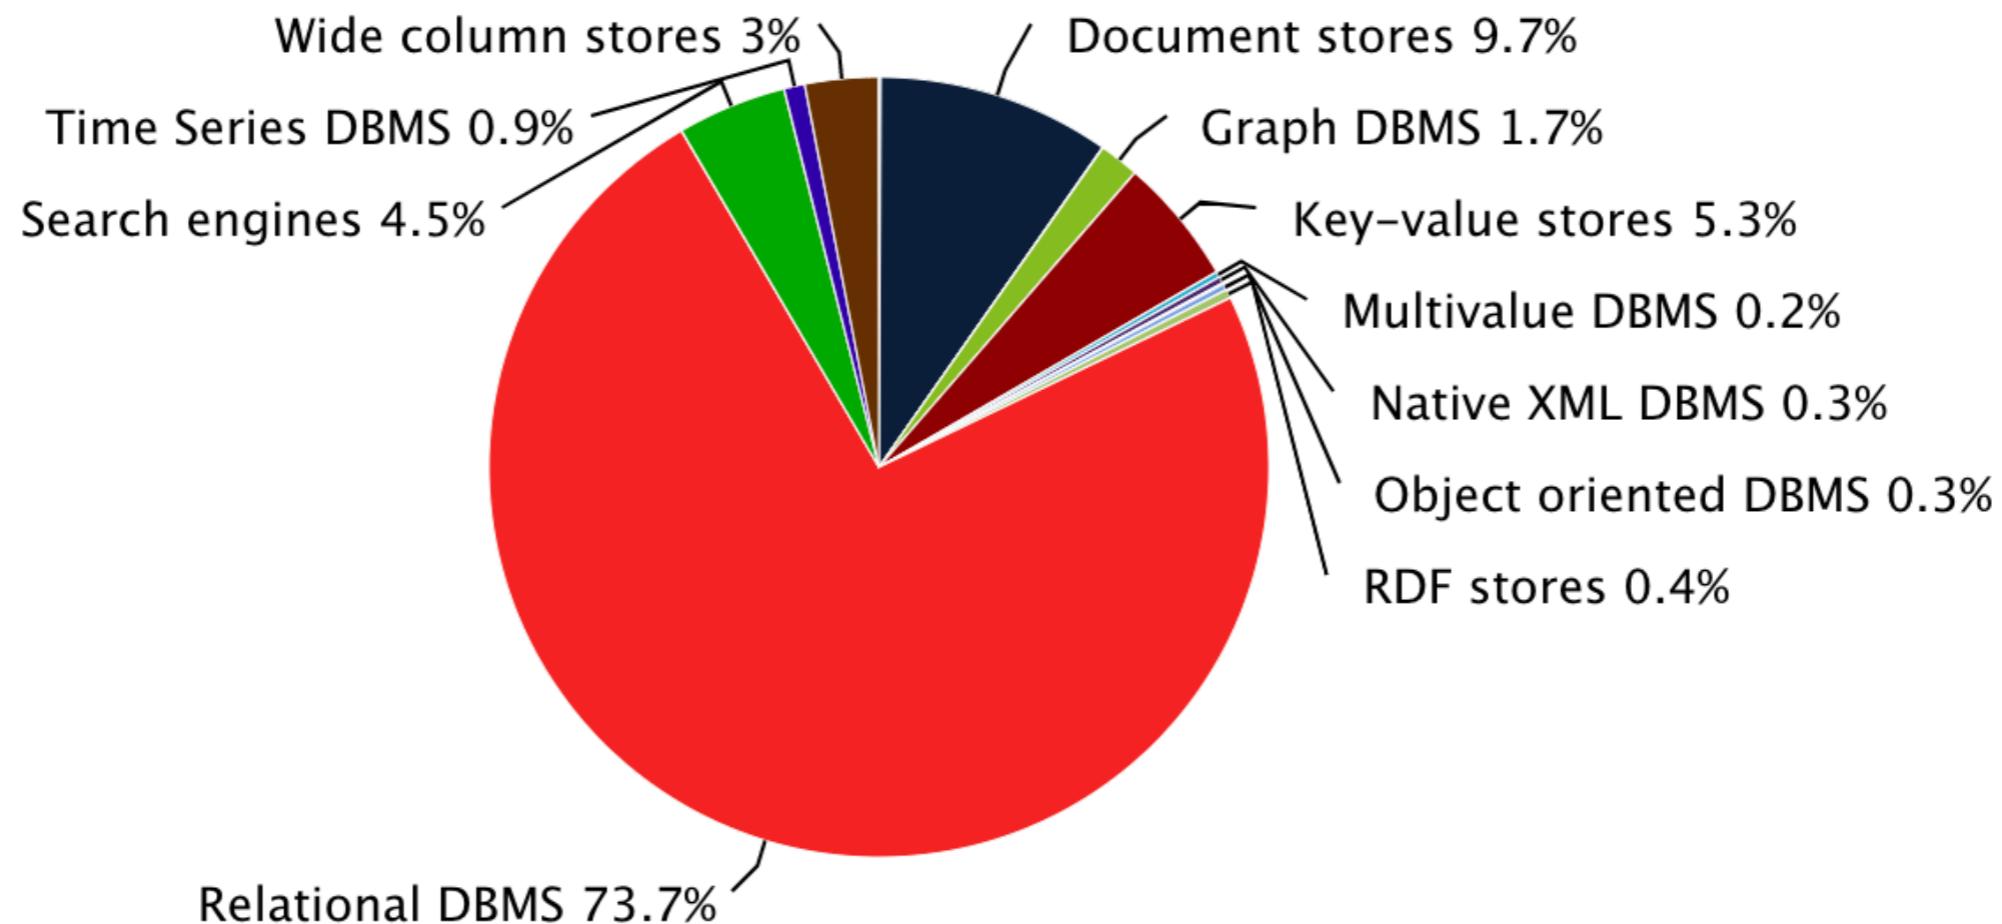
Other database models

Rank			DBMS	Database Model
Dec 2020	Nov 2020	Dec 2019		
1.	1.	1.	Oracle 	Relational, Multi-model 
2.	2.	2.	MySQL 	Relational, Multi-model 
3.	3.	3.	Microsoft SQL Server 	Relational, Multi-model 
4.	4.	4.	PostgreSQL 	Relational, Multi-model 
5.	5.	5.	MongoDB 	Document, Multi-model 
6.	6.	6.	IBM Db2 	Relational, Multi-model 
7.	7.	8.	Redis 	Key-value, Multi-model 
8.	8.	7.	Elasticsearch 	Search engine, Multi-model 
9.	9.	11.	SQLite 	Relational
10.	10.	10.	Cassandra 	Wide column
11.	11.	9.	Microsoft Access	Relational
12.	12.	13.	MariaDB 	Relational, Multi-model 
13.	13.	12.	Splunk	Search engine
14.	14.	15.	Teradata 	Relational, Multi-model 
15.	15.	14.	Hive	Relational

<https://db-engines.com/en/ranking>



Database models



<https://db-engines.com/en/ranking>



Why ?



ACID ?

Atomicity
Consistency
Isolation
Durability



SQL ?

Structure Query Language



SQL

SQL uses human-readable syntax
Used to access data in RDBMS



SQL

customer_id	fname	lname	last_purchase_date	status_level
1234	Alice	Smith	23-Jan-2016	silver
2345	Bob	Washington	14-July-2018	bronze
2352	Charlie	Johnson	20-Dec-2018	gold
2678	Chendong	Mao	19-Dec-2017	gold
3422	Javier	Valdez	23-Dec-2018	gold
3577	Avery	Winston	5-Jan-2019	silver
4240	Charles	Davis	15-Jan-2017	bronze
6235	Vihaan	Patel	12-Nov-2018	gold
2352	Katherine	Coltrane	04-Apr-2016	silver



```
SELECT *
FROM SOME_TABLE
WHERE 1=1
HAVING ...
GROUP BY ...
```



History



1980

1990

2000

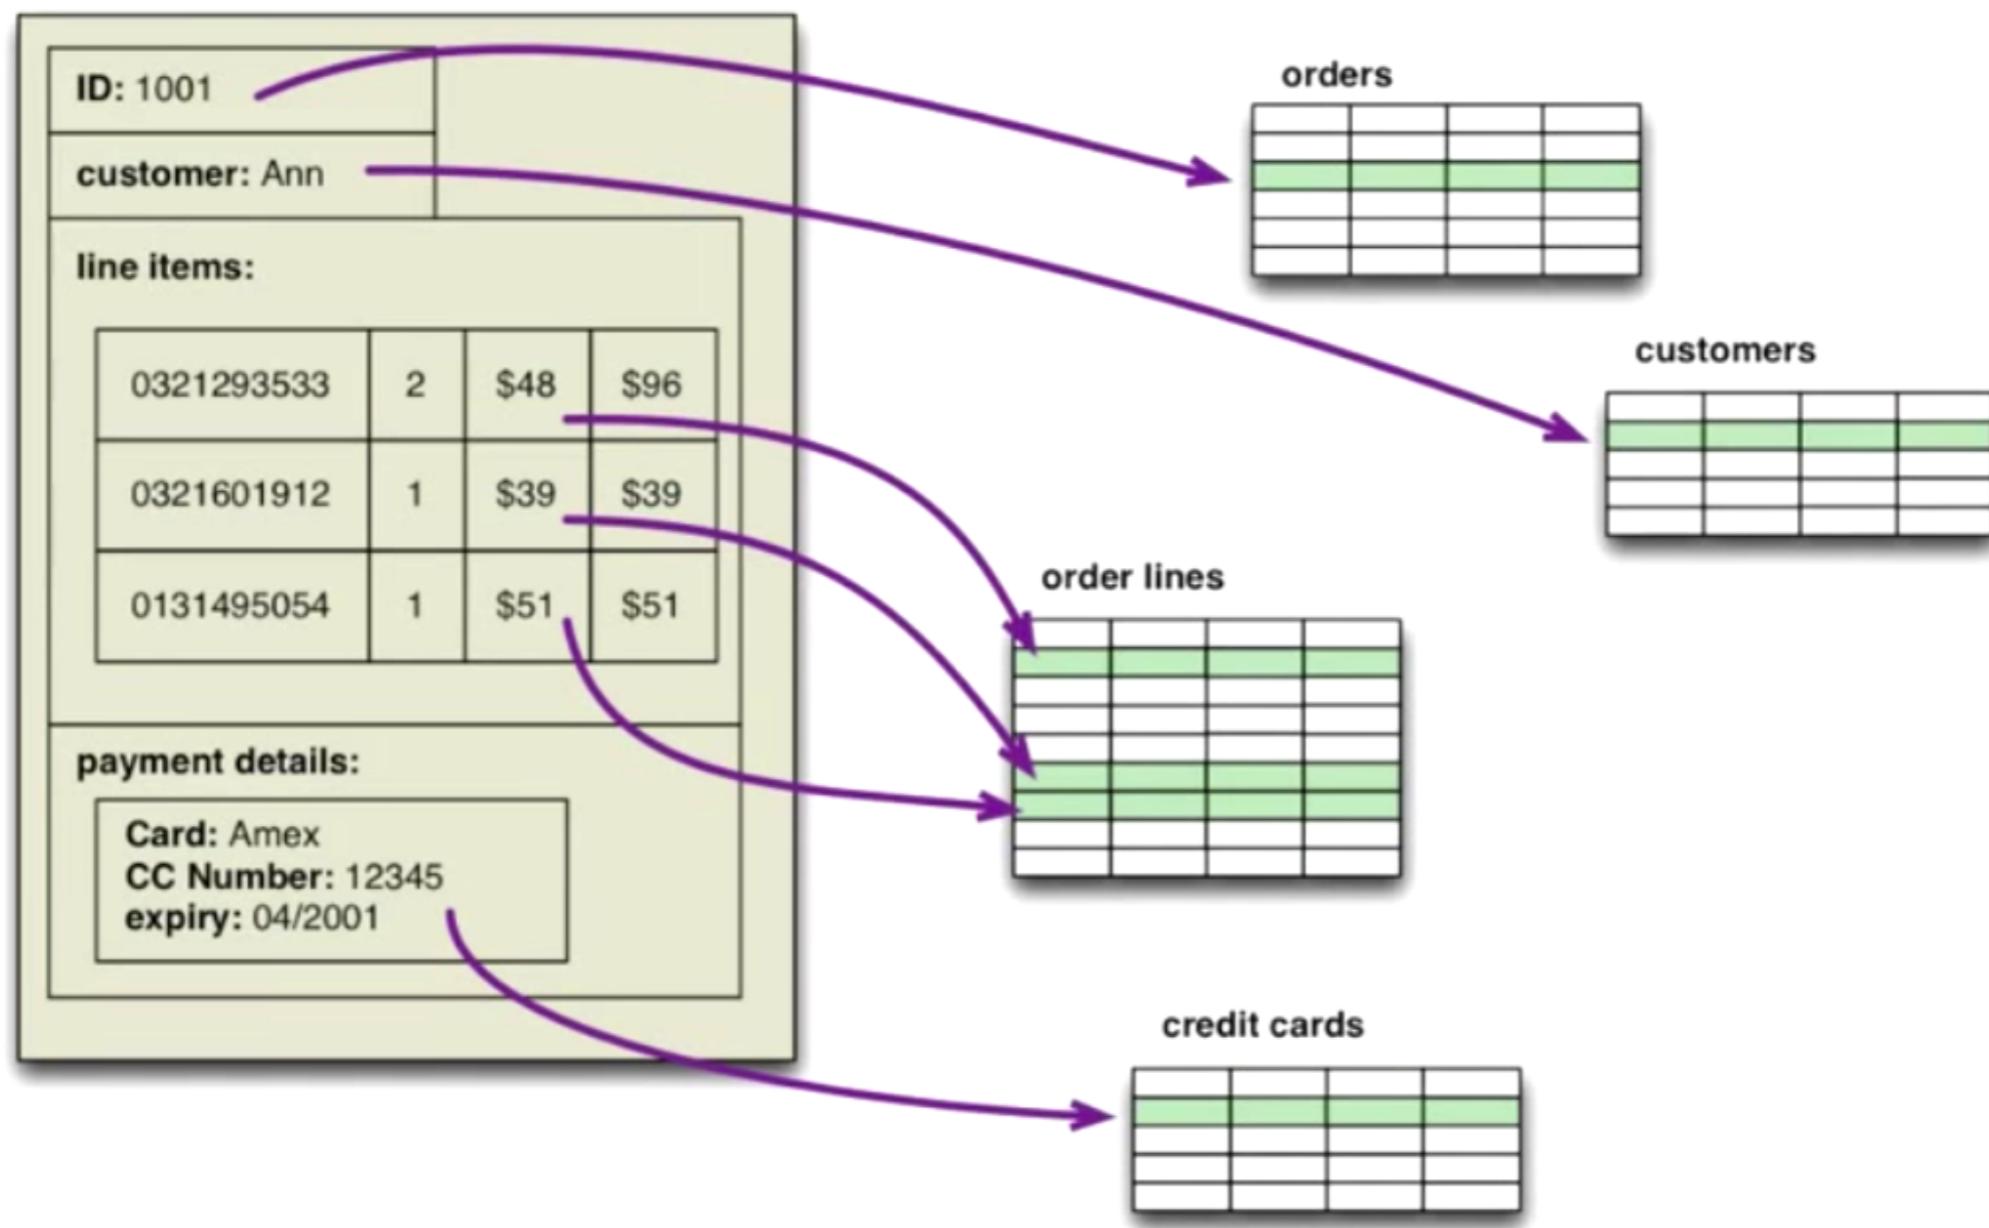
2010



RDBMS

Persistence
Integration
SQL
Transaction

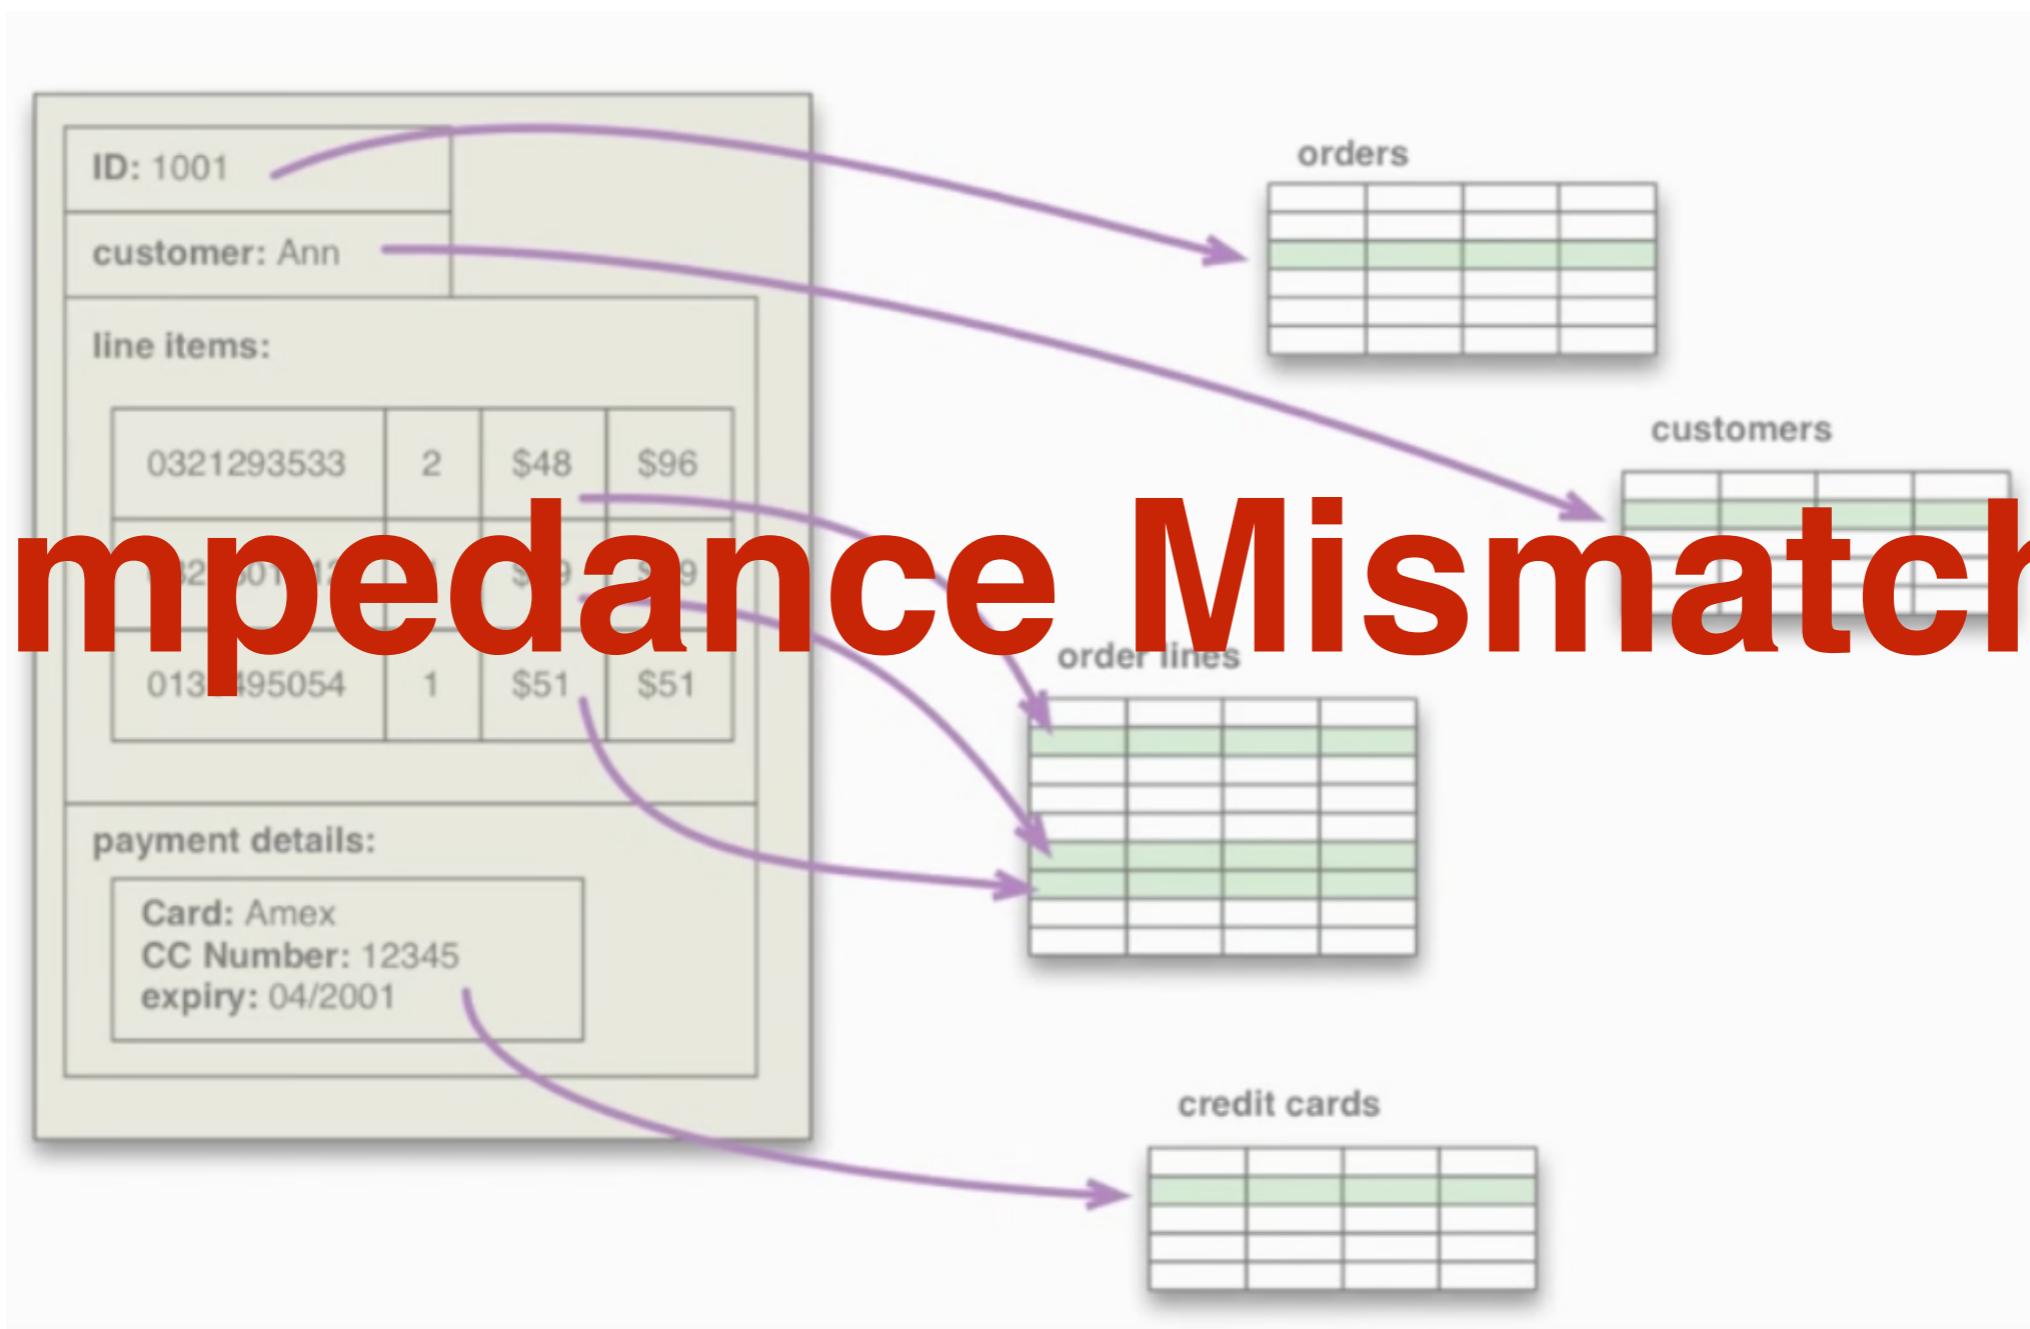




Object Oriented Programming

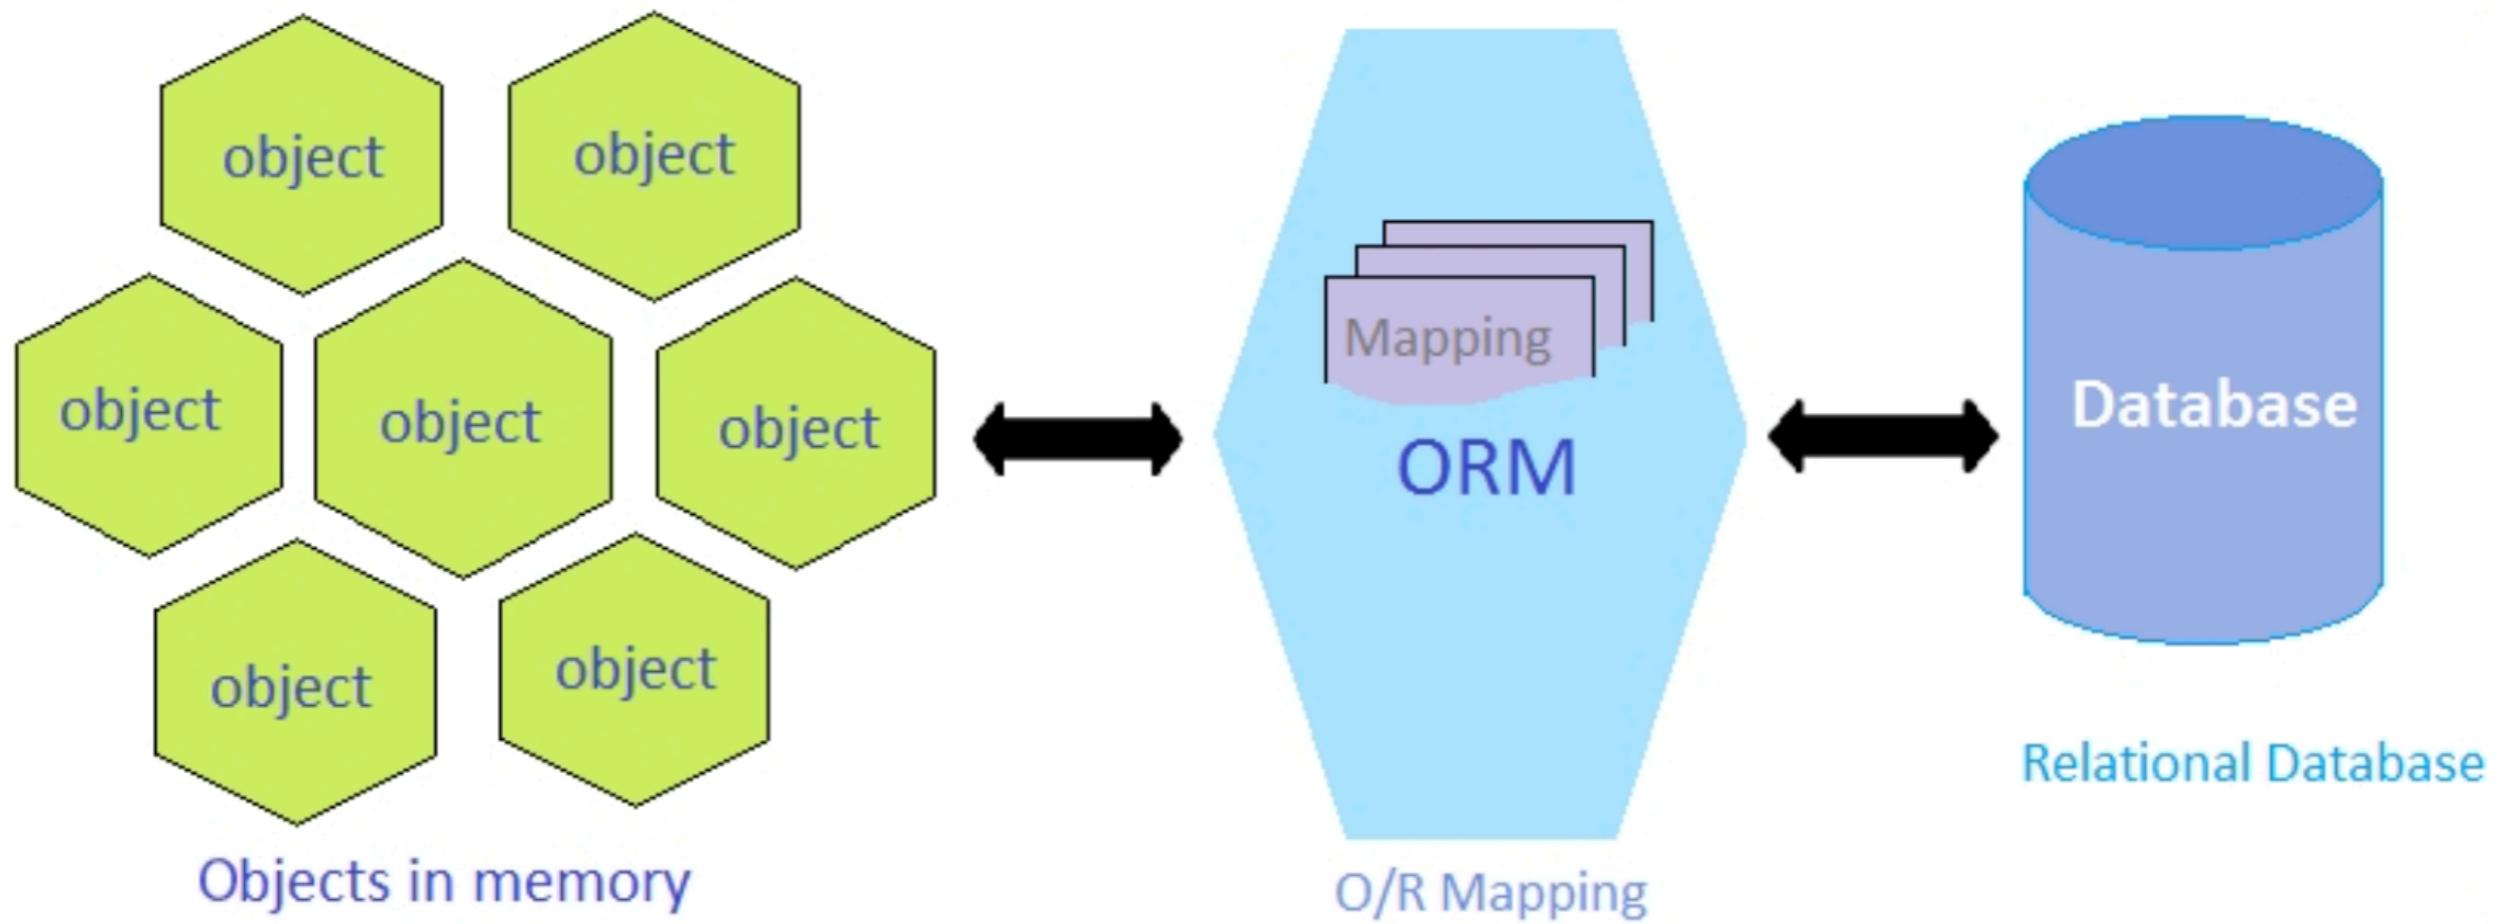


Impedance Mismatch



Object Relation Mapping







JPA

Java Persistence API



1980

1990

2000

2010



Object Database



Back to RDBMS and ORM



1980

1990

2000

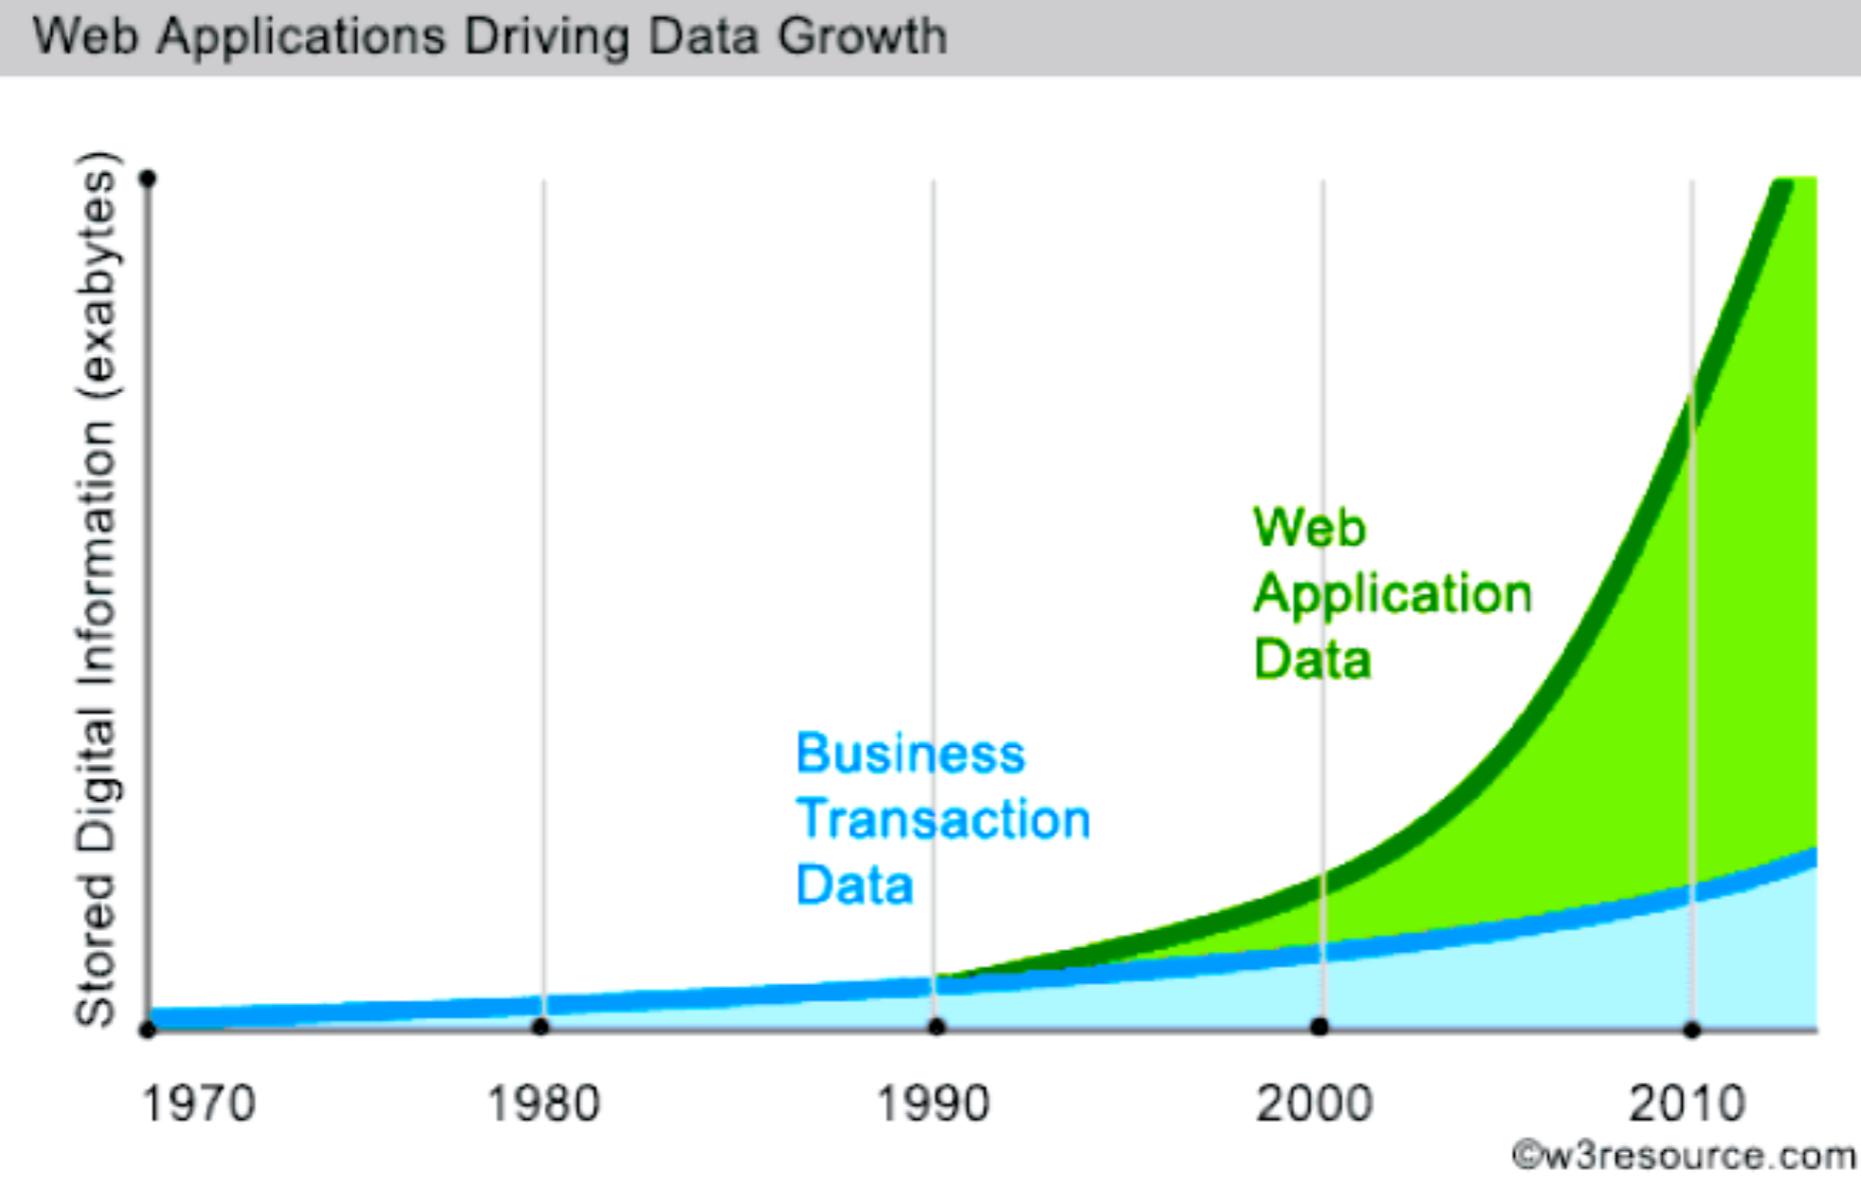
2010



Problem !!



Web Application + Internet

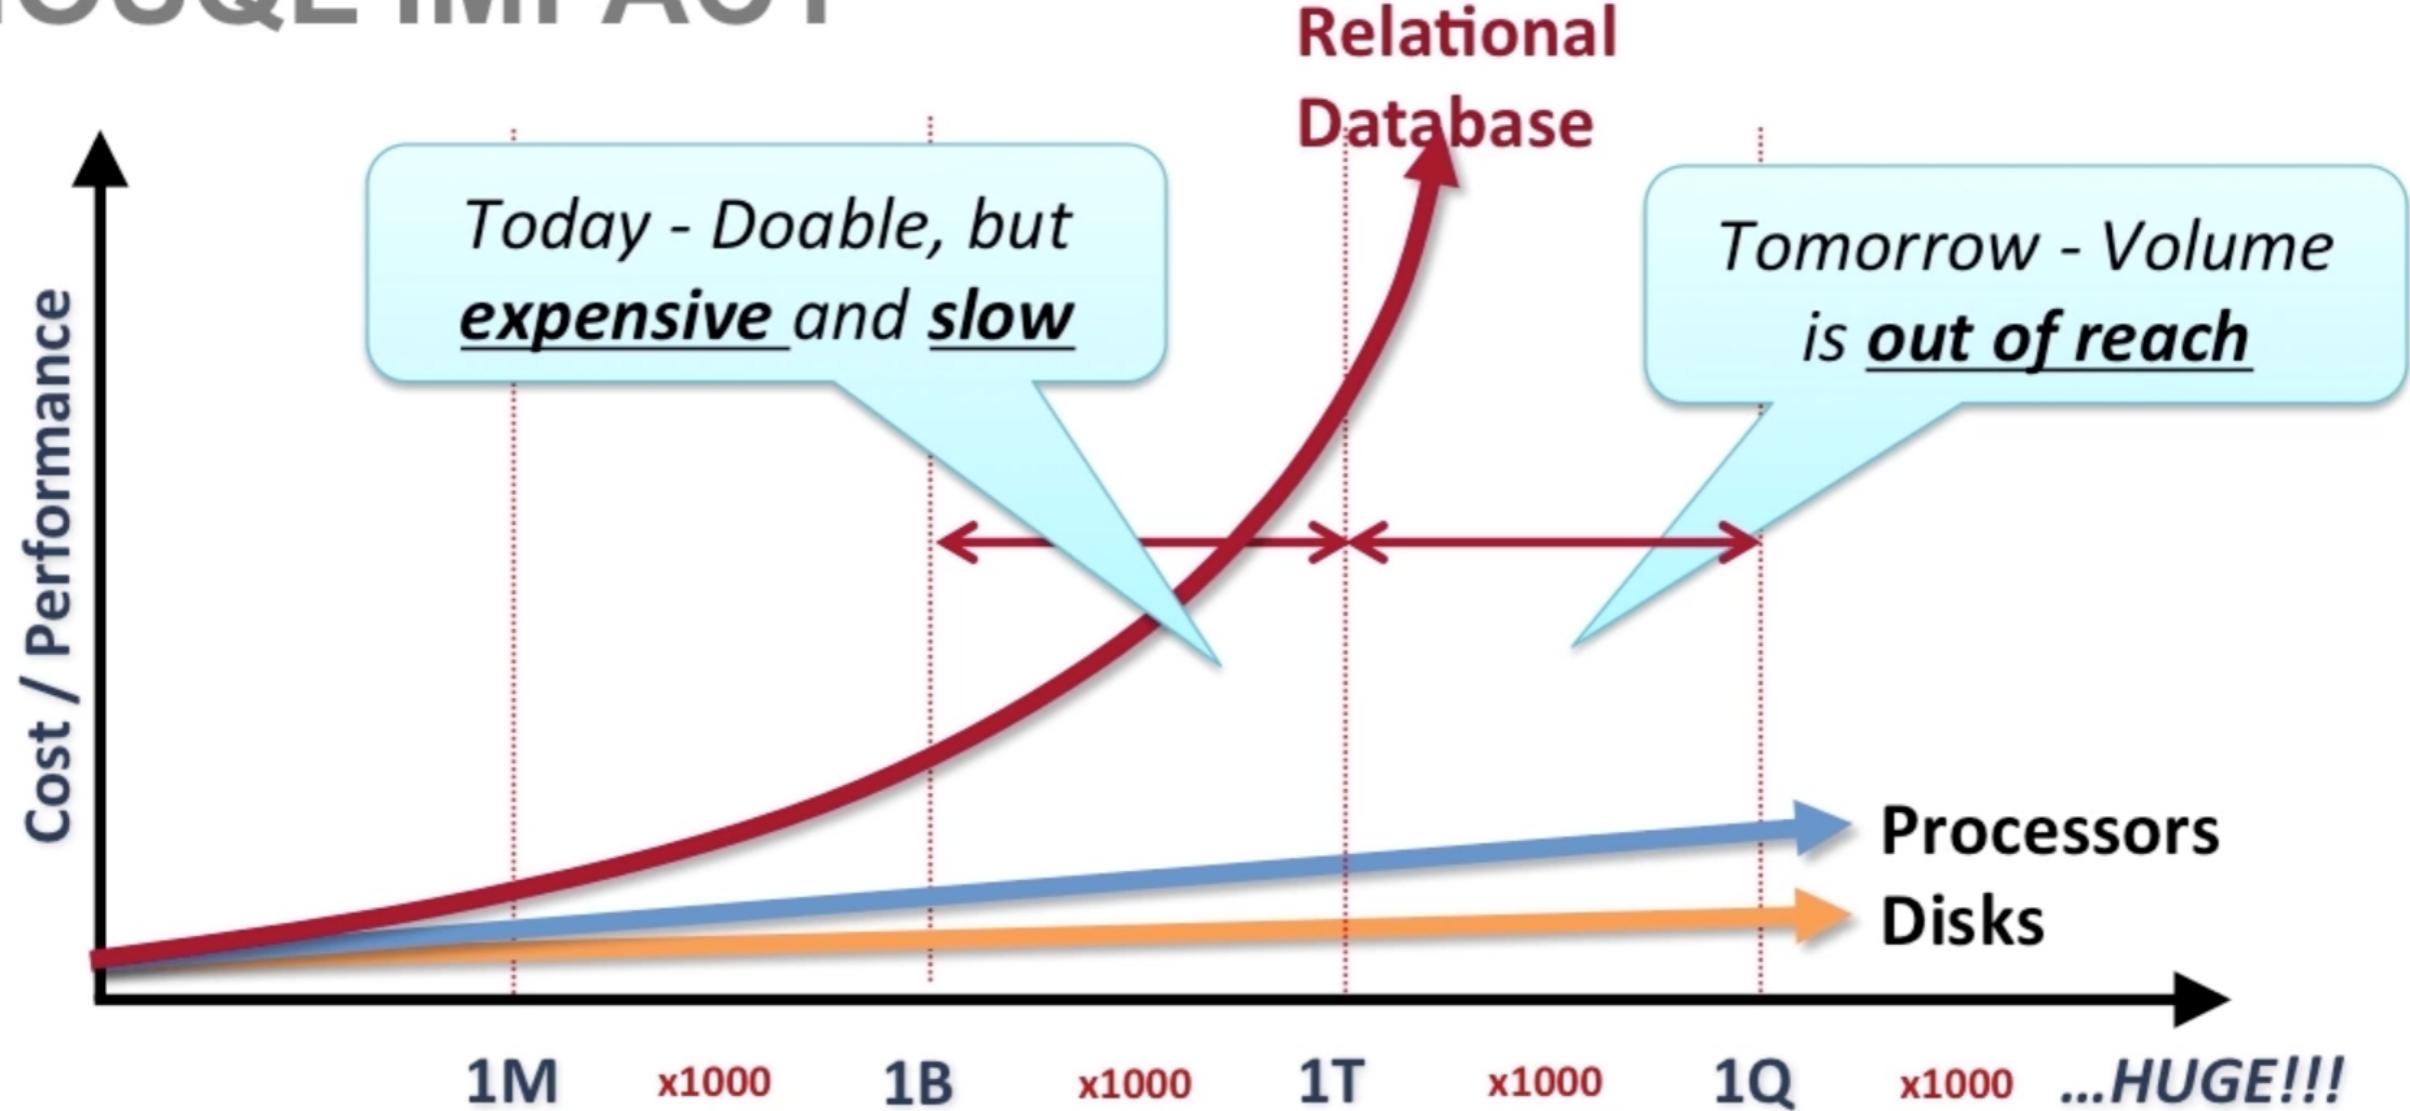


Data challenges

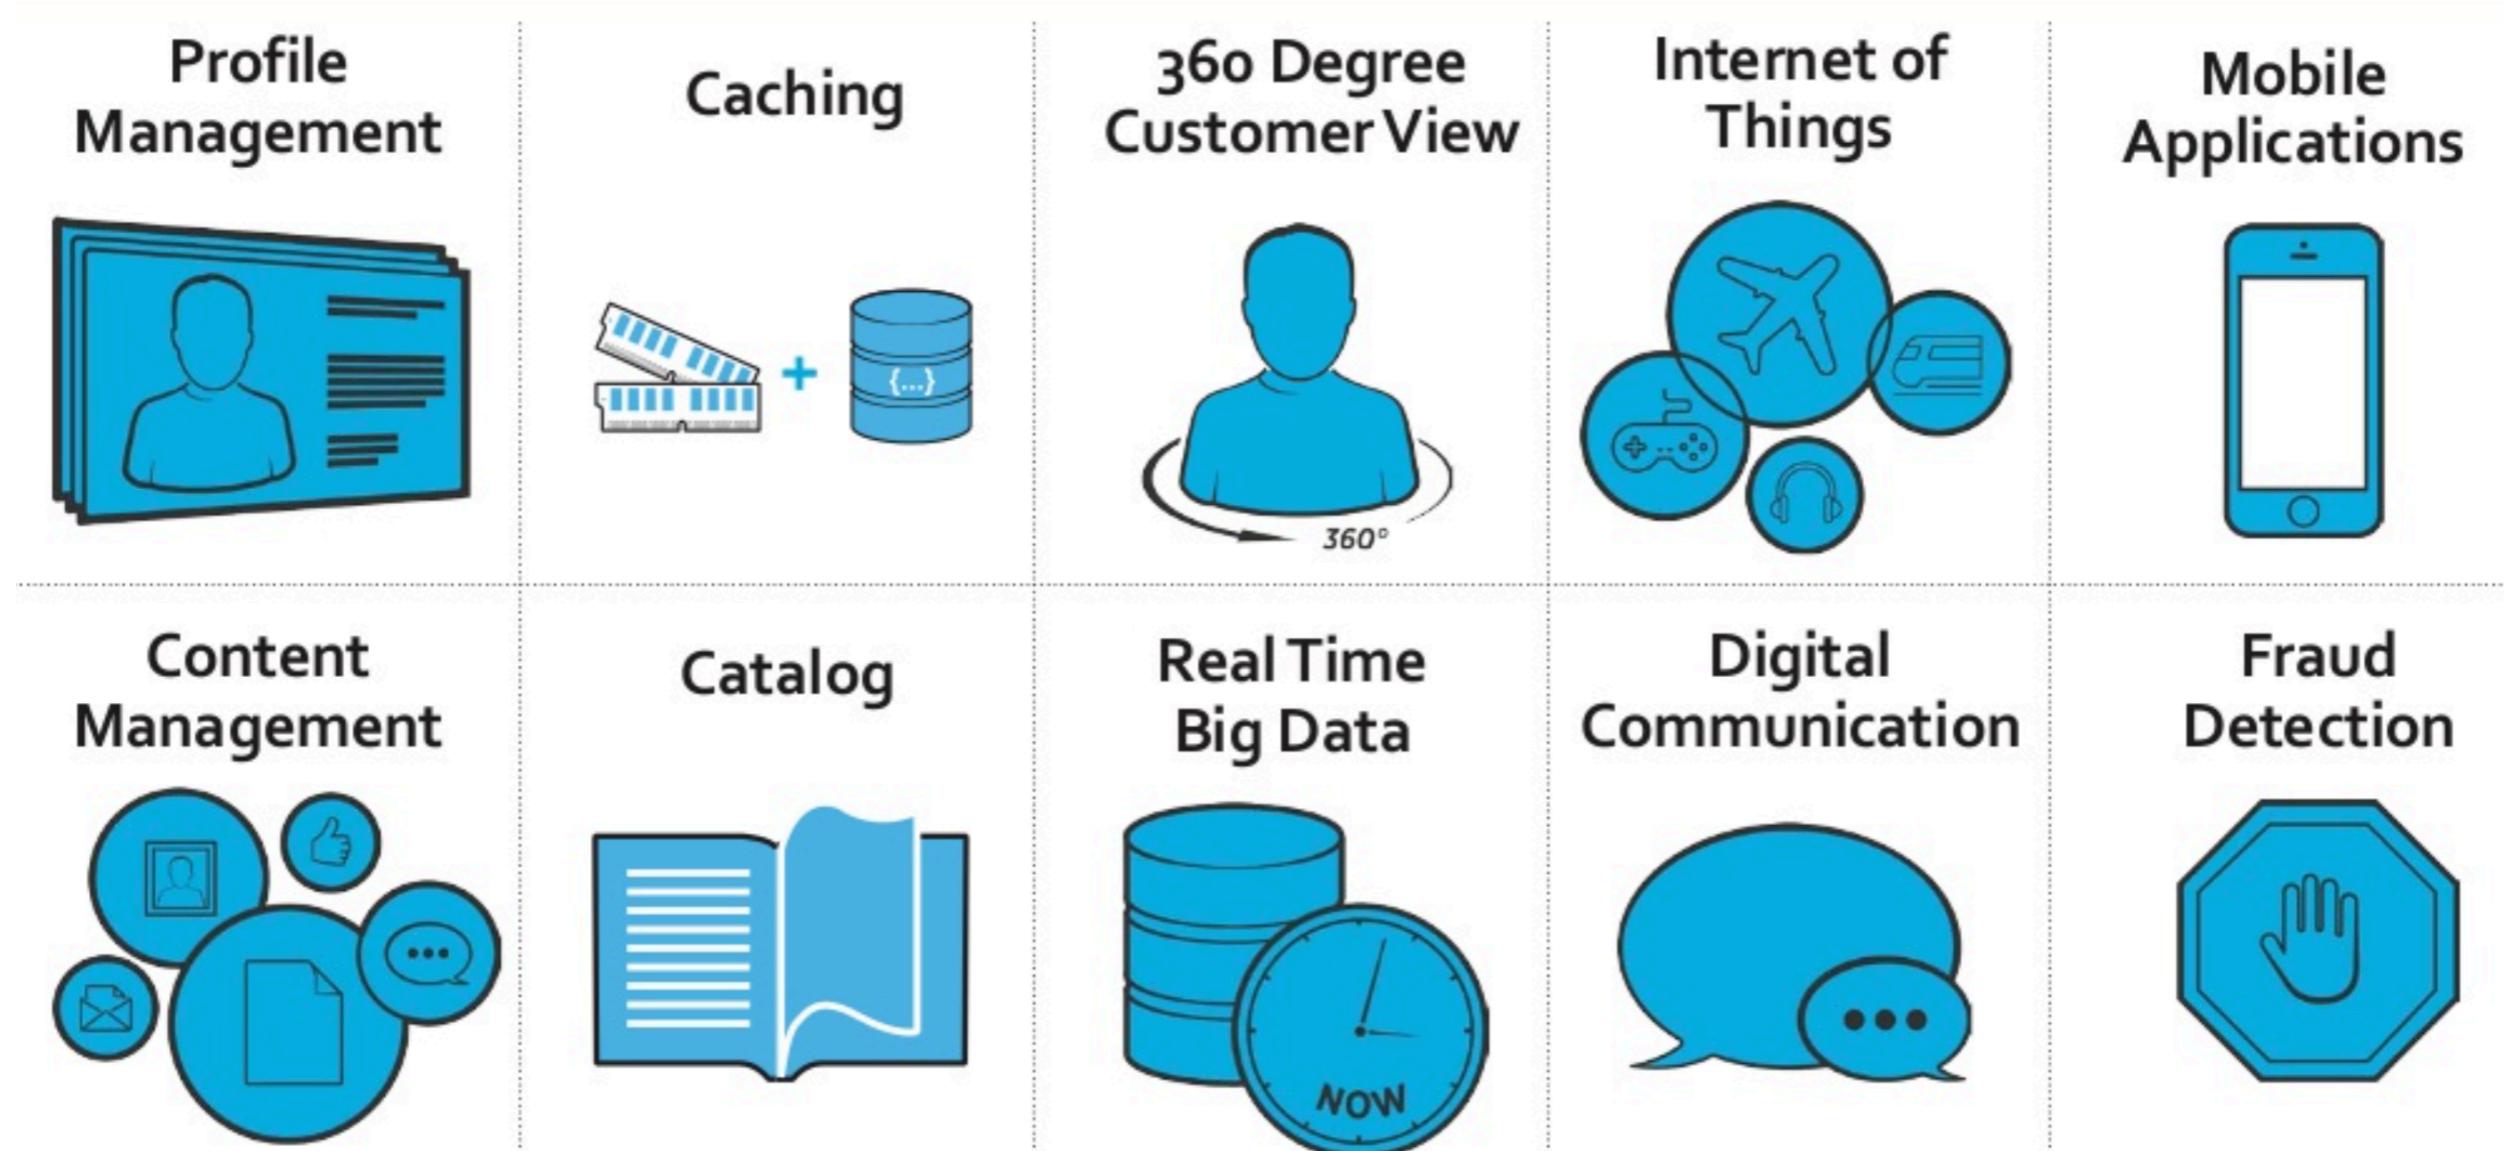
Data volume
Data velocity
Data variety
Data valence



NOSQL IMPACT



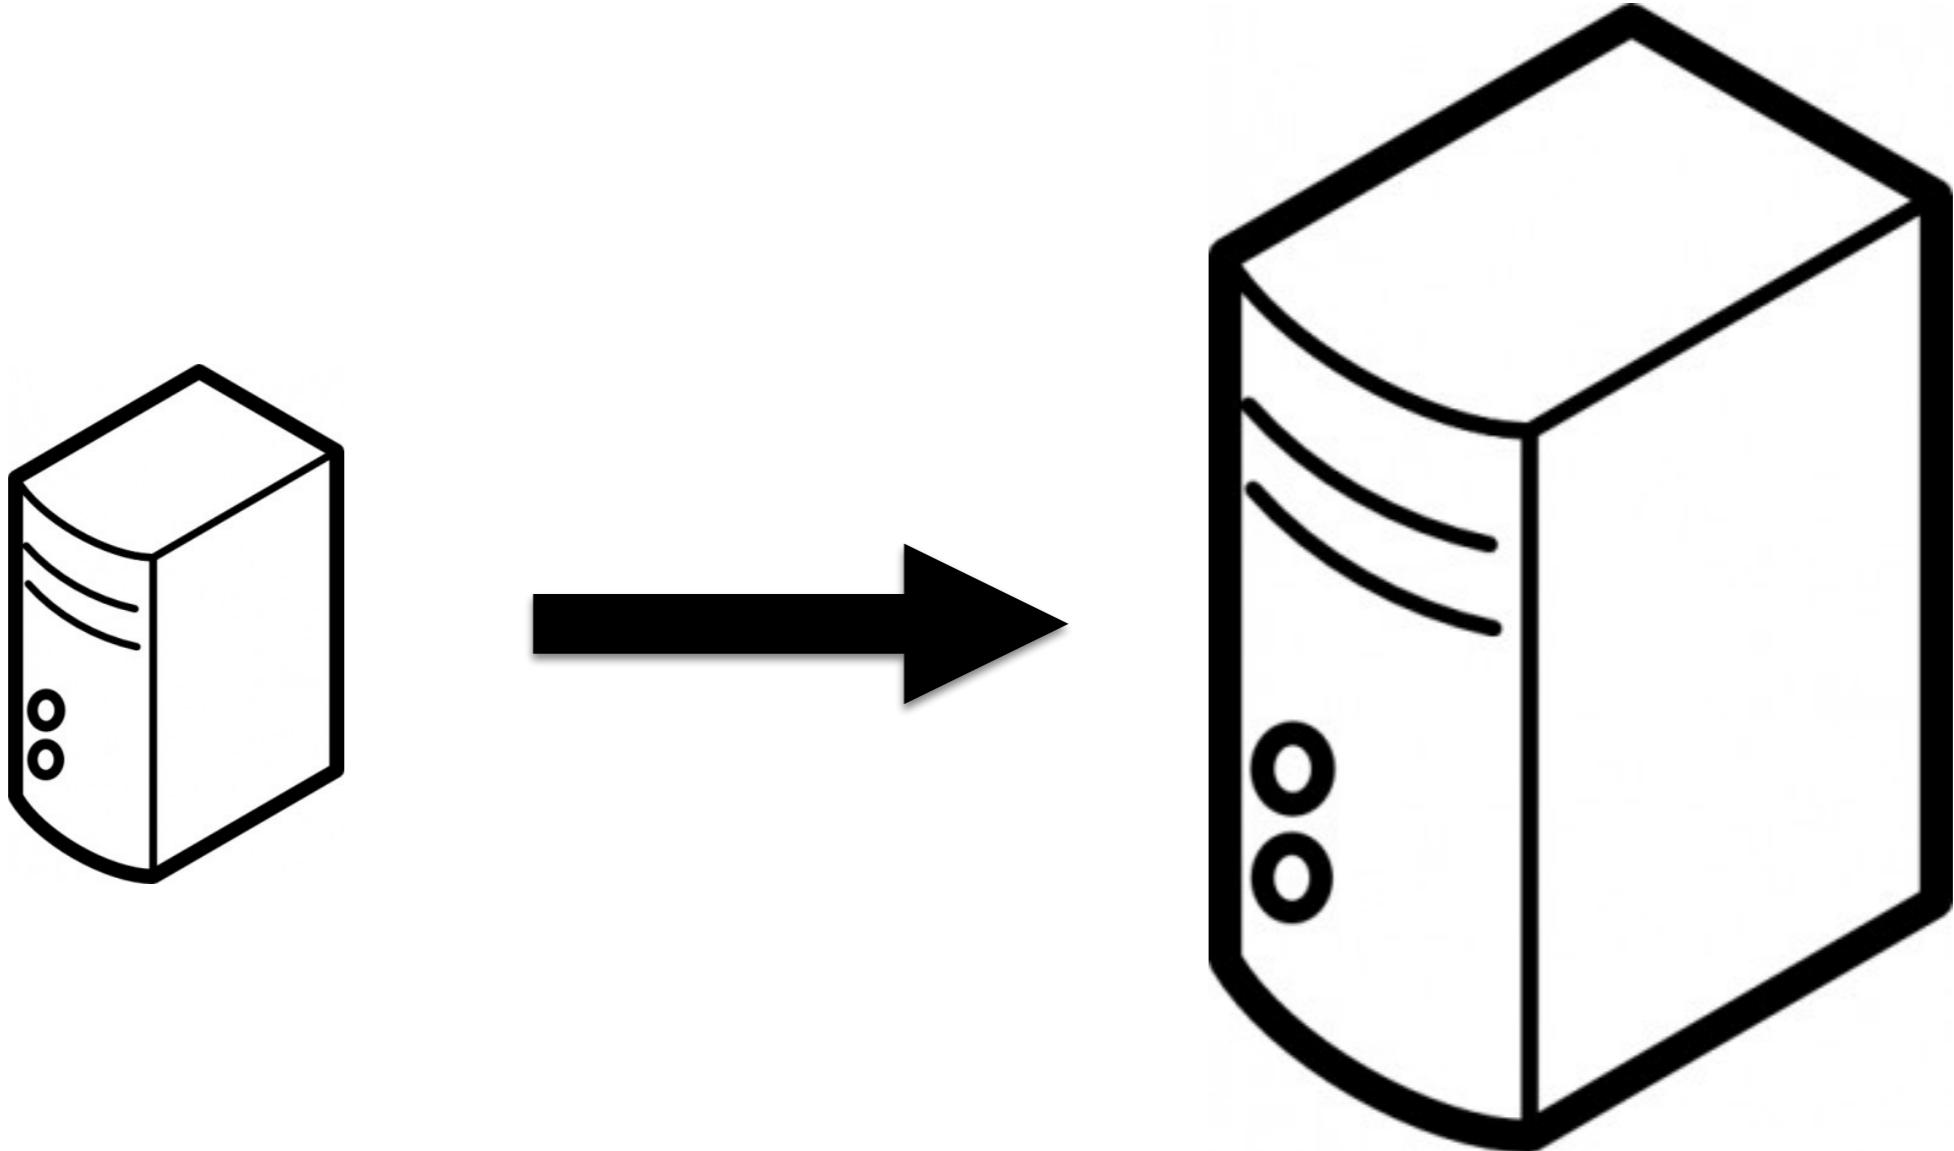
Requirements !!



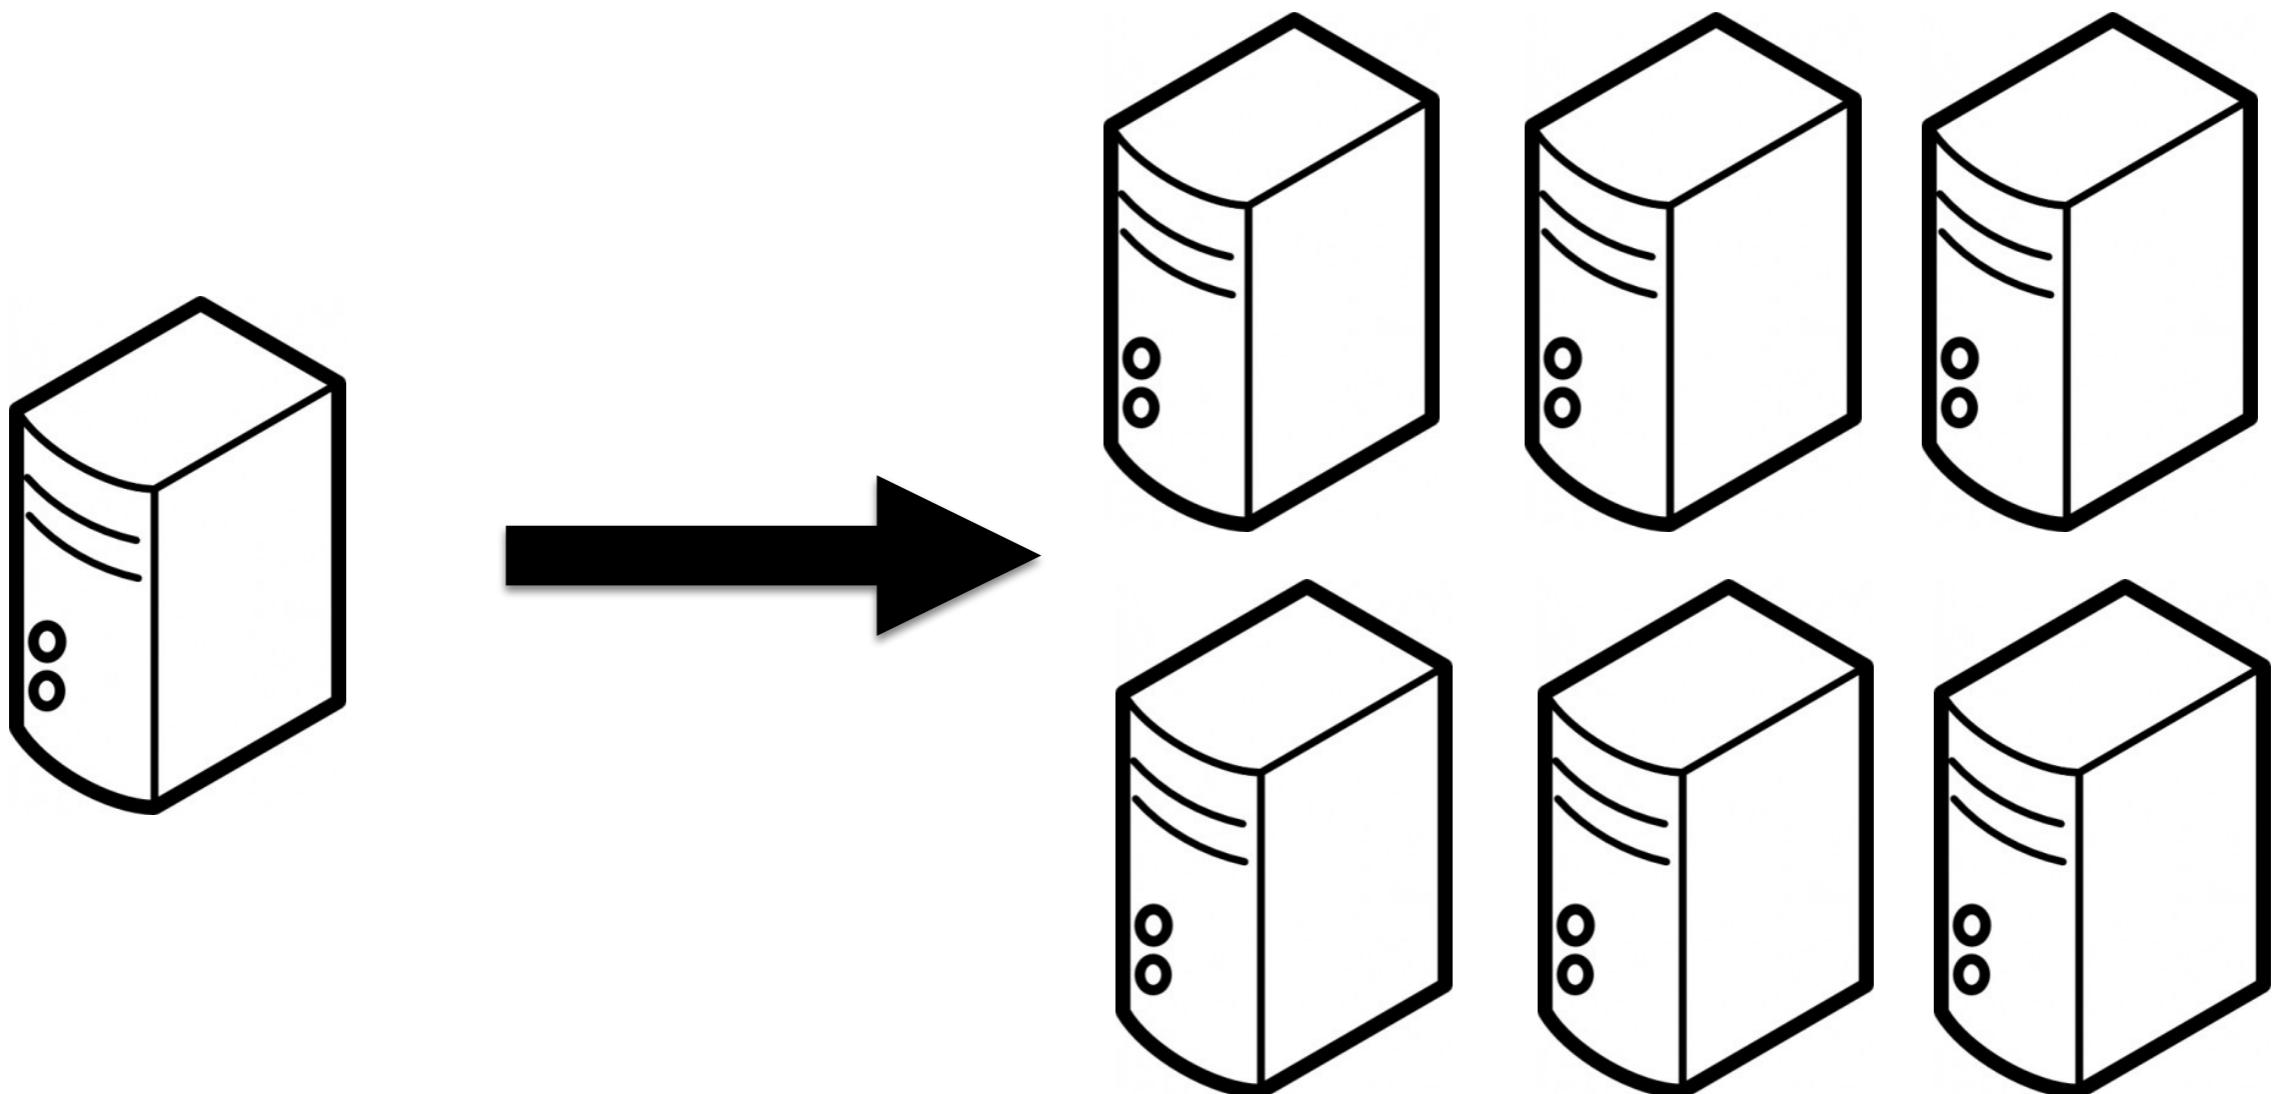
How to scale RDBMS ?



Scale up

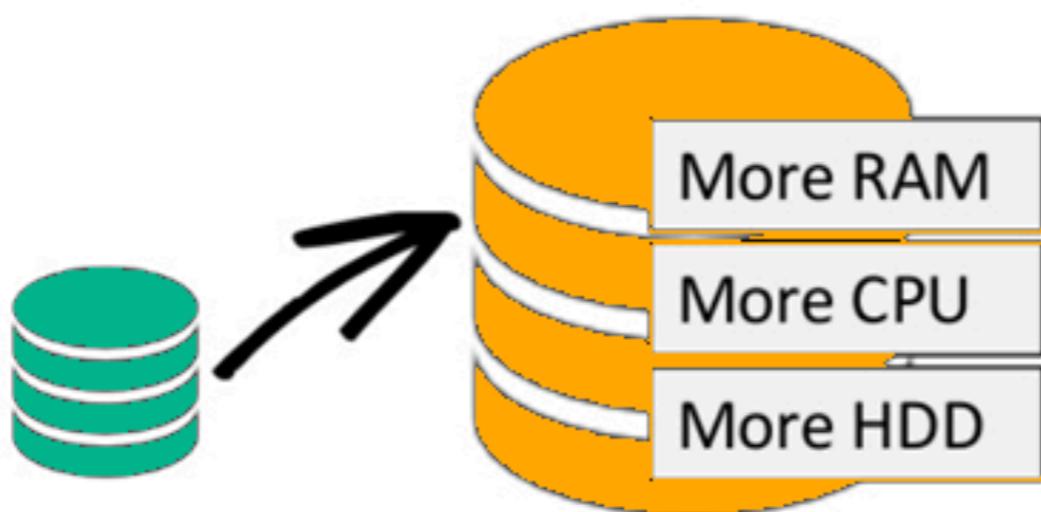


Scale out

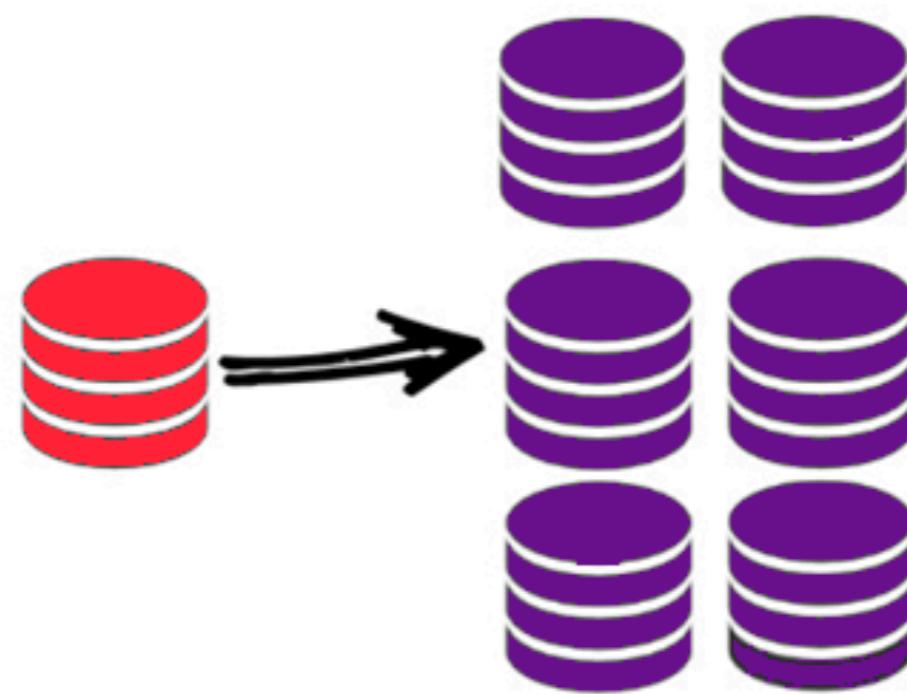


Scale up vs out ?

Scale-Up (vertical scaling):

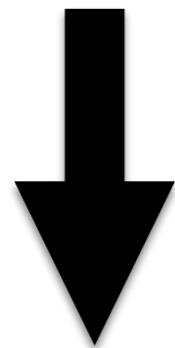


Scale-Out (horizontal scaling):



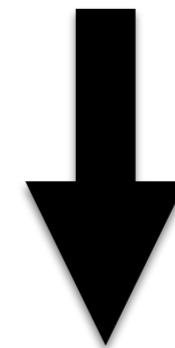


Google



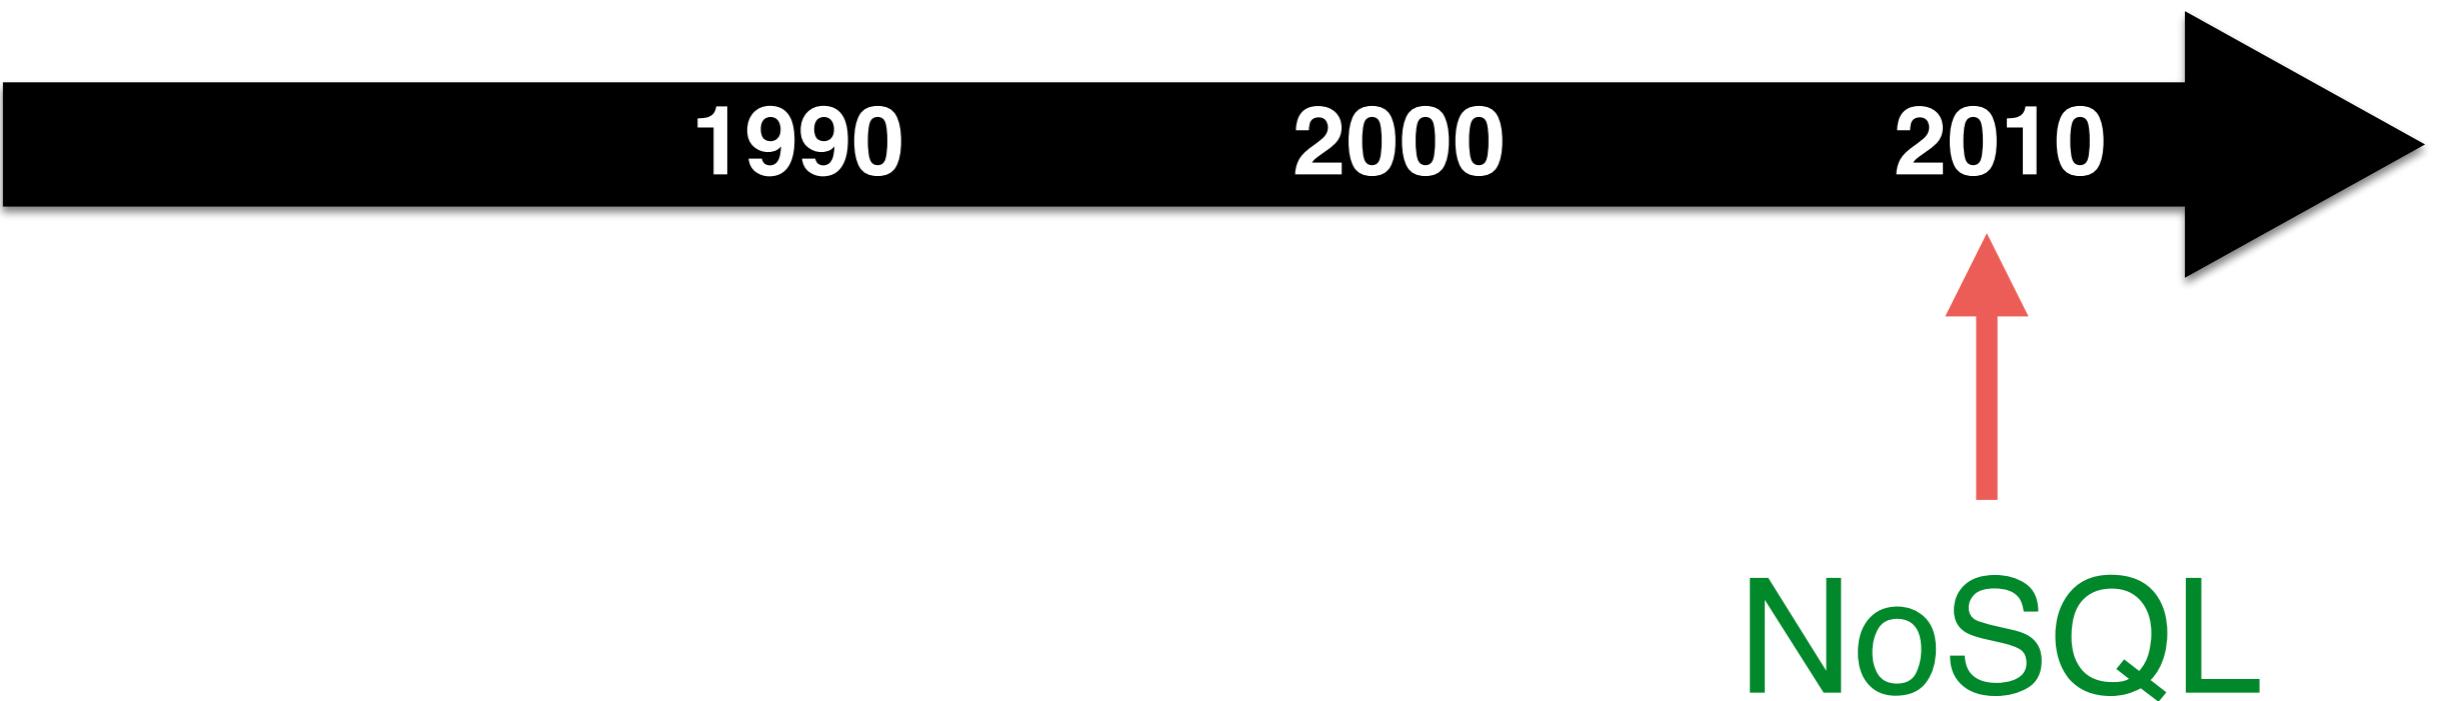
Bigtable

amazon



Dynamo





NoSQL

No SQL

No No SQL

Not Only SQL

Non-Relational



NoSQL properties

Non-relational
Opensource
Cluster friendly
Schema-less
For new web app

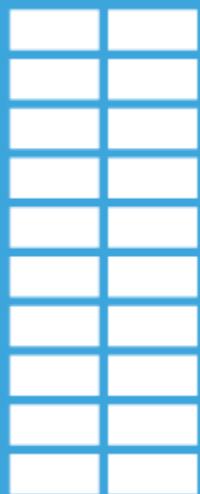


Data models

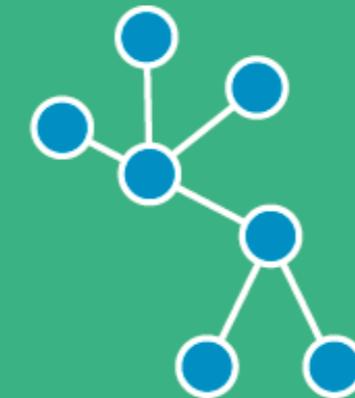
Key-value
Document
Column
Graph
Search
Time-series



Key-Value



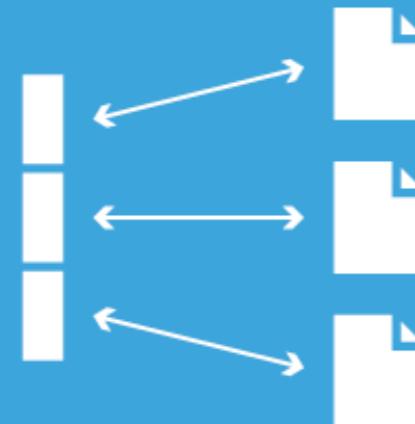
Graph DB



Column Family

1			1
	1	1	
			1
1		1	
1			1
	1		1
			1

Document





APACHE
HBASE

Column Family



Couchbase



mongoDB

Document



Graph



Project Voldemort
A distributed database.

Key-Value

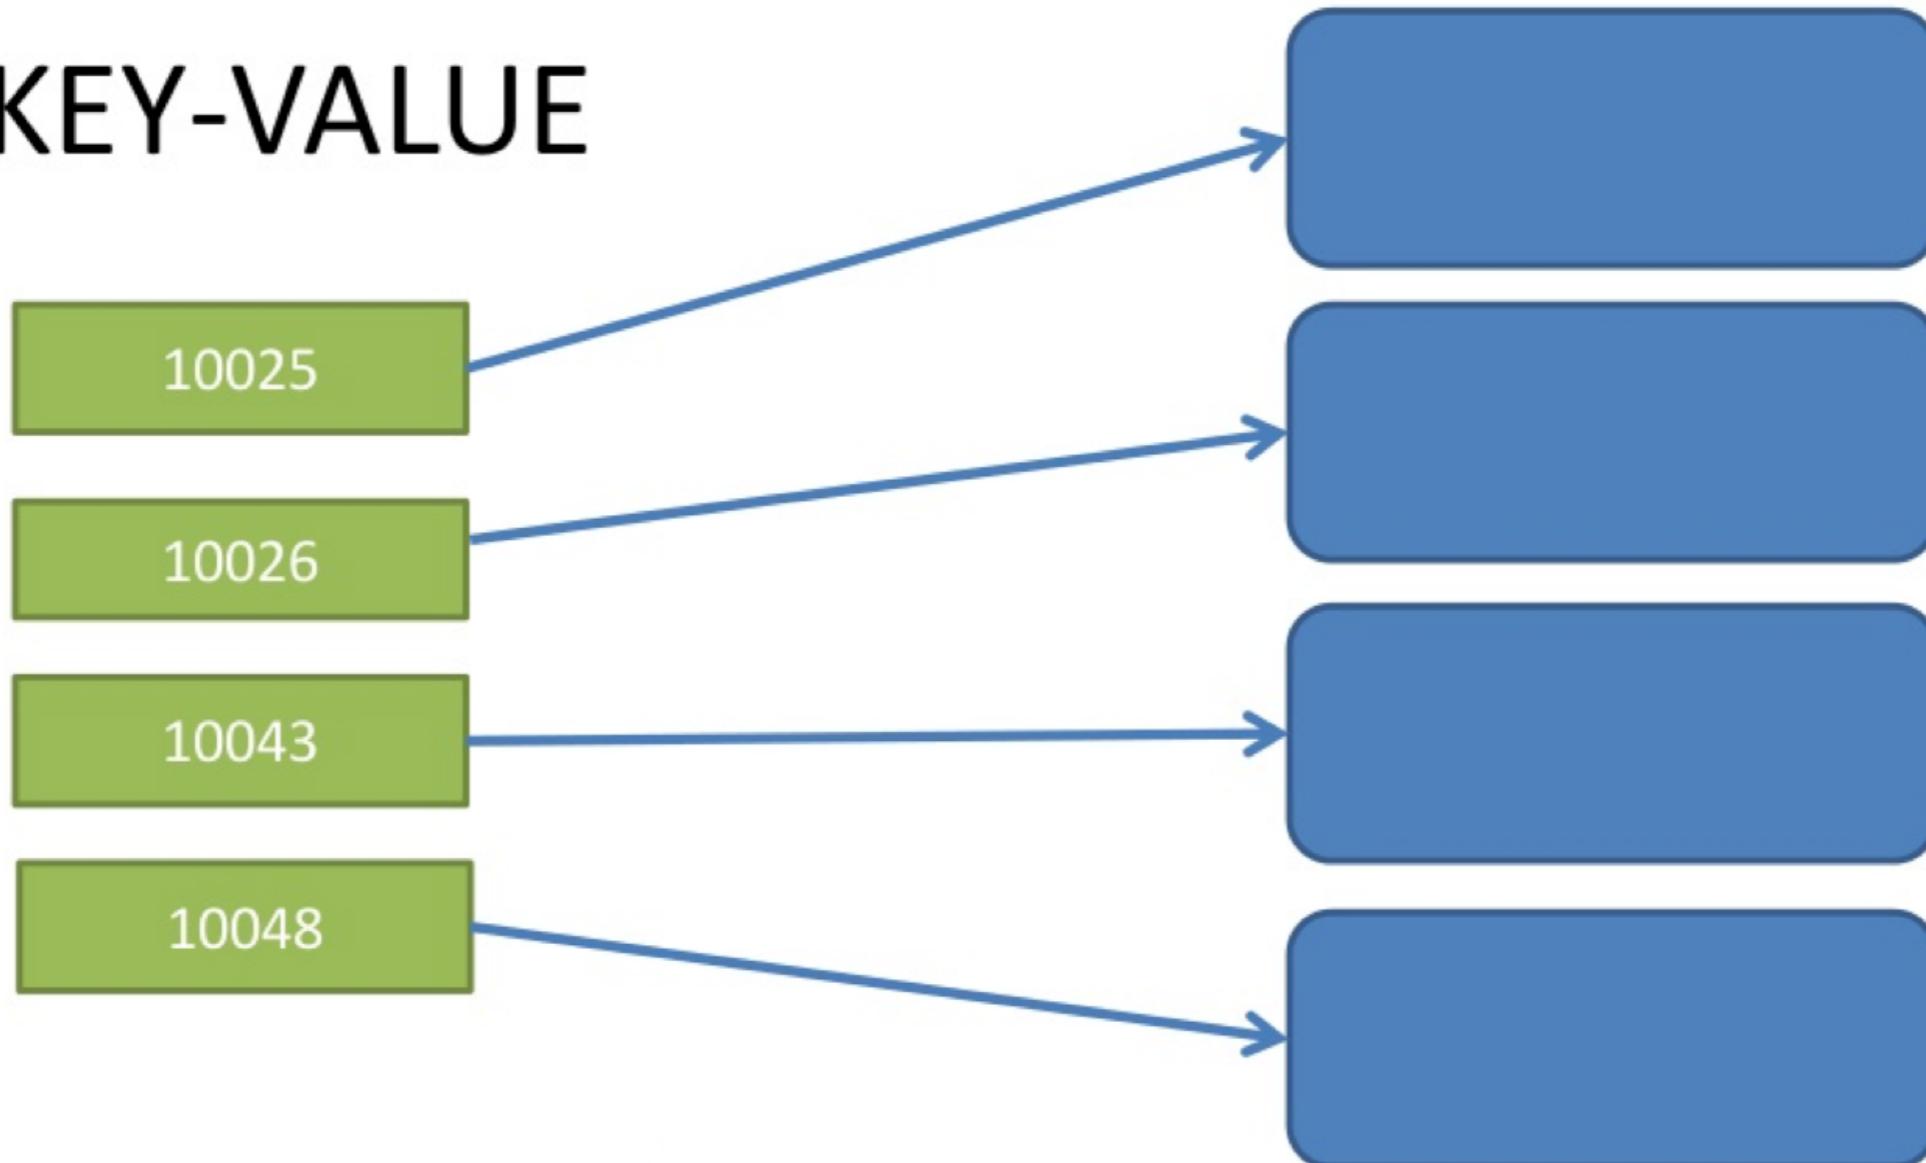


redis

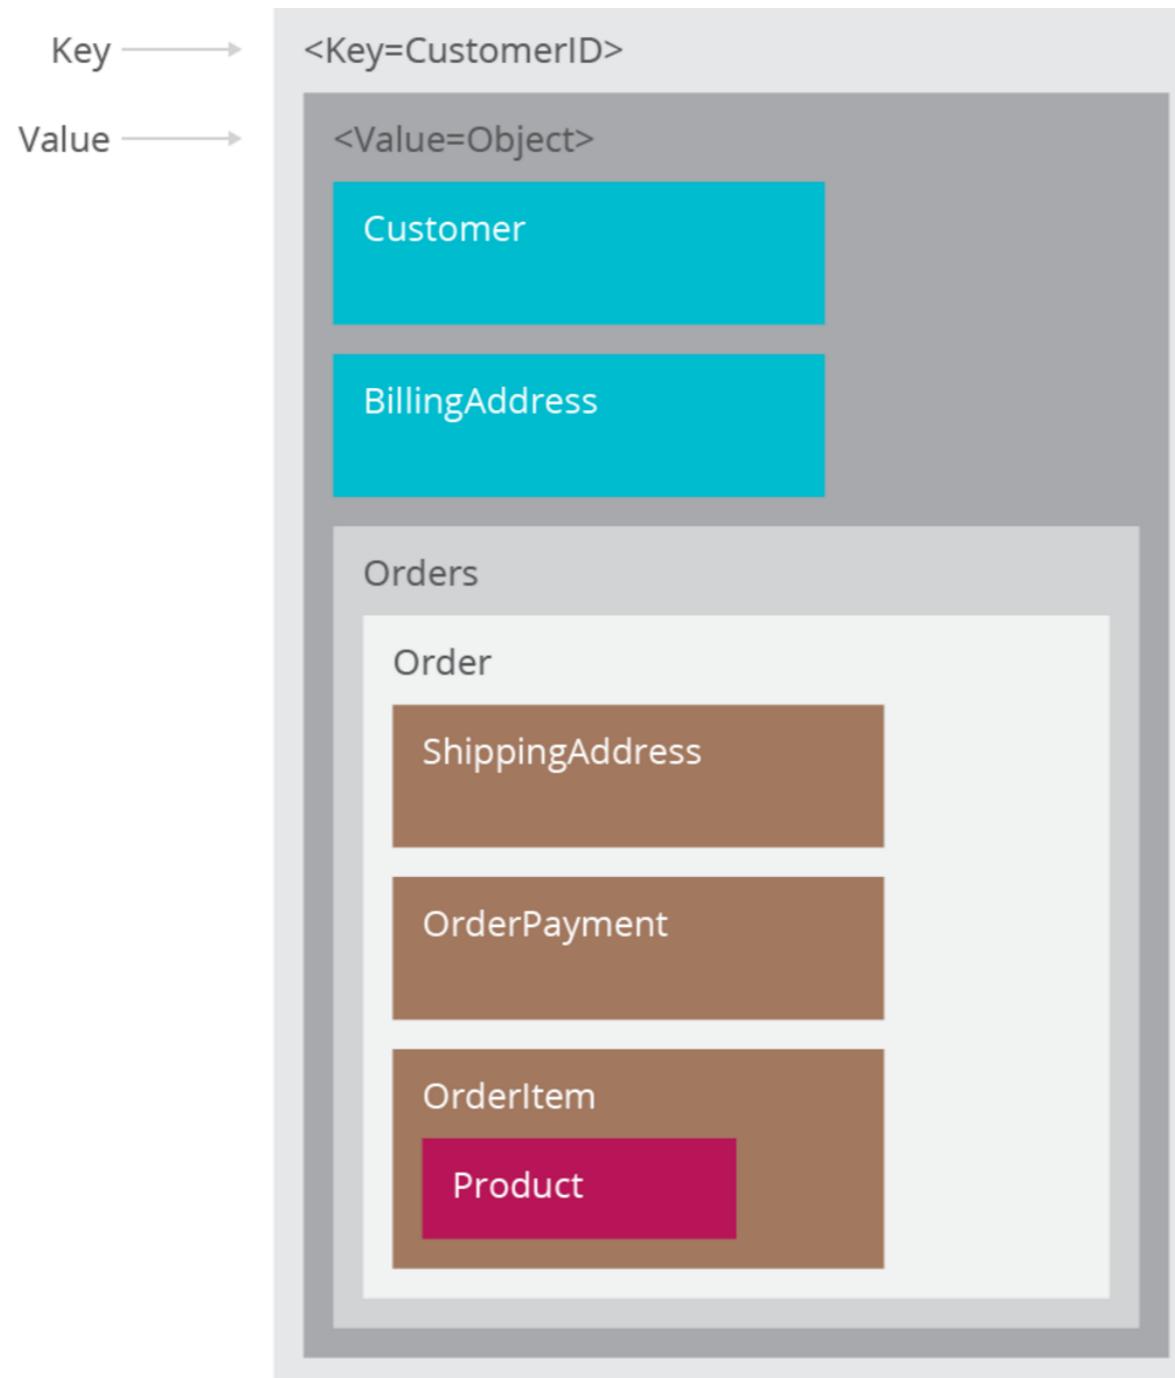


Key-value

KEY-VALUE

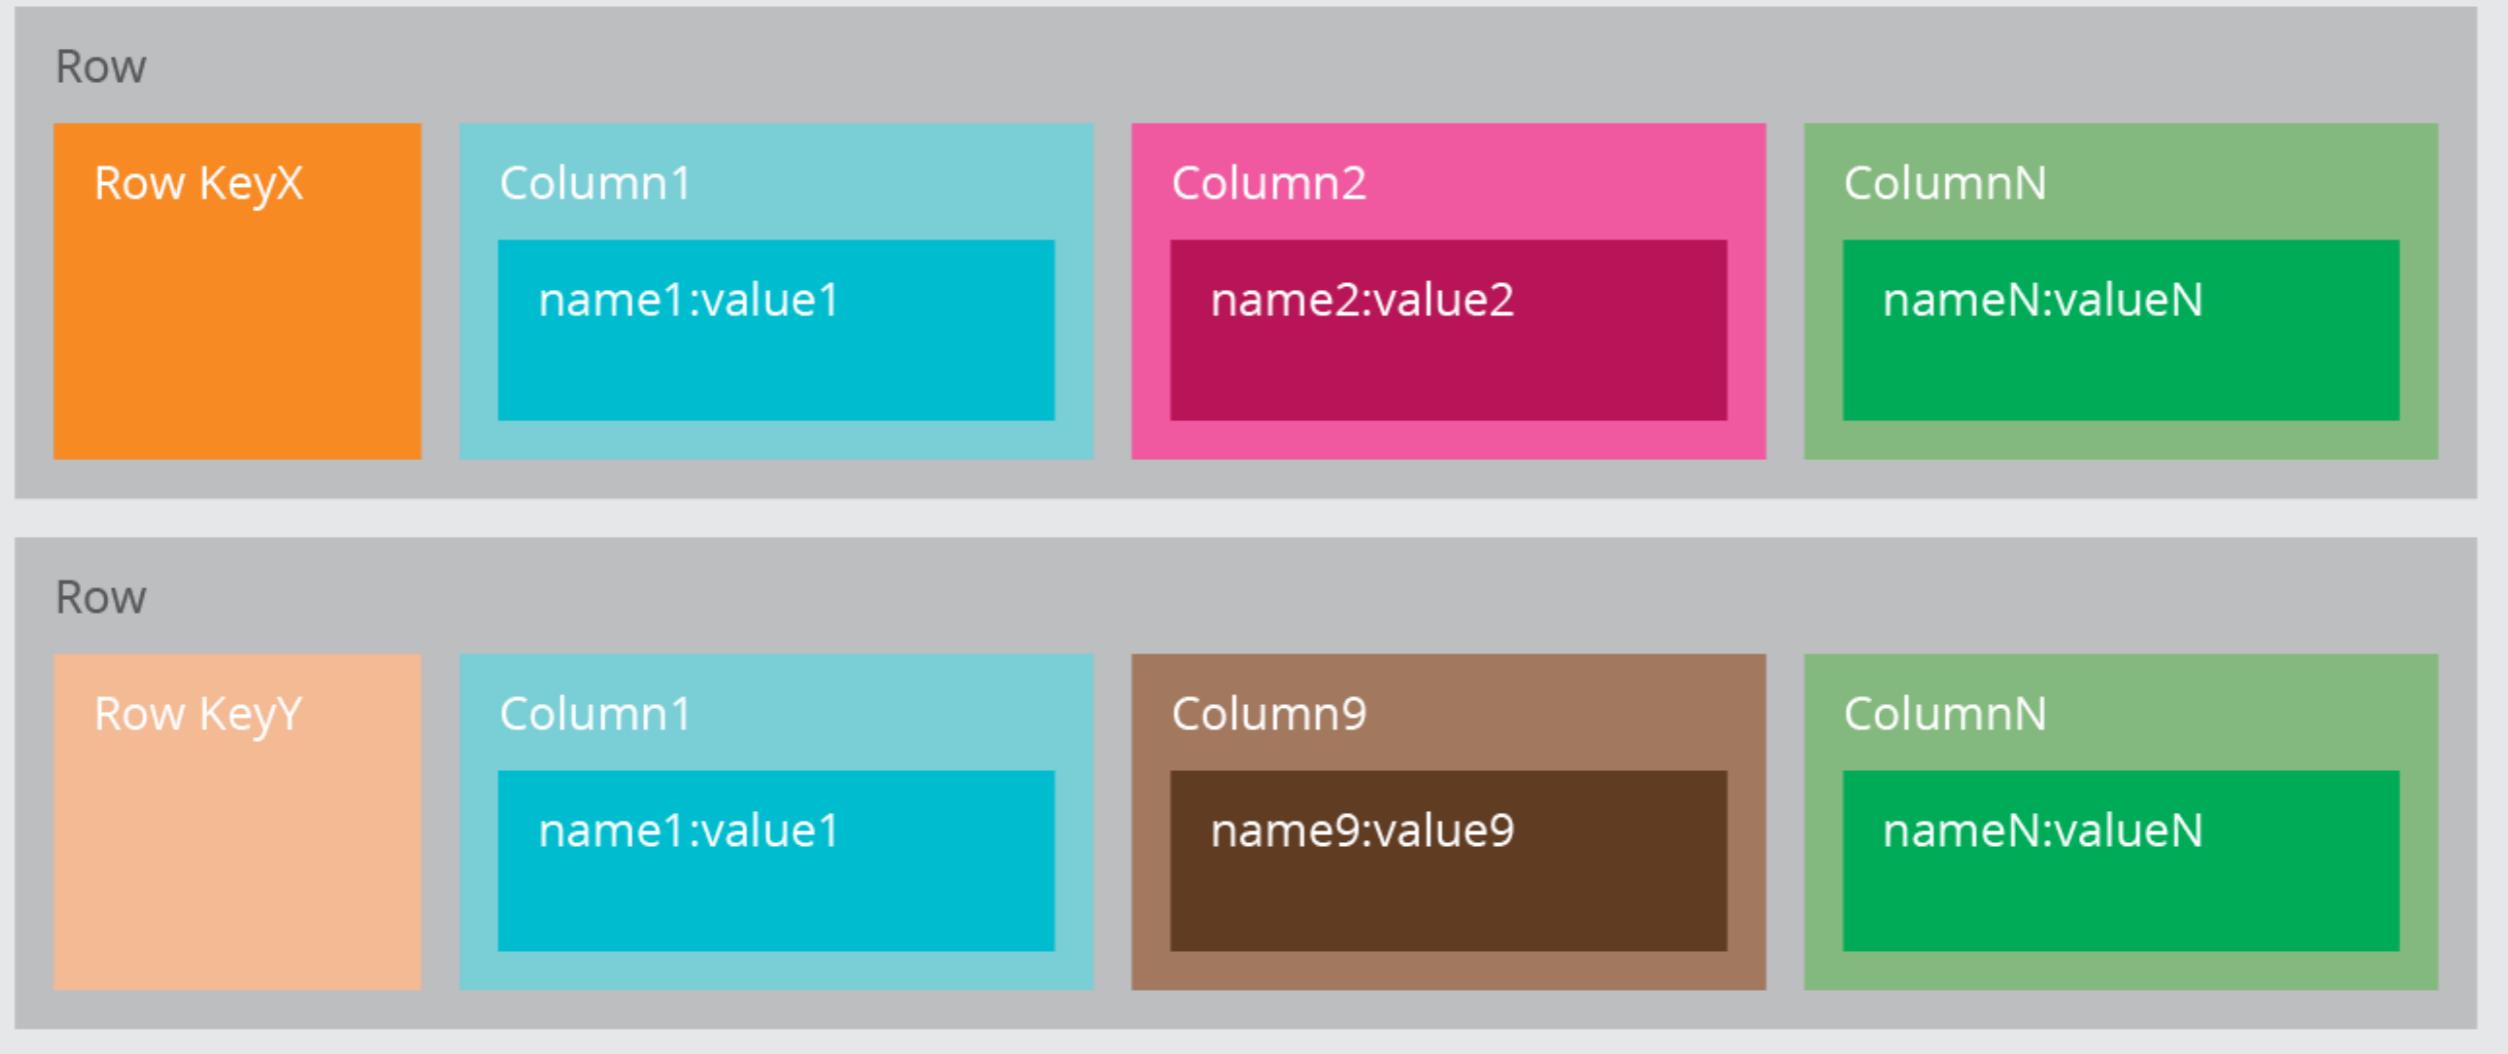


Document

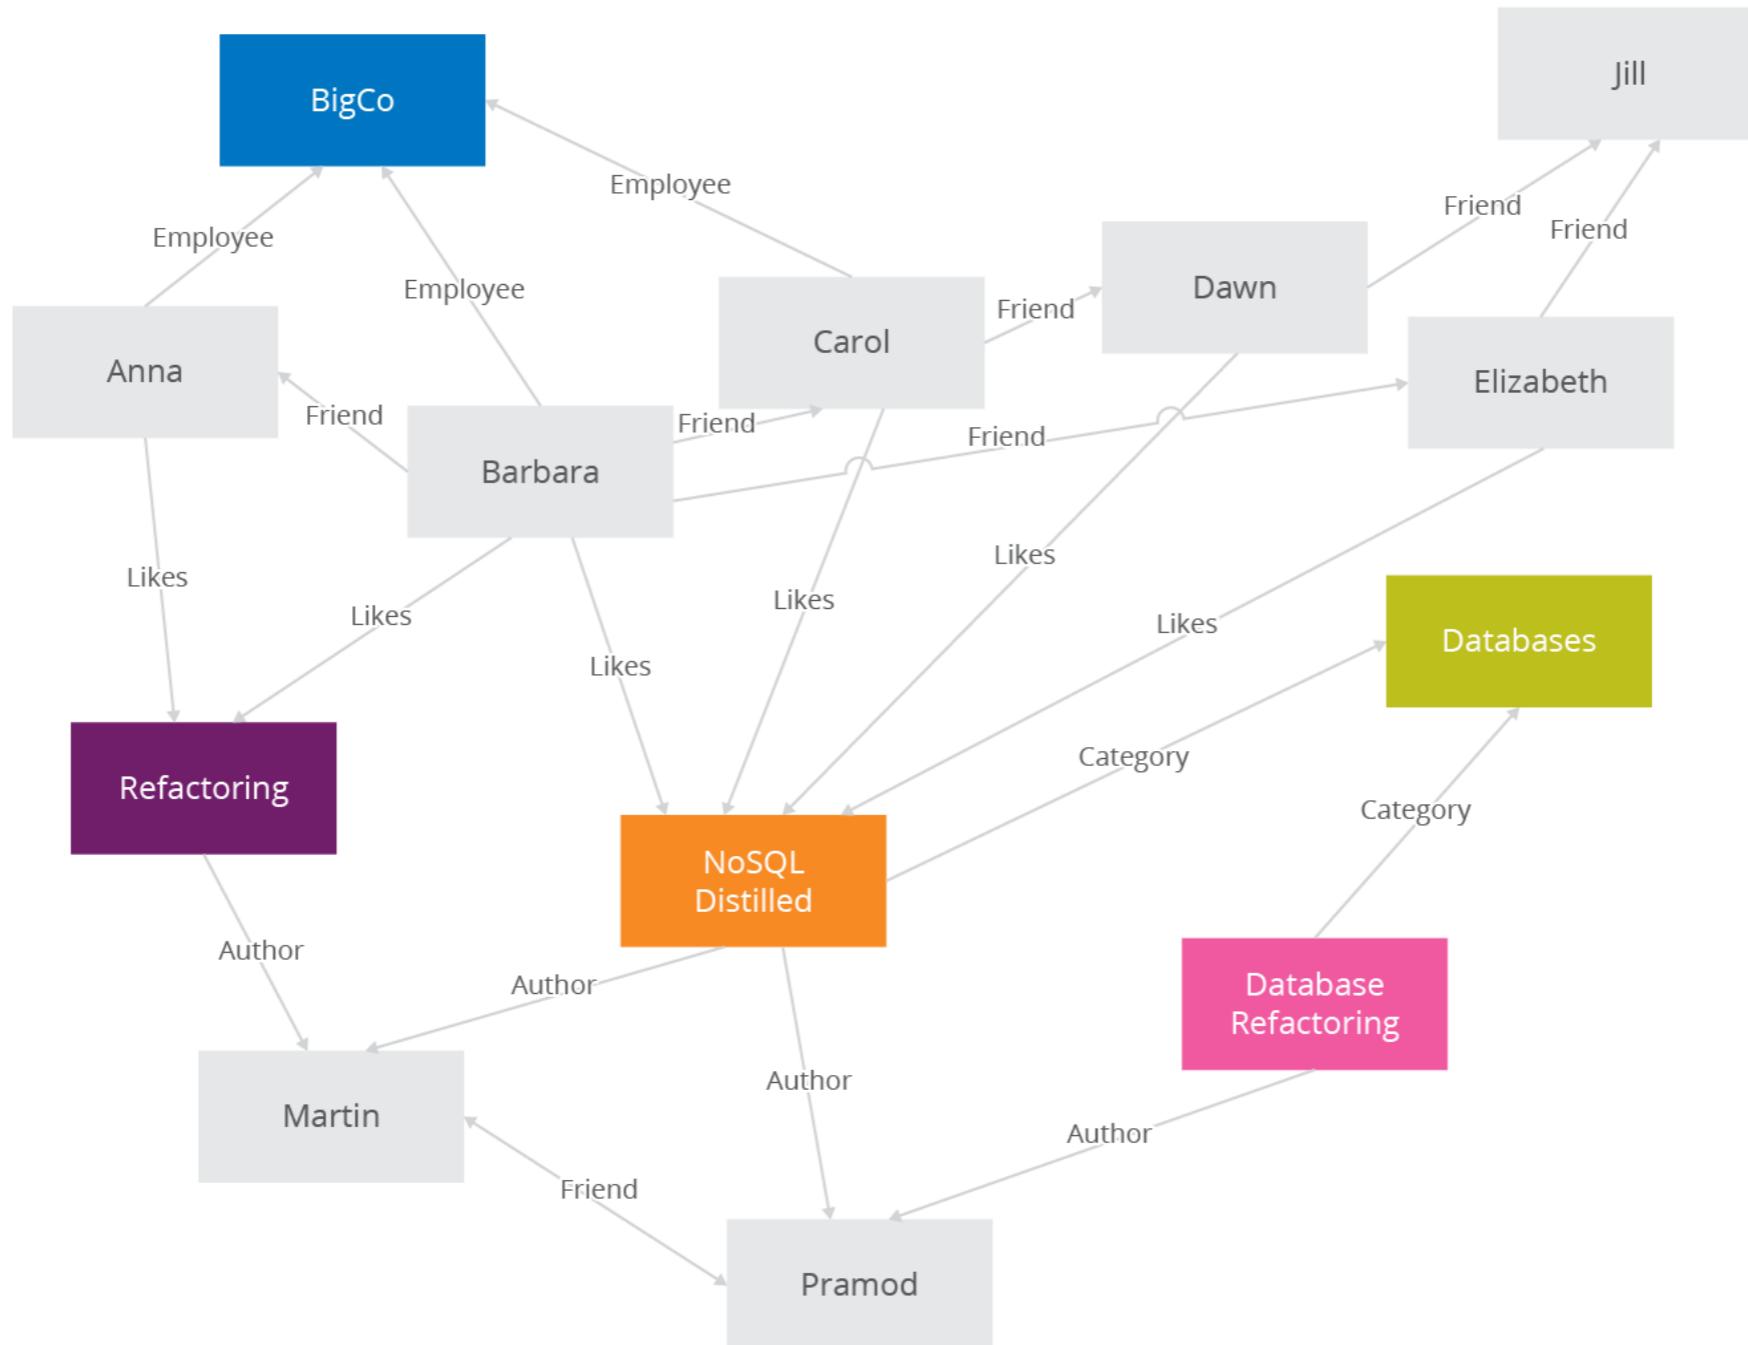


Column

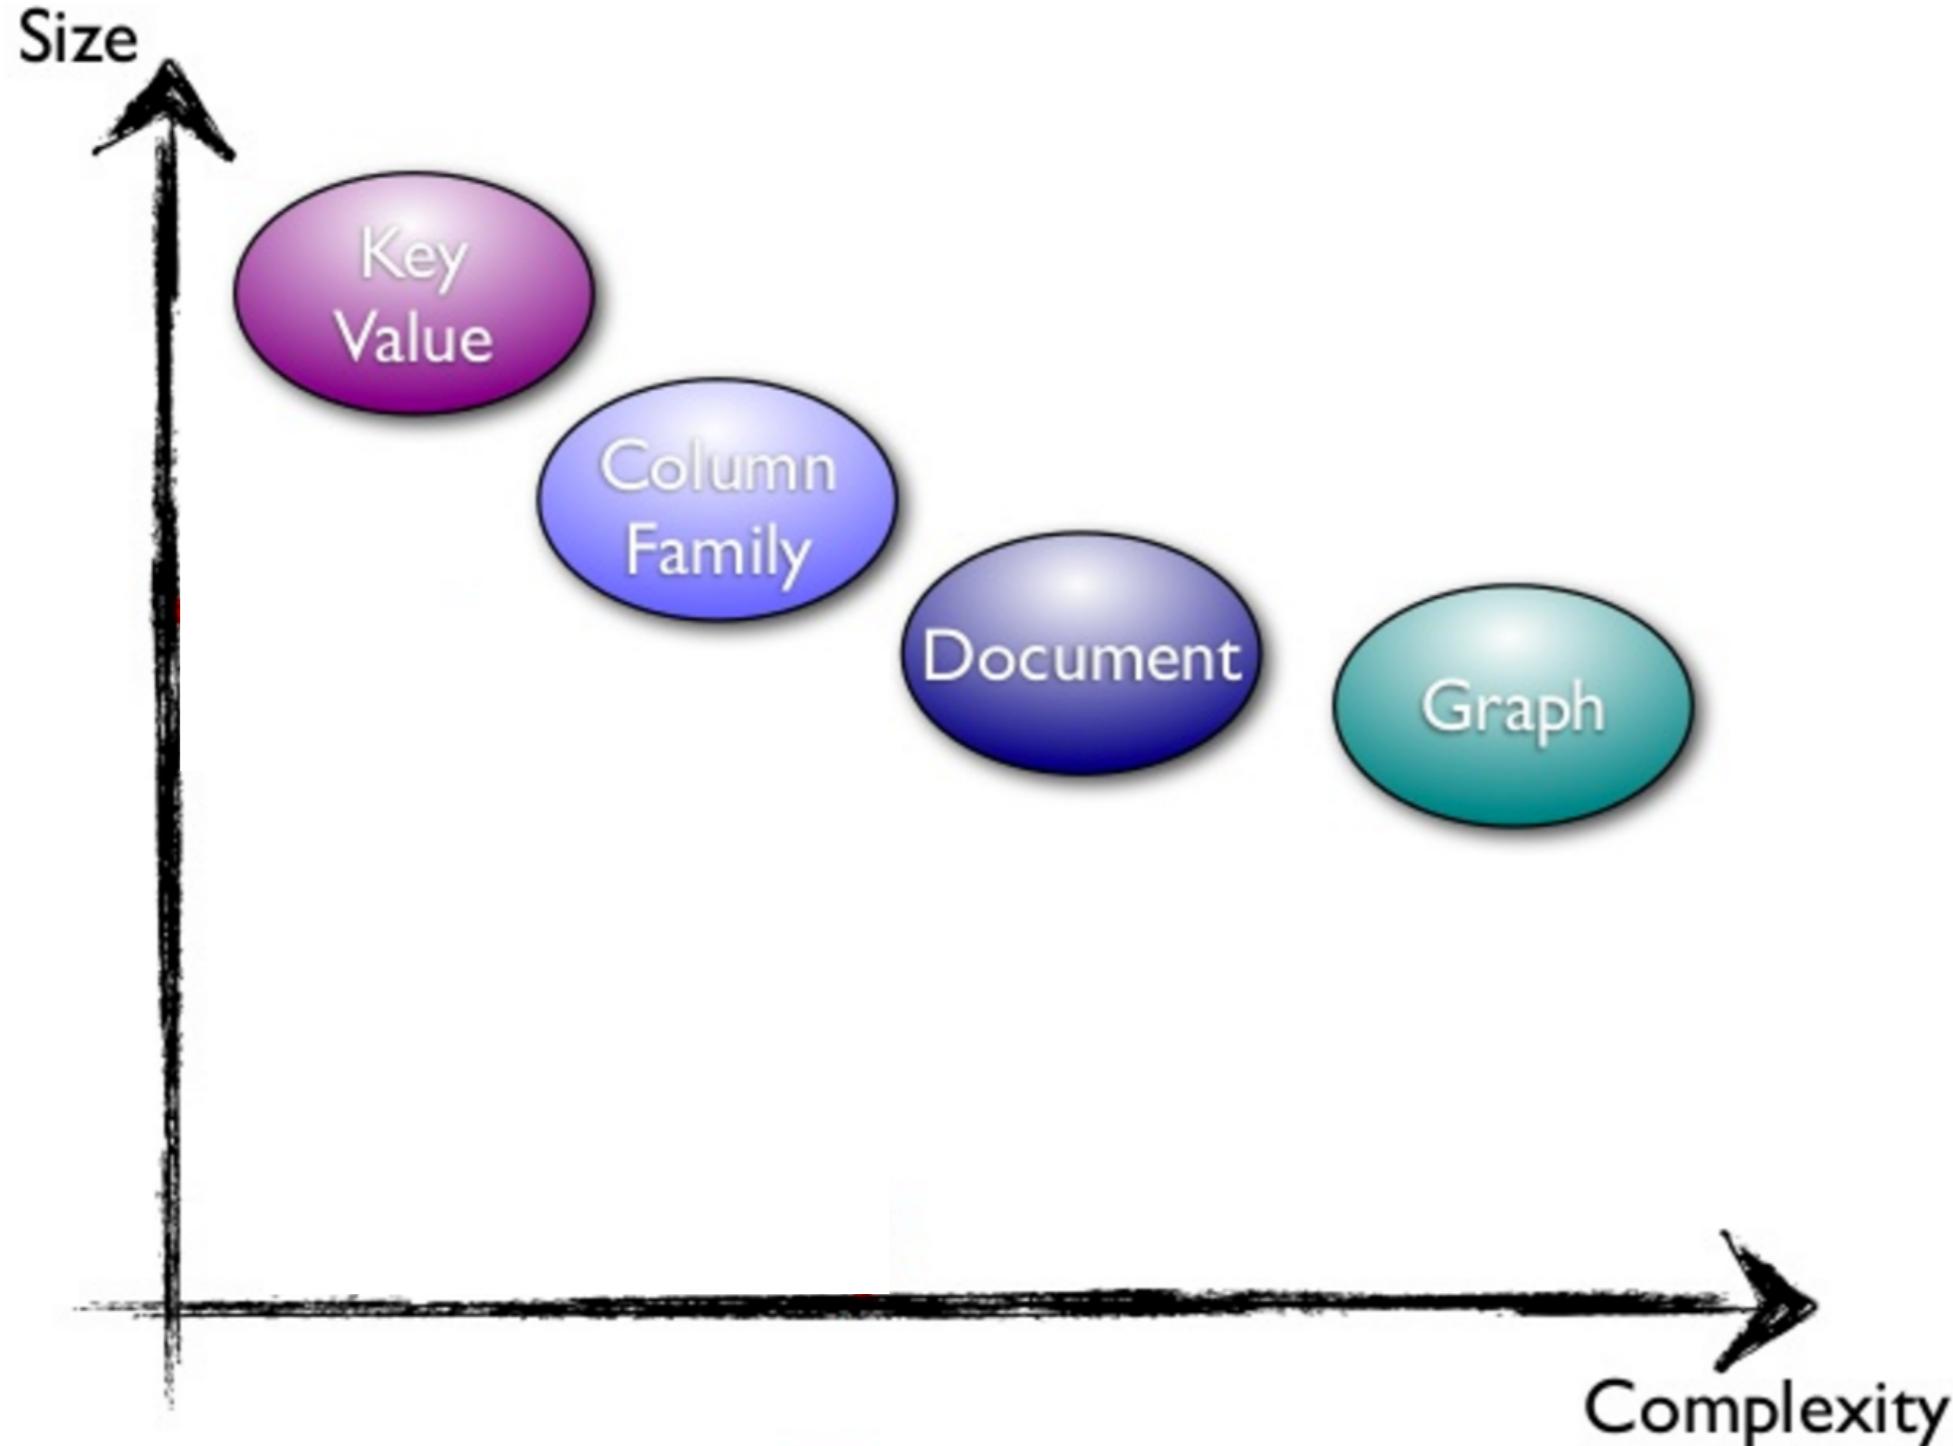
Column Family



Graph



Data Size vs Complexity



NoSQL and Consistency



RDBMS => ACID
NoSQL => BASE



BASE

Basic Availability

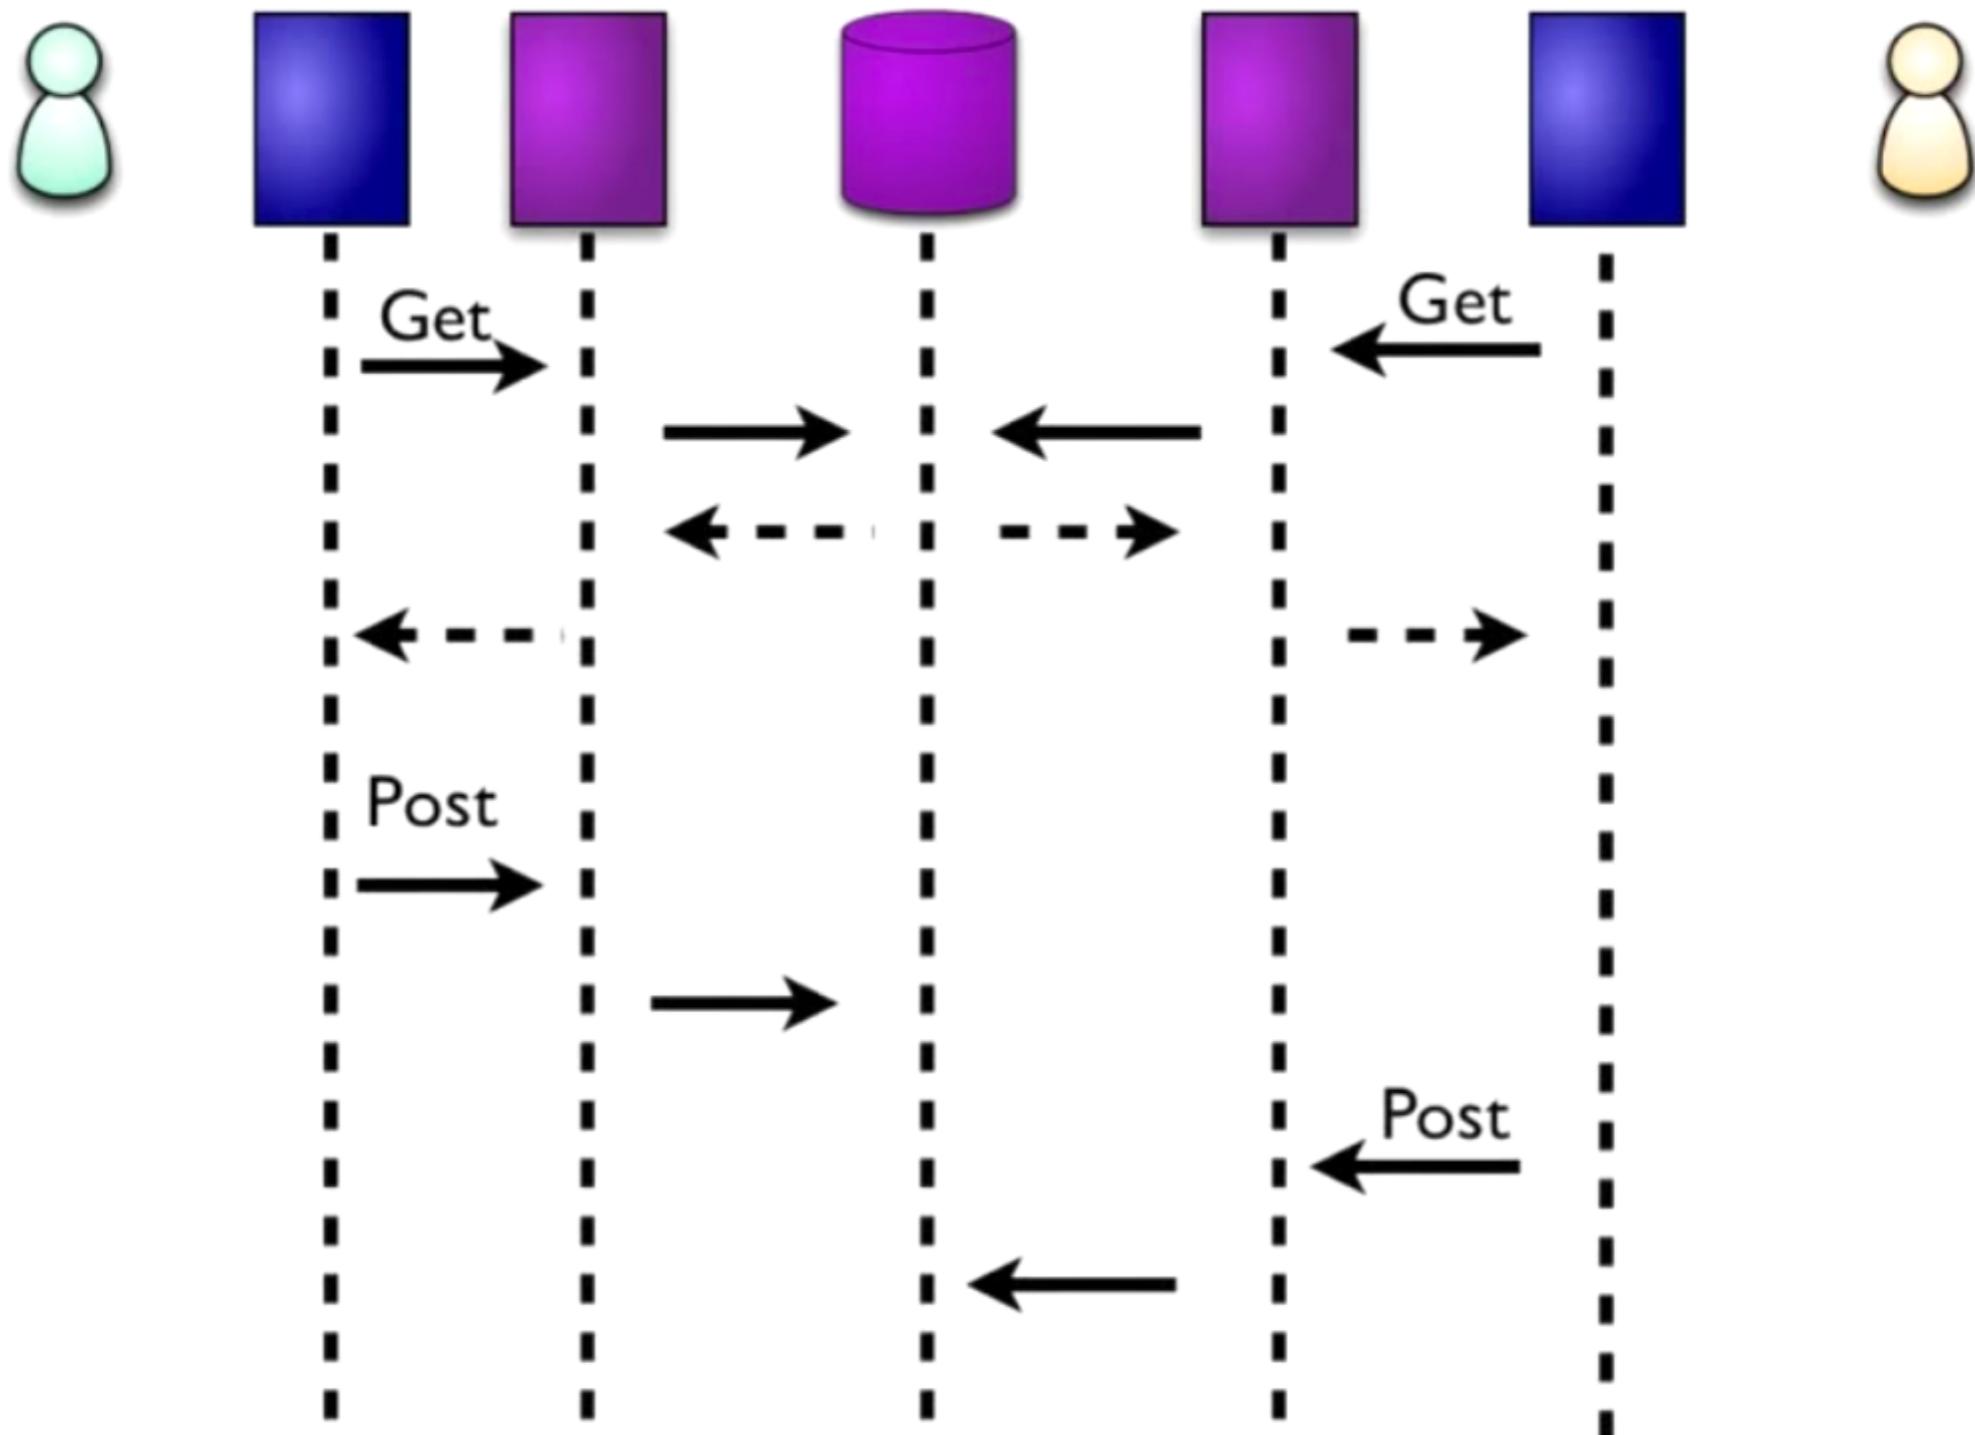
Soft state

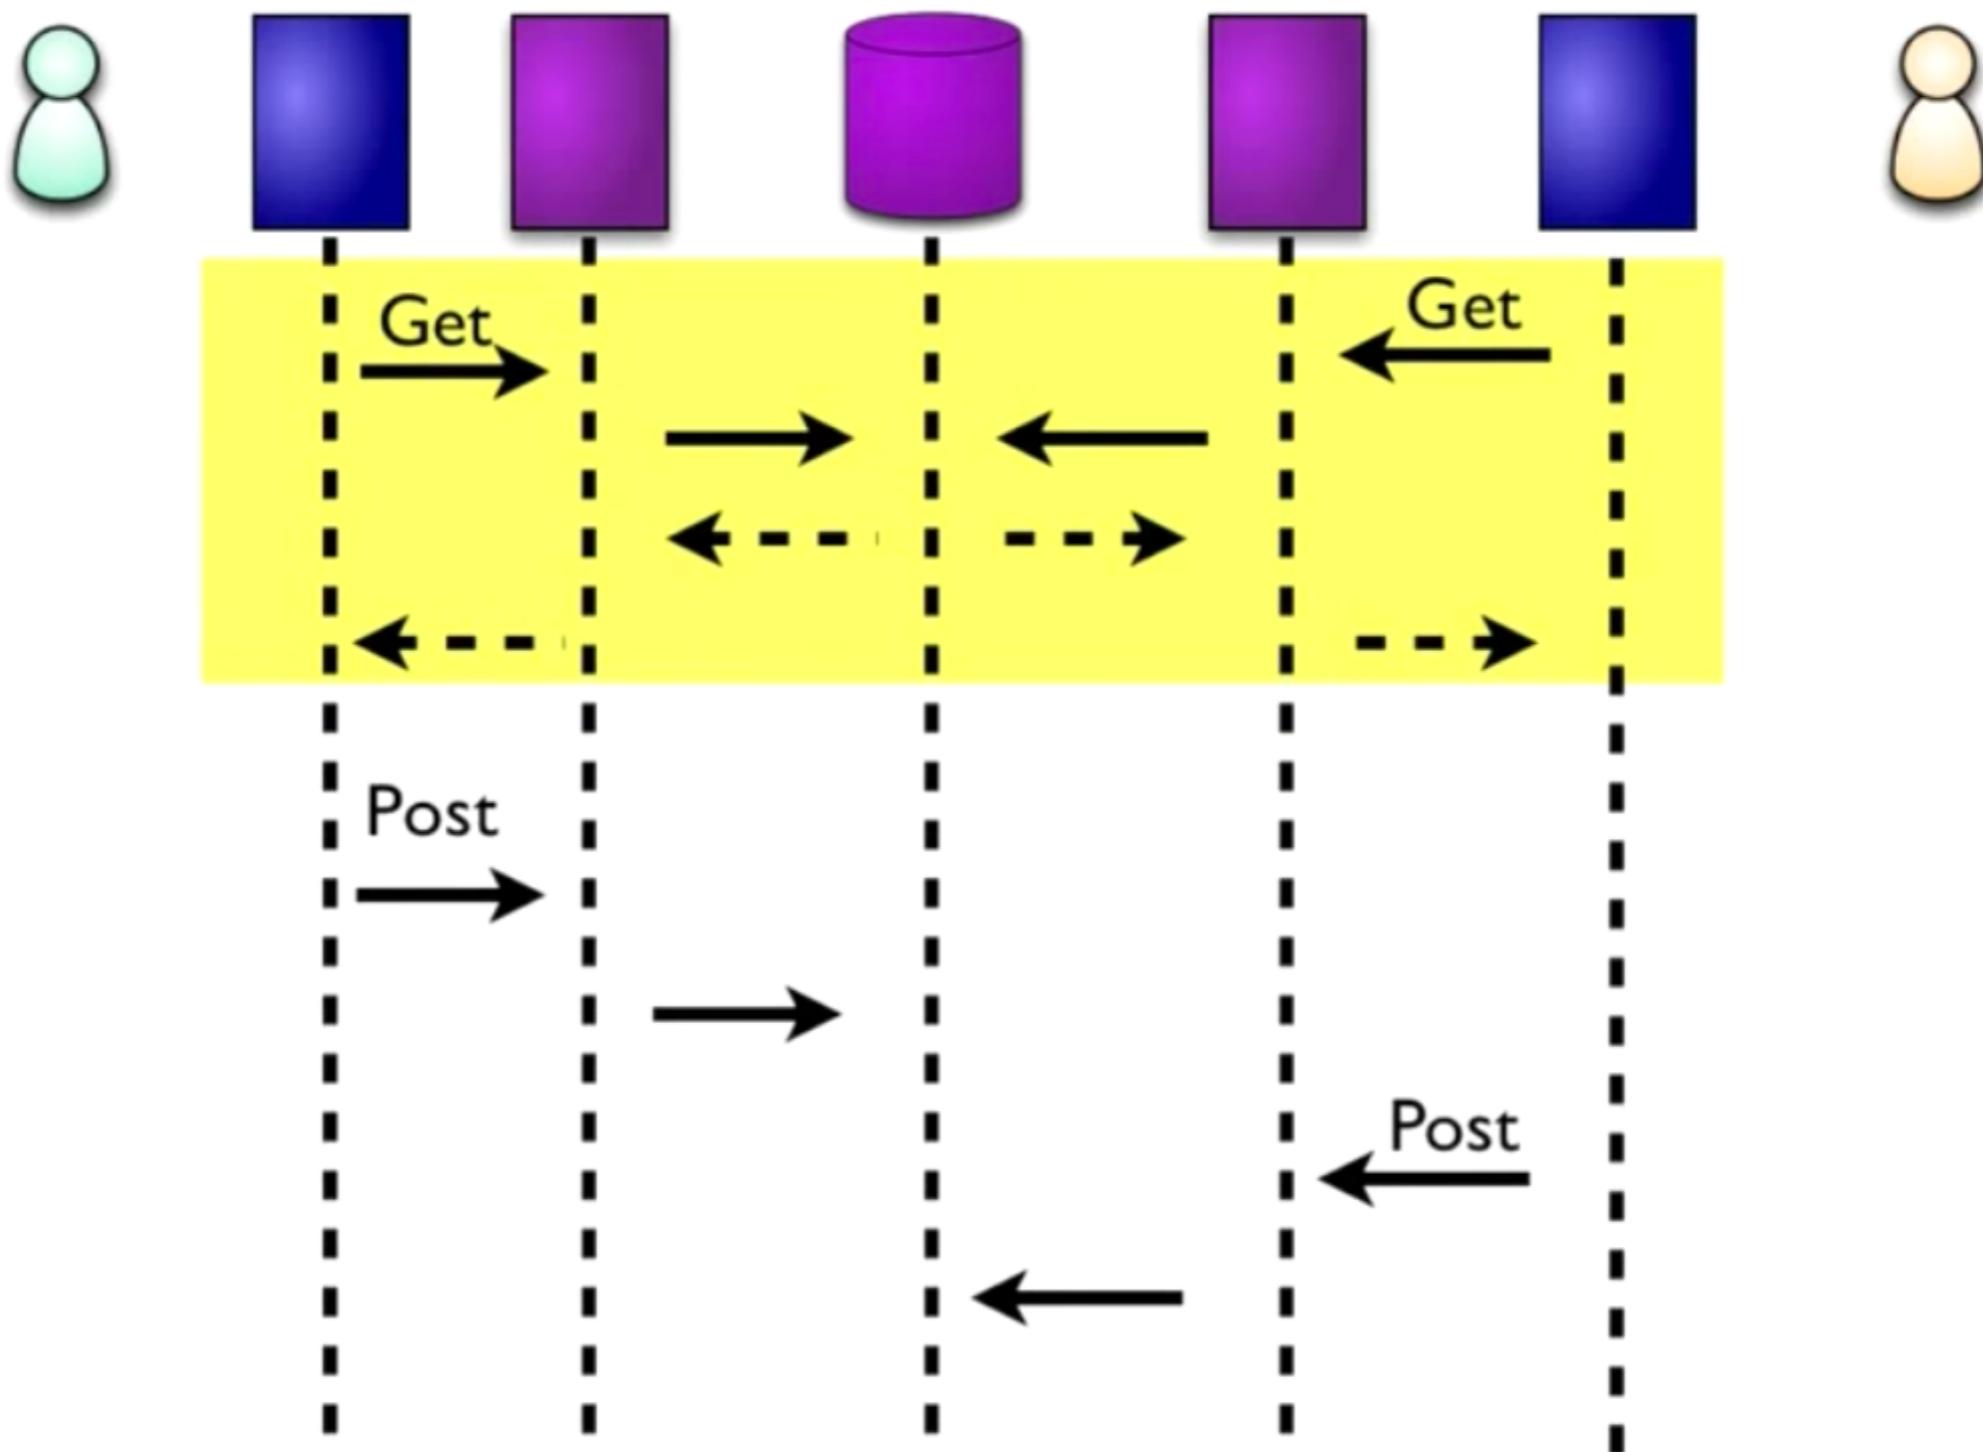
Eventual consistency

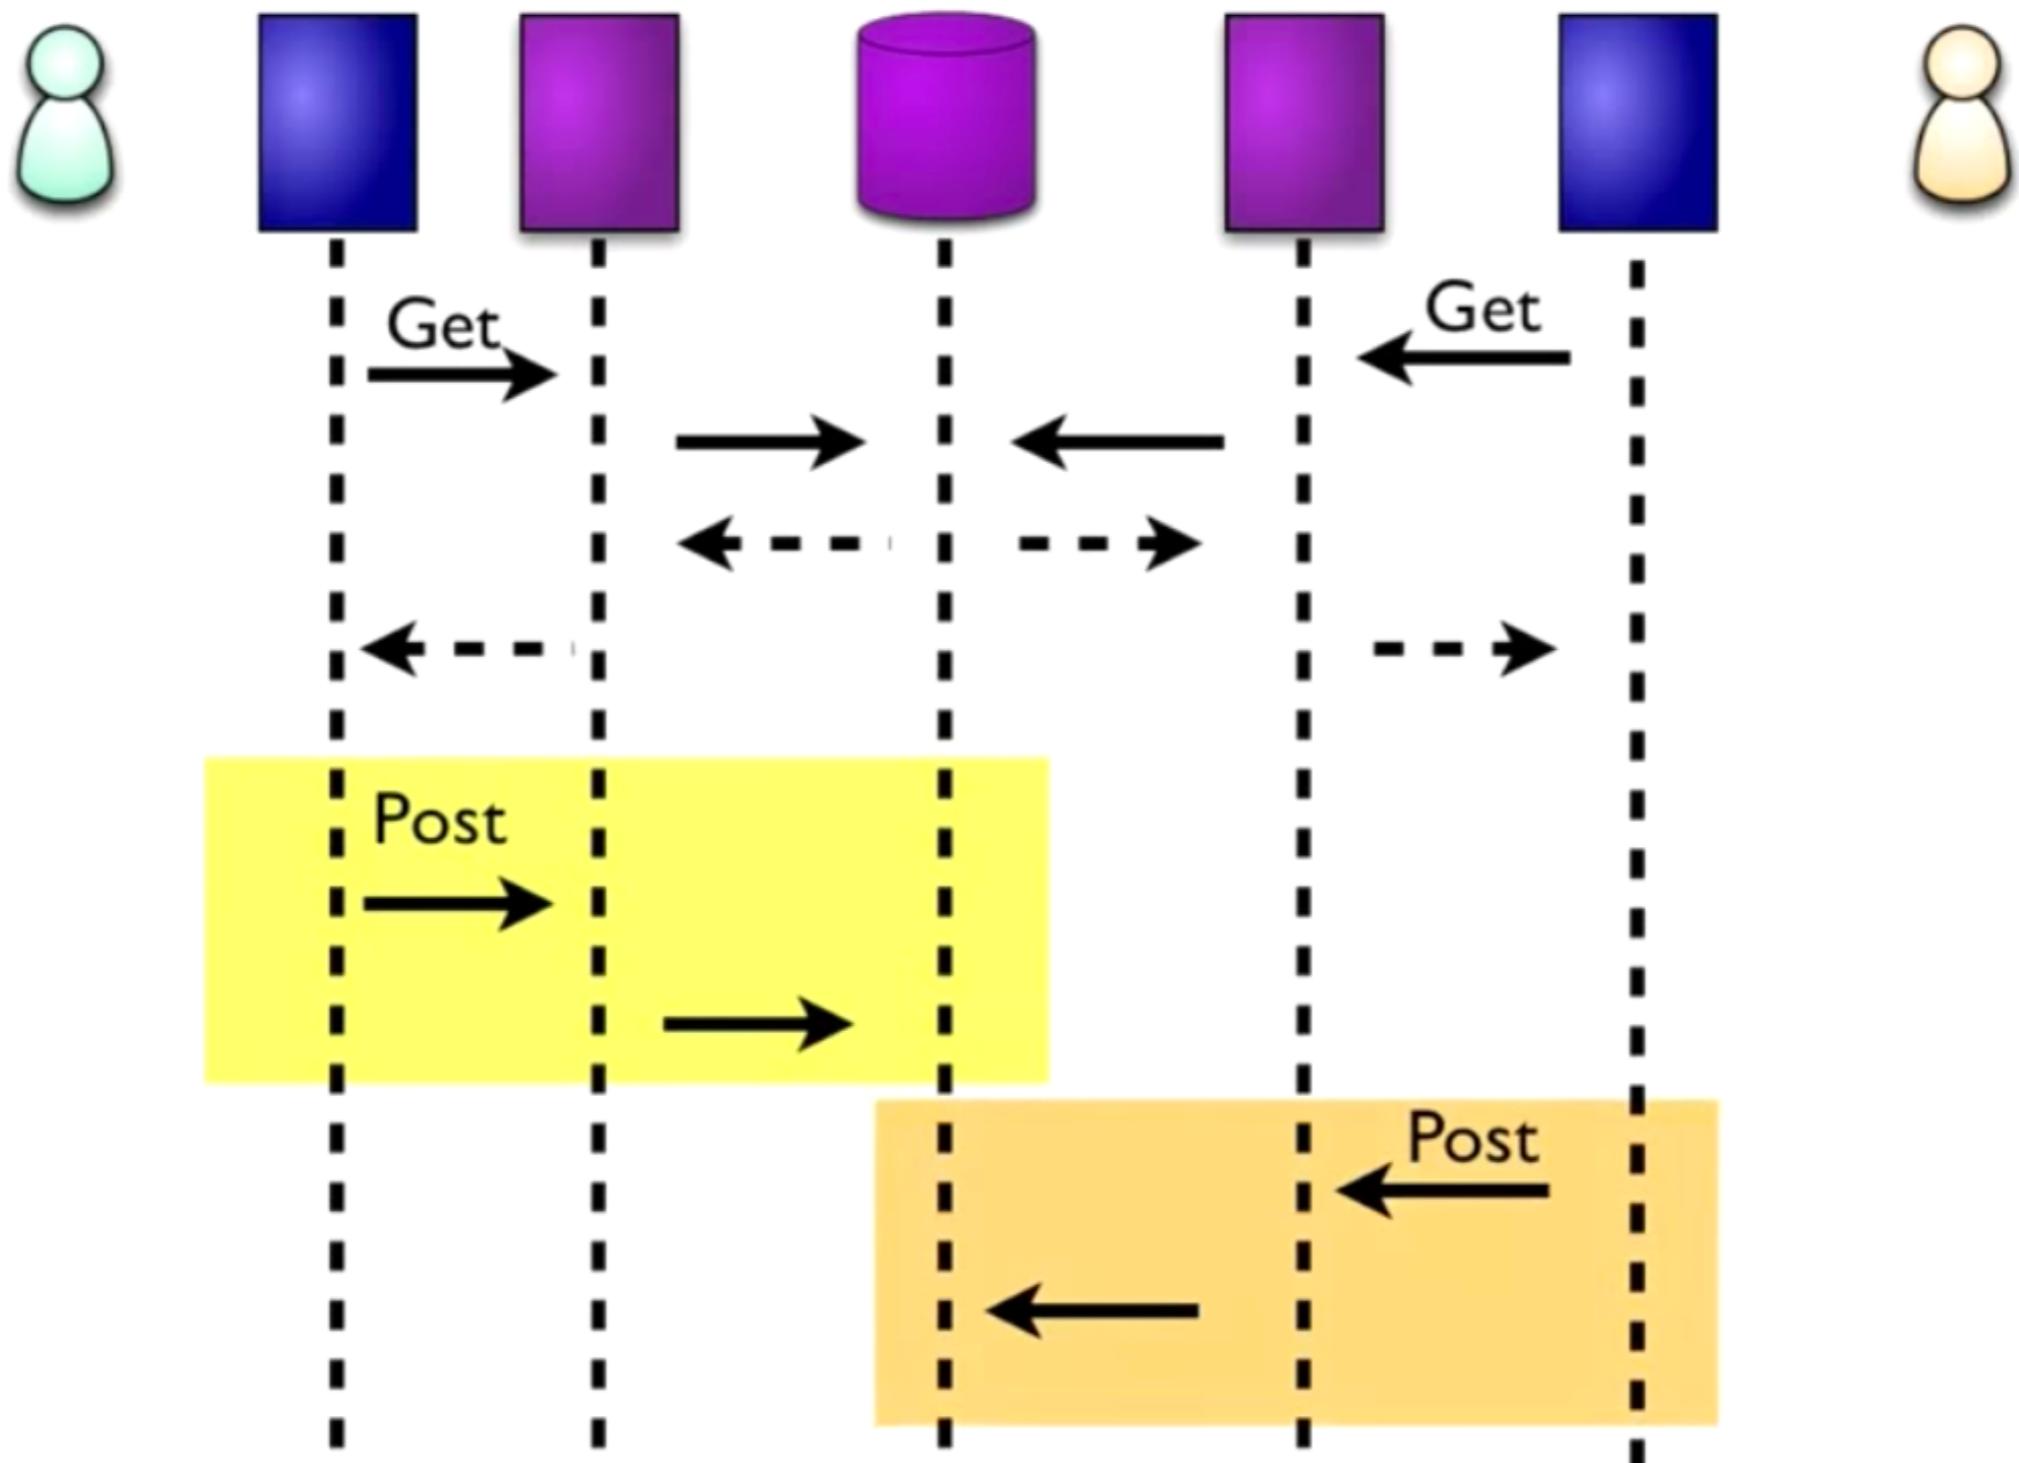


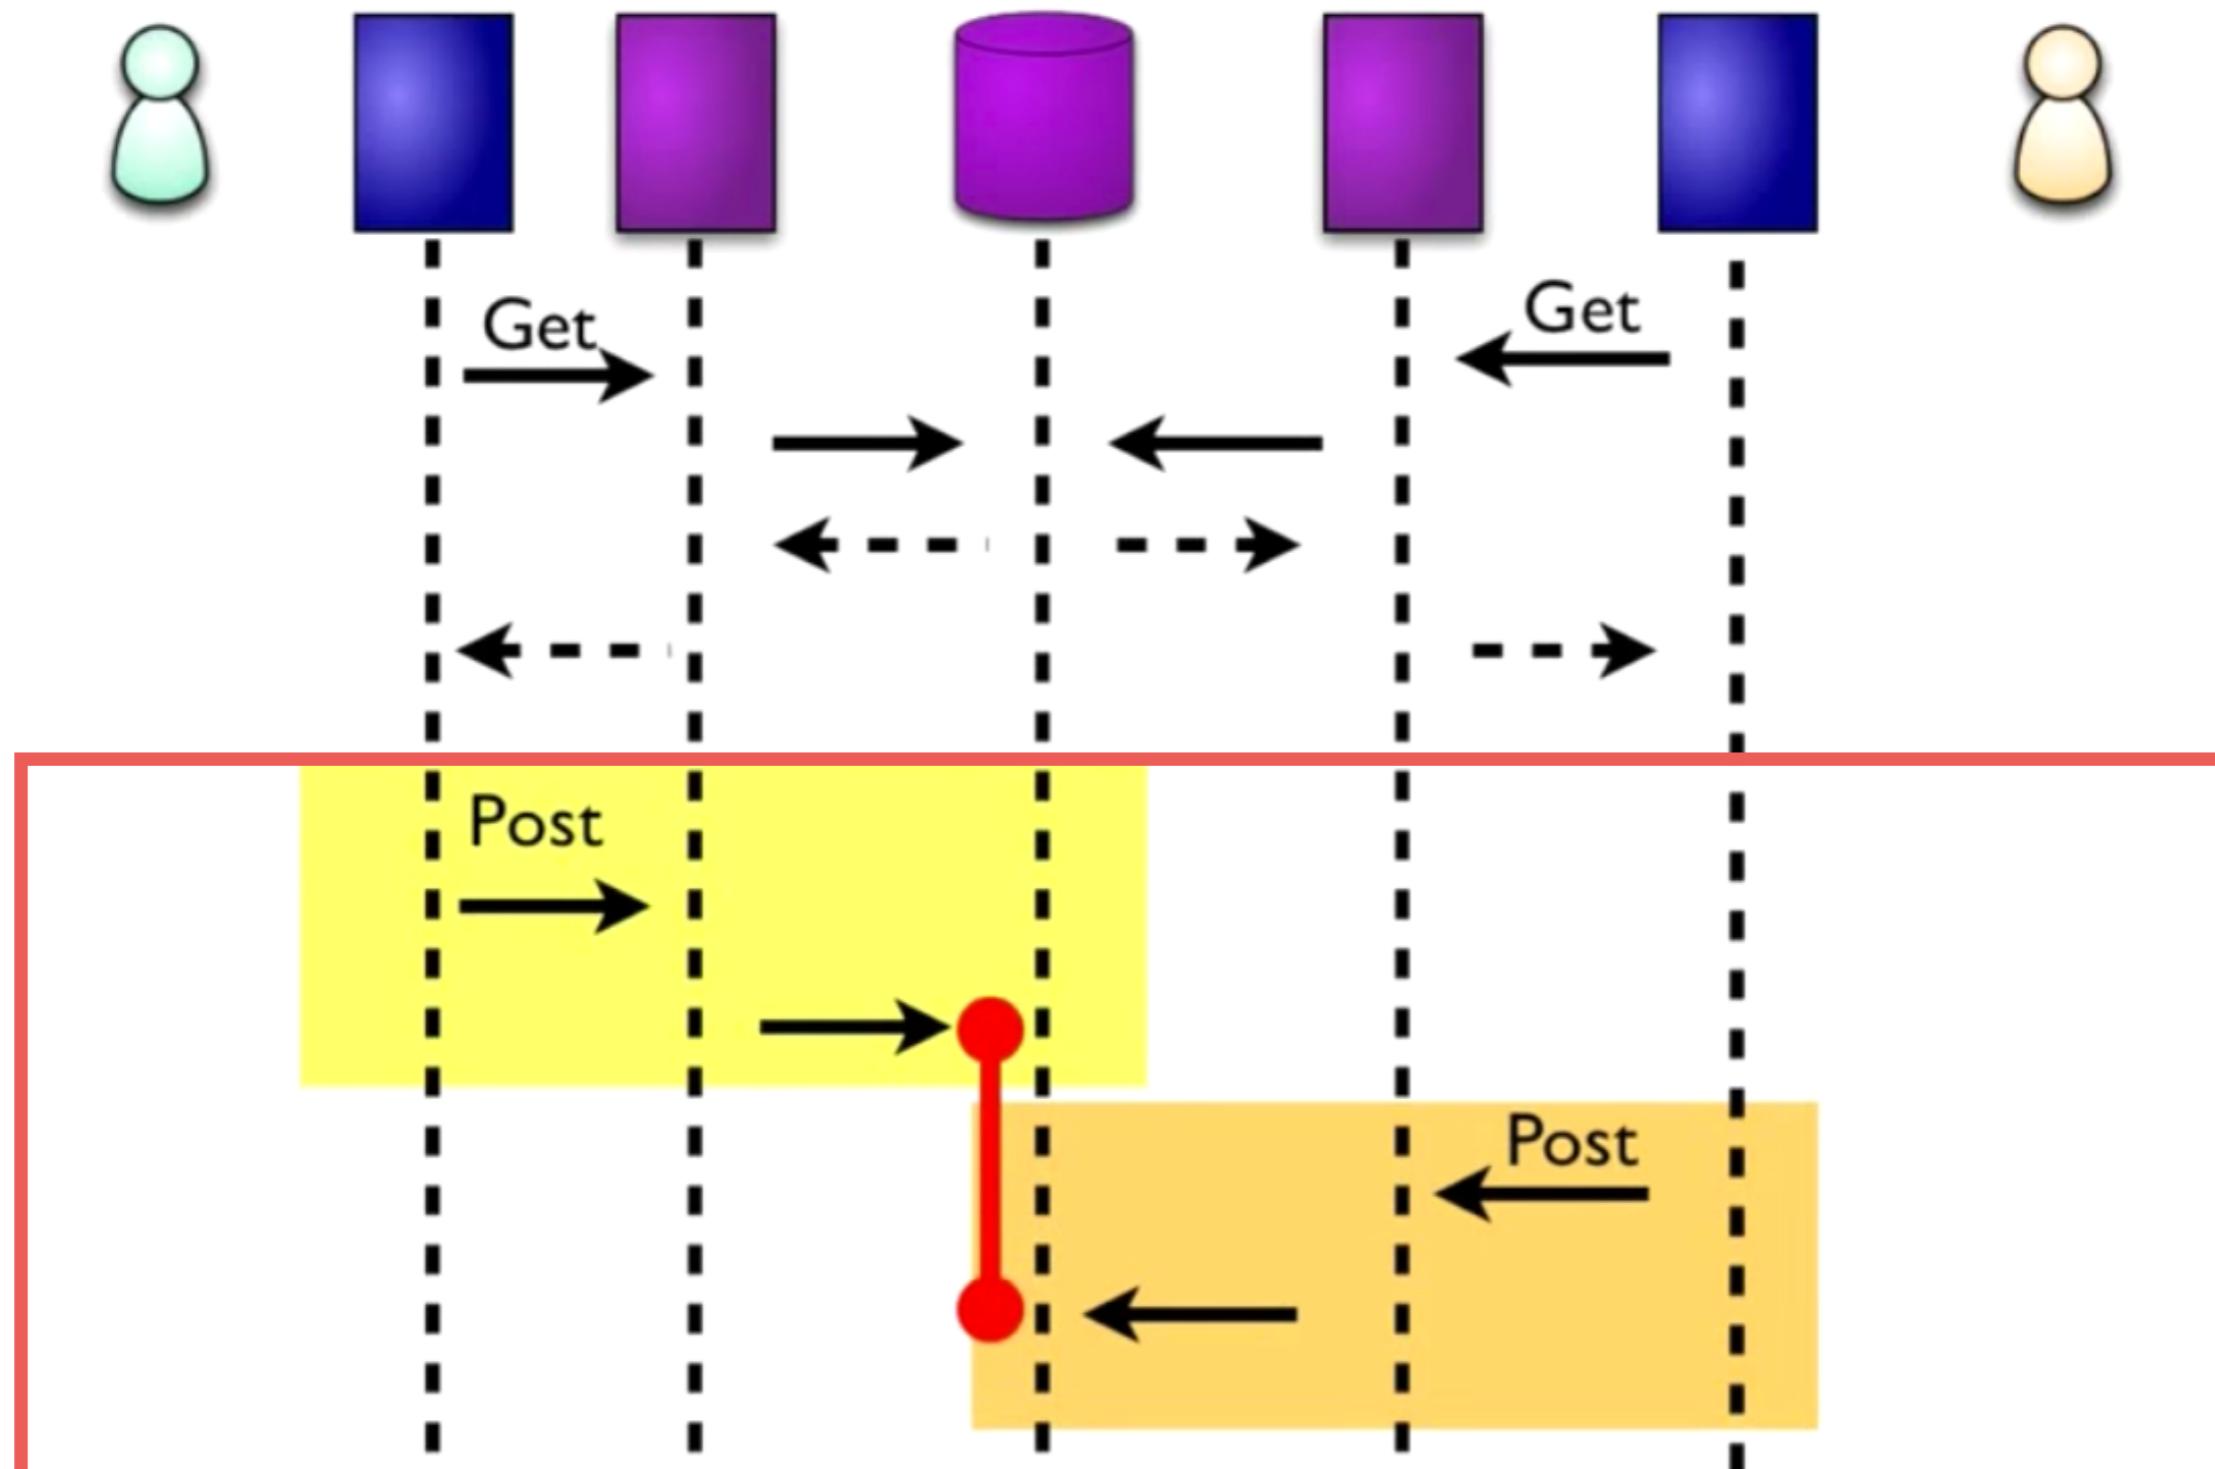
Consistency problem











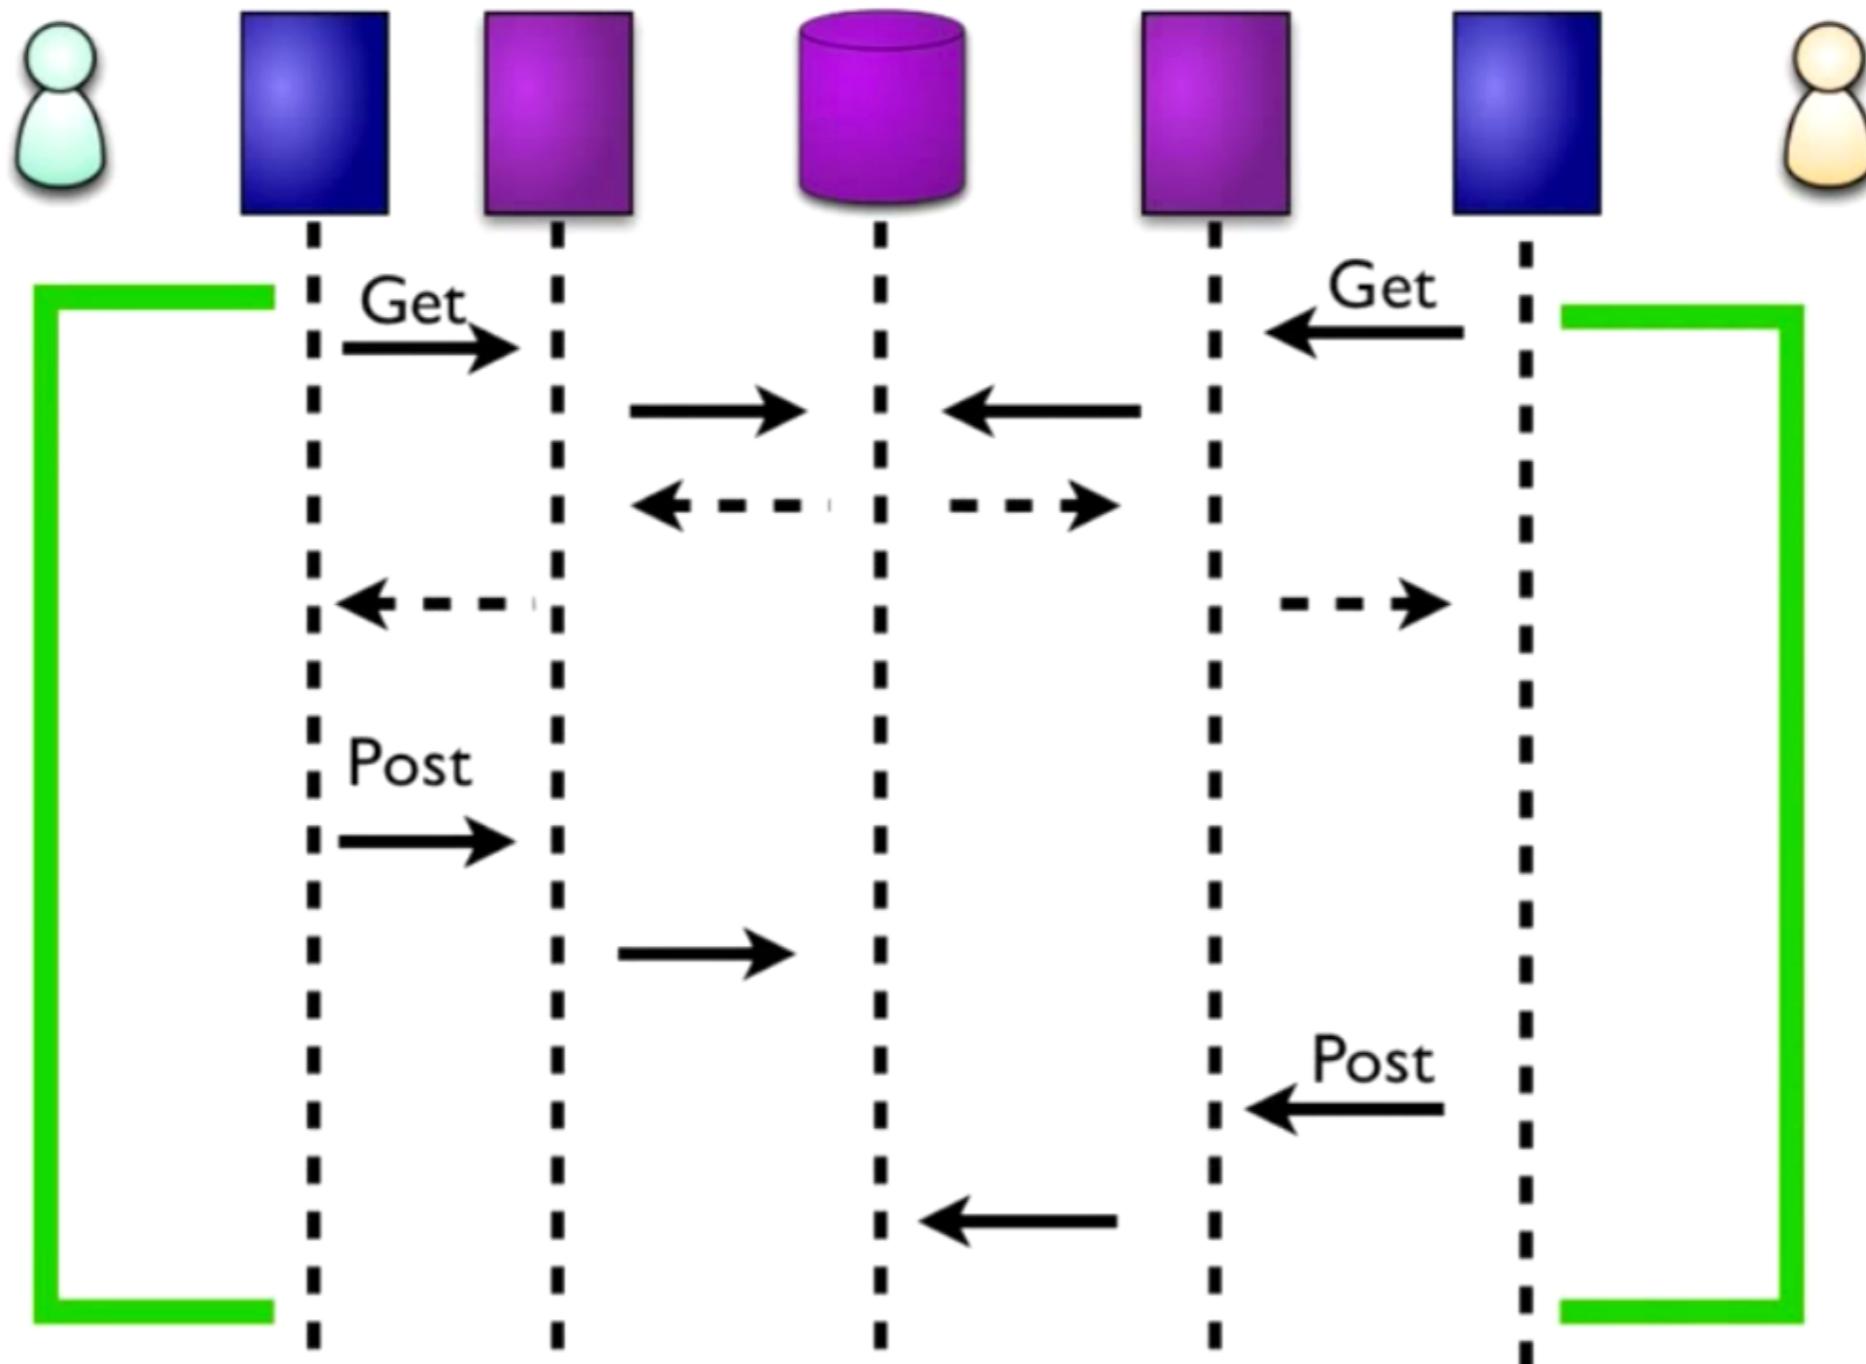
Problem ?



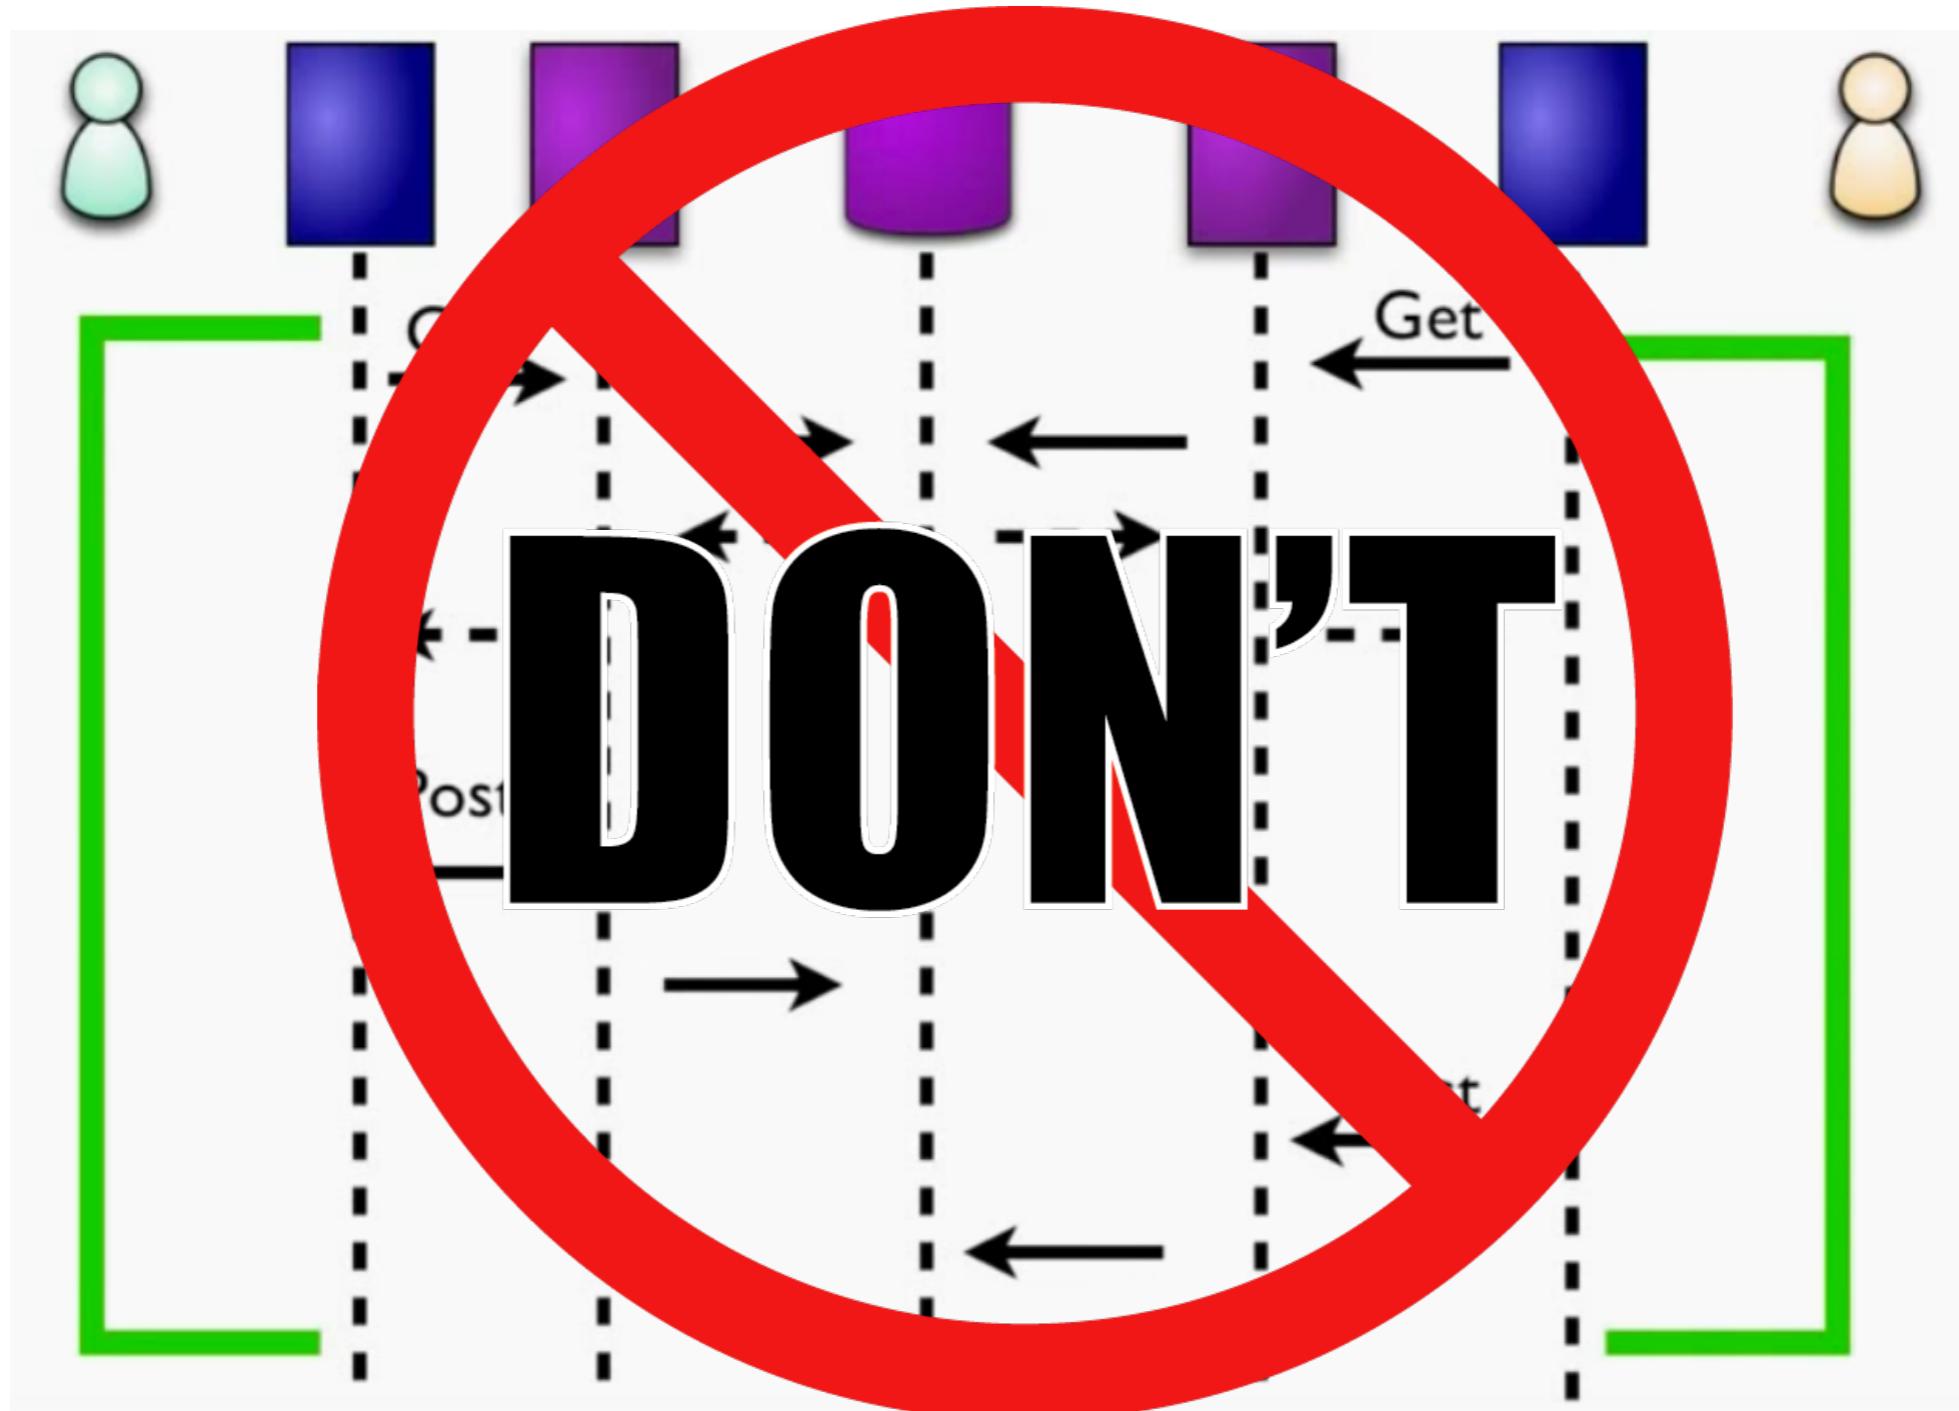
Solution ?



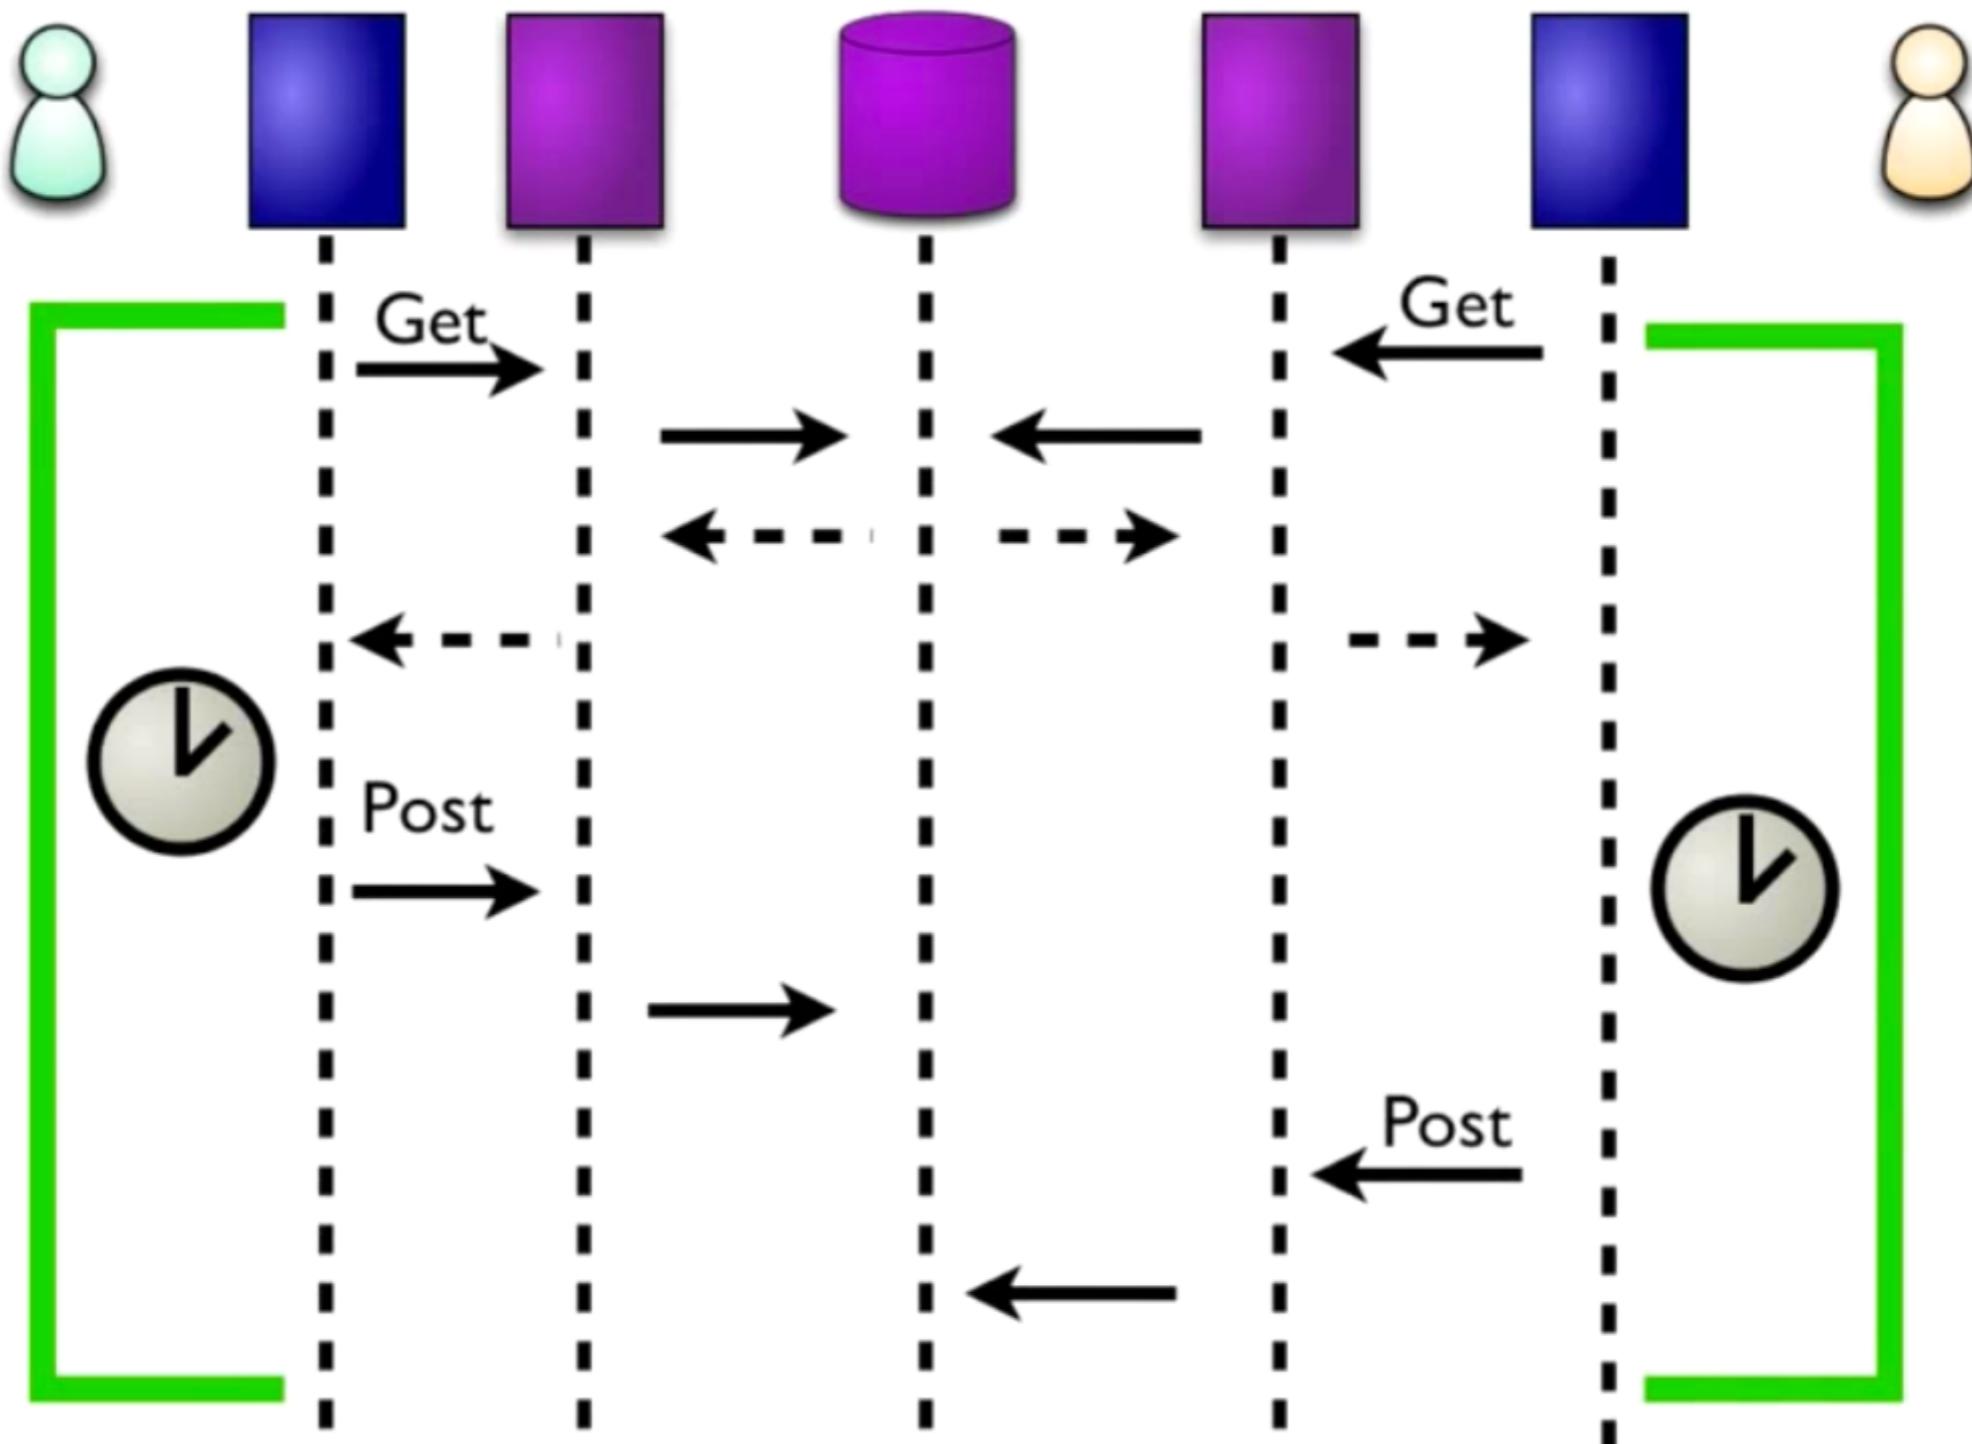
Transaction ?



Transaction ?



Deadlock

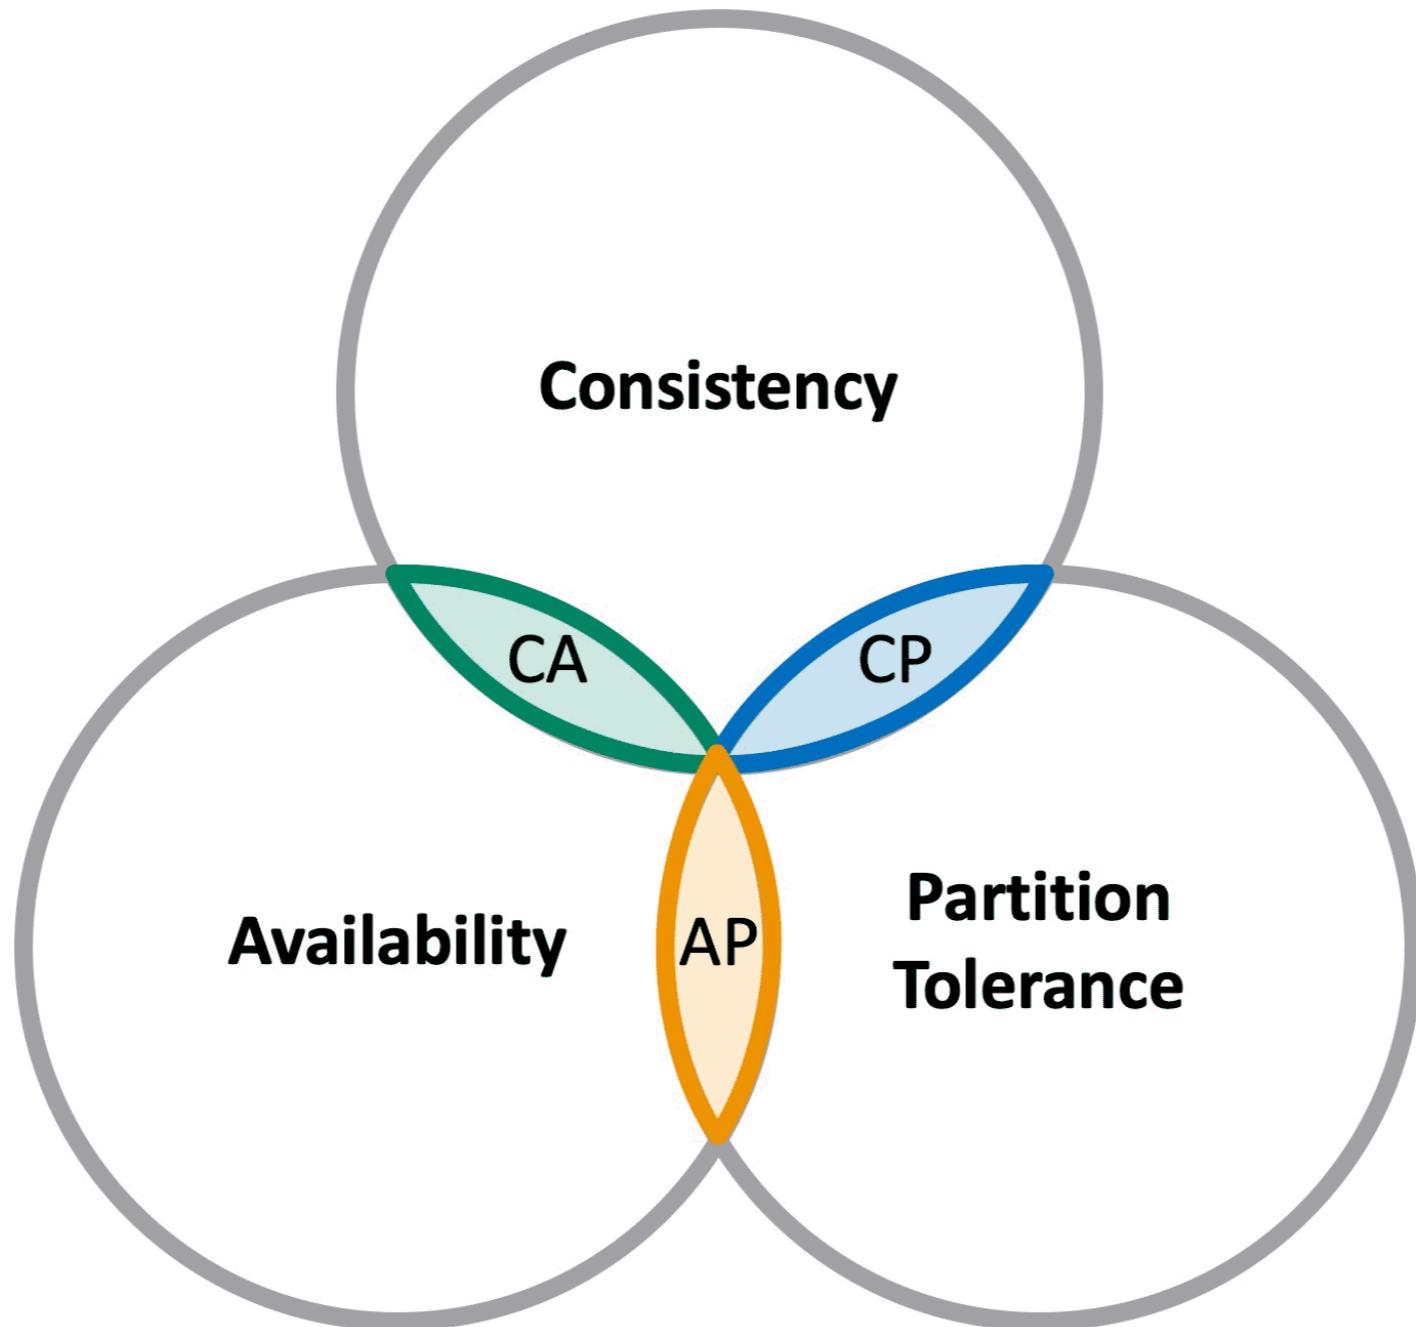


New way ?



CAP Theorem





Consistency

CA Category

CP Category

There is a risk of some data becoming unavailable
Ex: MongoDB Hbase
Memcache Big table Redis

C

Network Problem might stop the system
Ex: RDBMS (Oracle SQL Server MySQL)

Pick two

Partition Tolerance

P

A

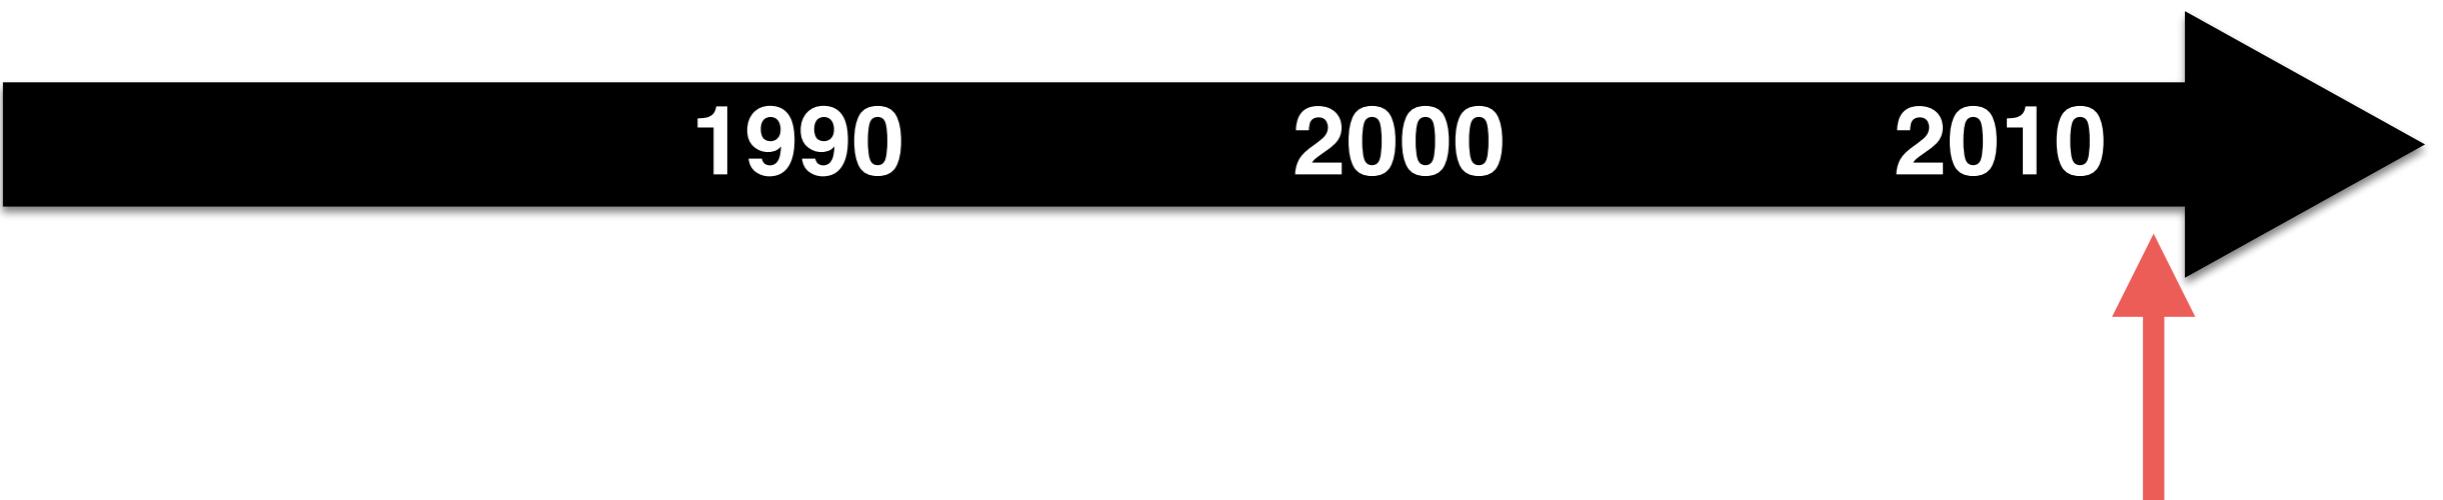
Availability

AP Category

Clients may read inconsistent data

Ex: Cassandra RIAK CouchDB





Polyglot persistence



The
Pragmatic
Programmers

Second Edition
**Seven Databases
in Seven Weeks**

A Guide to Modern
Databases and the
NoSQL Movement

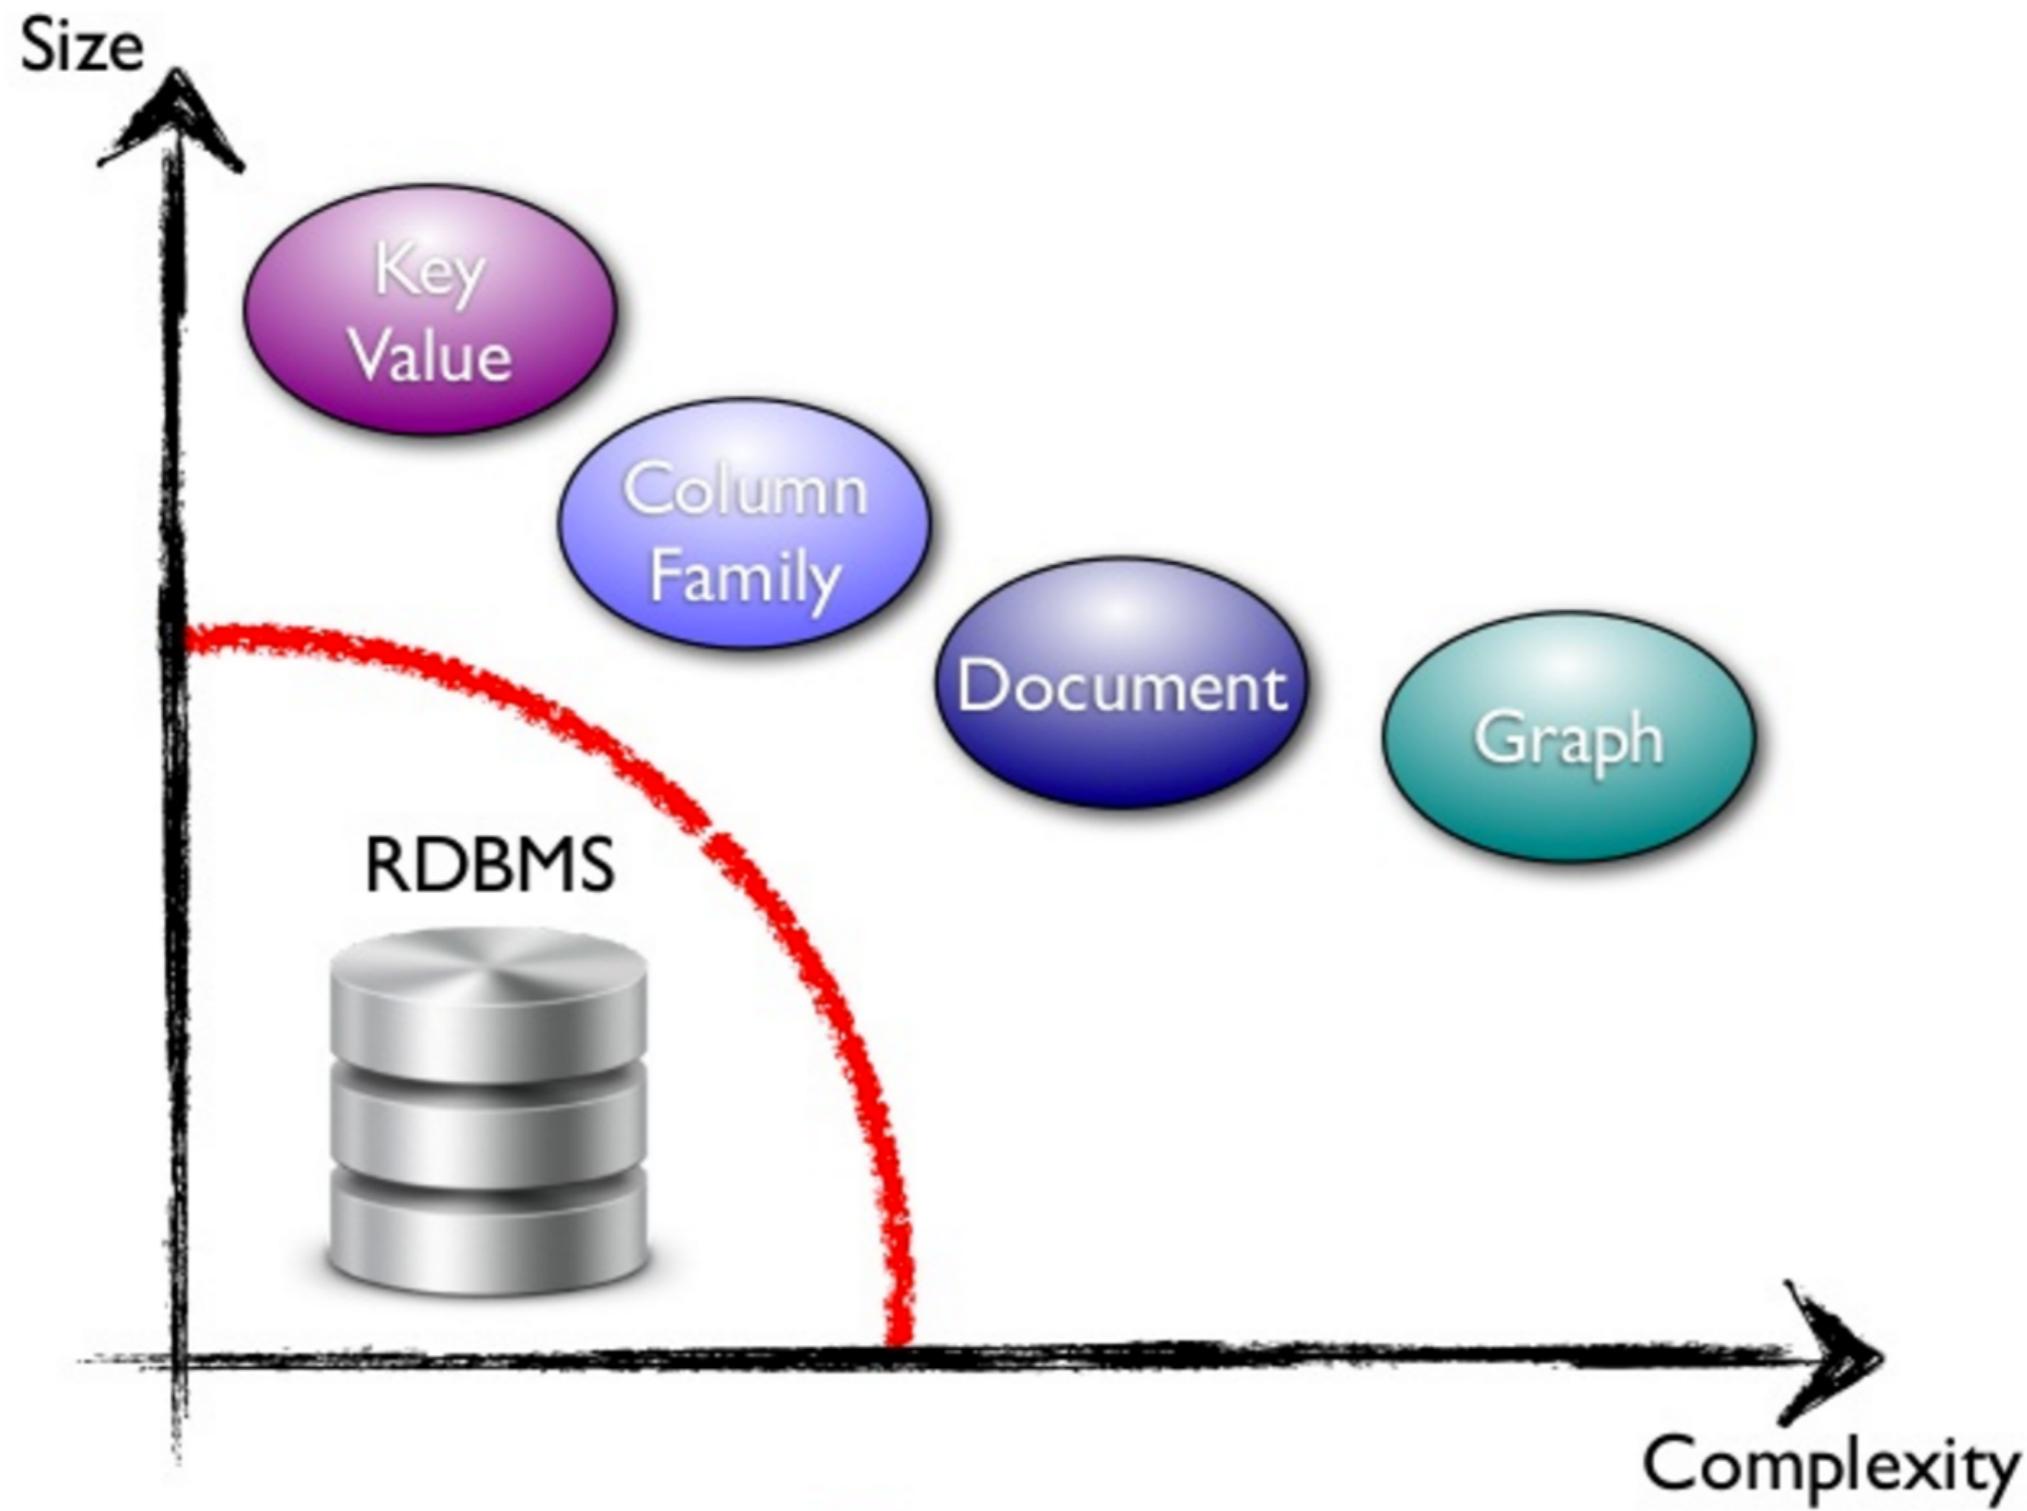
Luc Perkins
with Eric Redmond and Jim R. Wilson

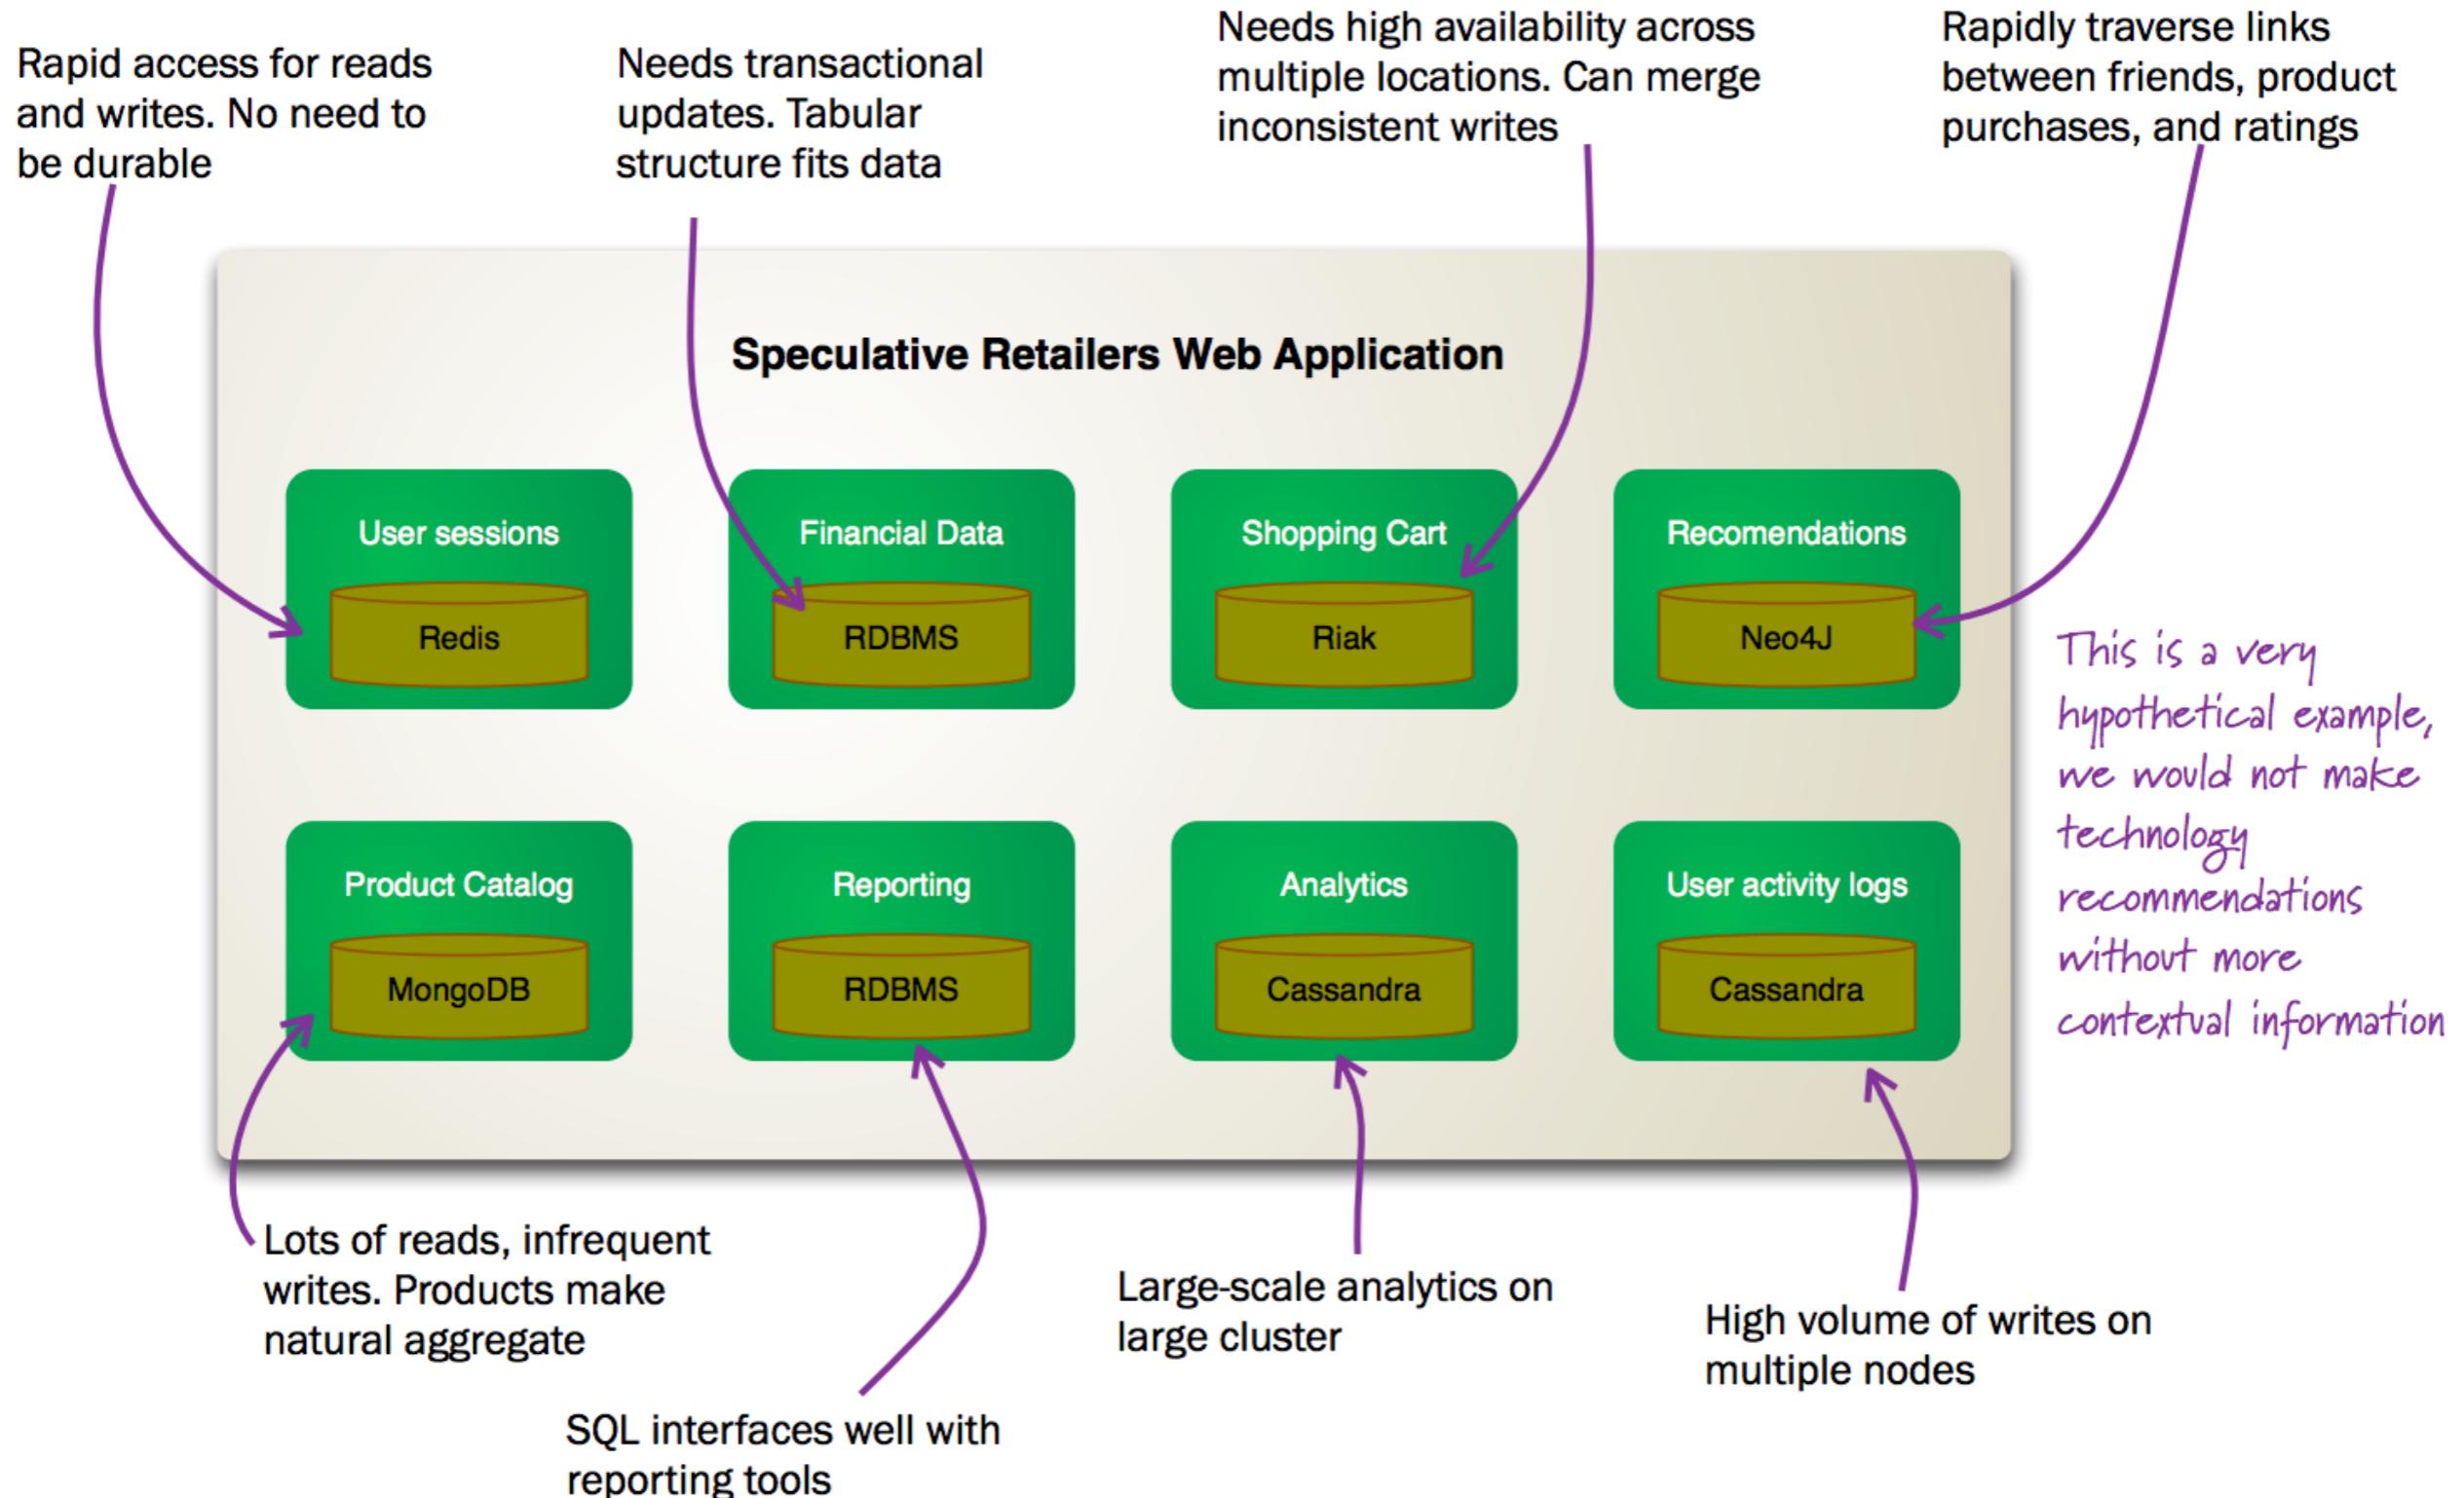
Series editor: Bruce A. Tate
Development editor: Jacquelyn Carter



SQL vs NoSQL

© 2017 - 2018 Siam Chamnankit Company Limited. All rights reserved.





Problems

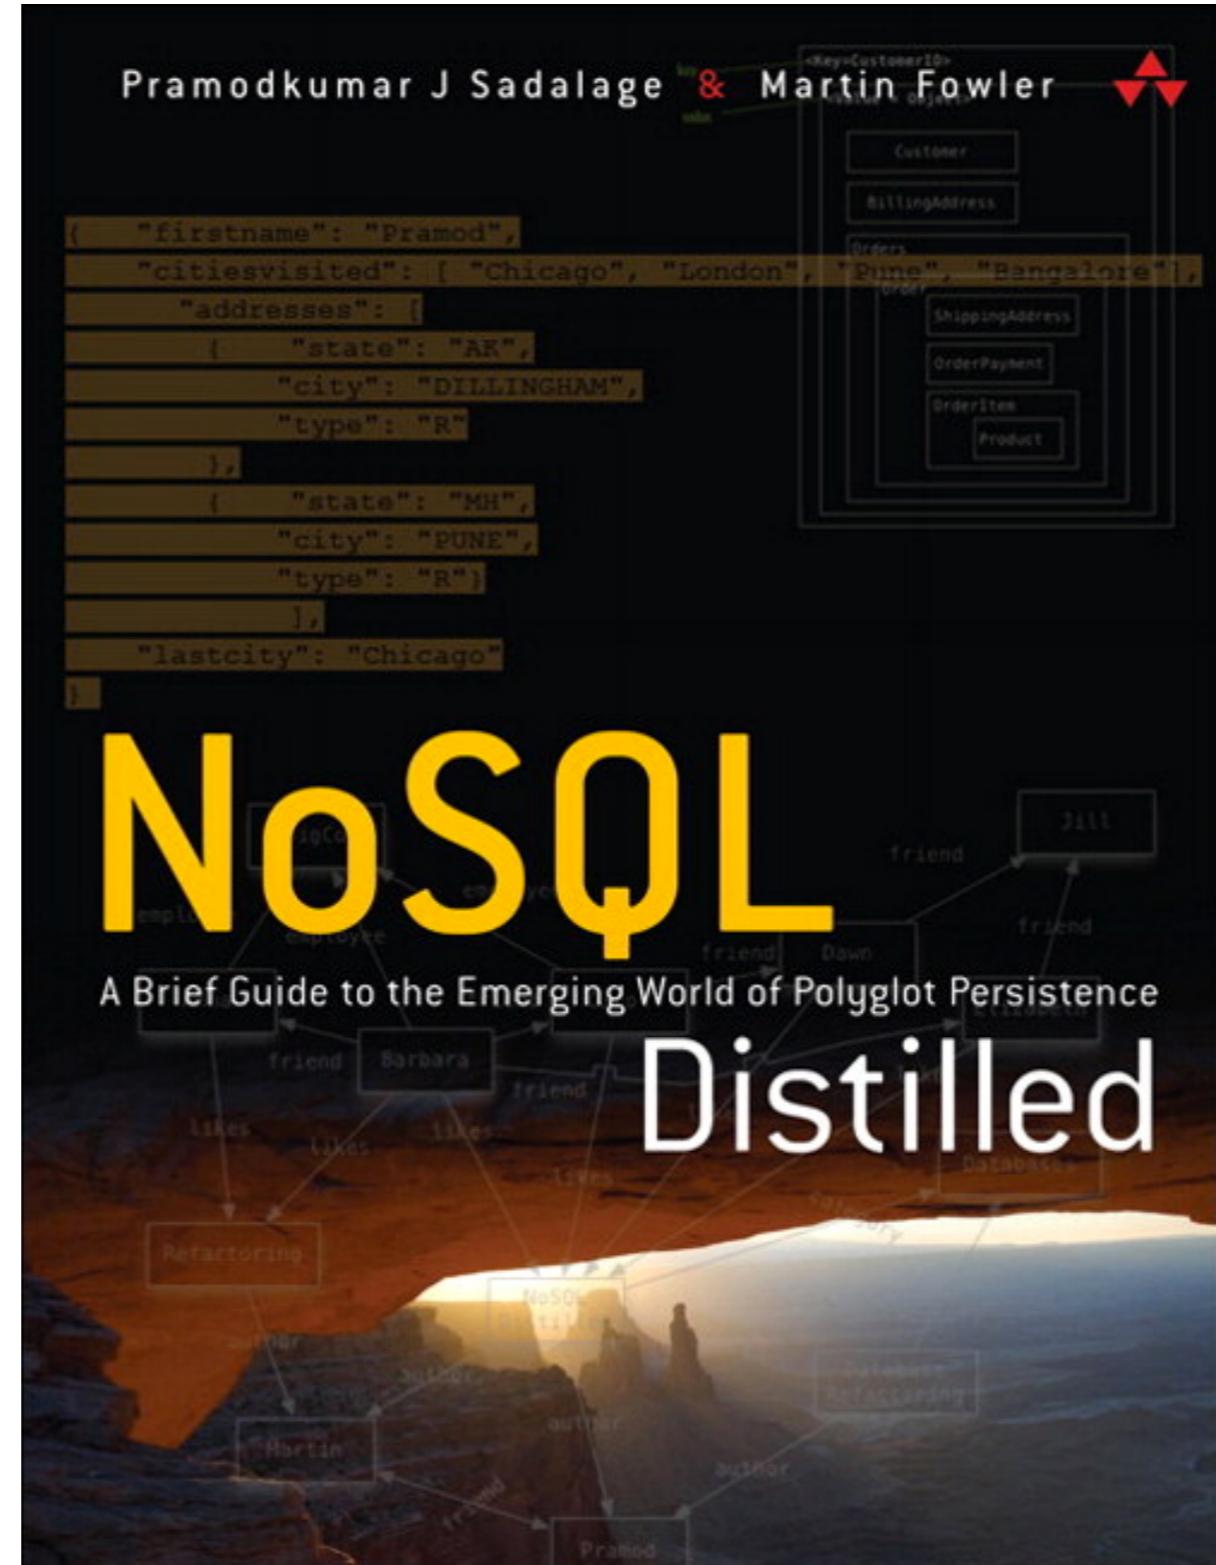
Decisions

Organization changes

Immaturity

Dealing with Eventual consistency





NoSQL + RDBMS





SCHOONER

NewSQL

JustOne_{DB}

VoltDB

xeround

NUODB

GENIEDB®

SQLFire

TRANS|ATTICE

H-Store

Clustrix





**KEEP
CALM**

AND

**use the right tool
for the right job**



Workshop



Workshop

MariaDB
PostgreSQL
Redis

Dealing with Eventual consistency



RDBMS and NoSQL



Redis

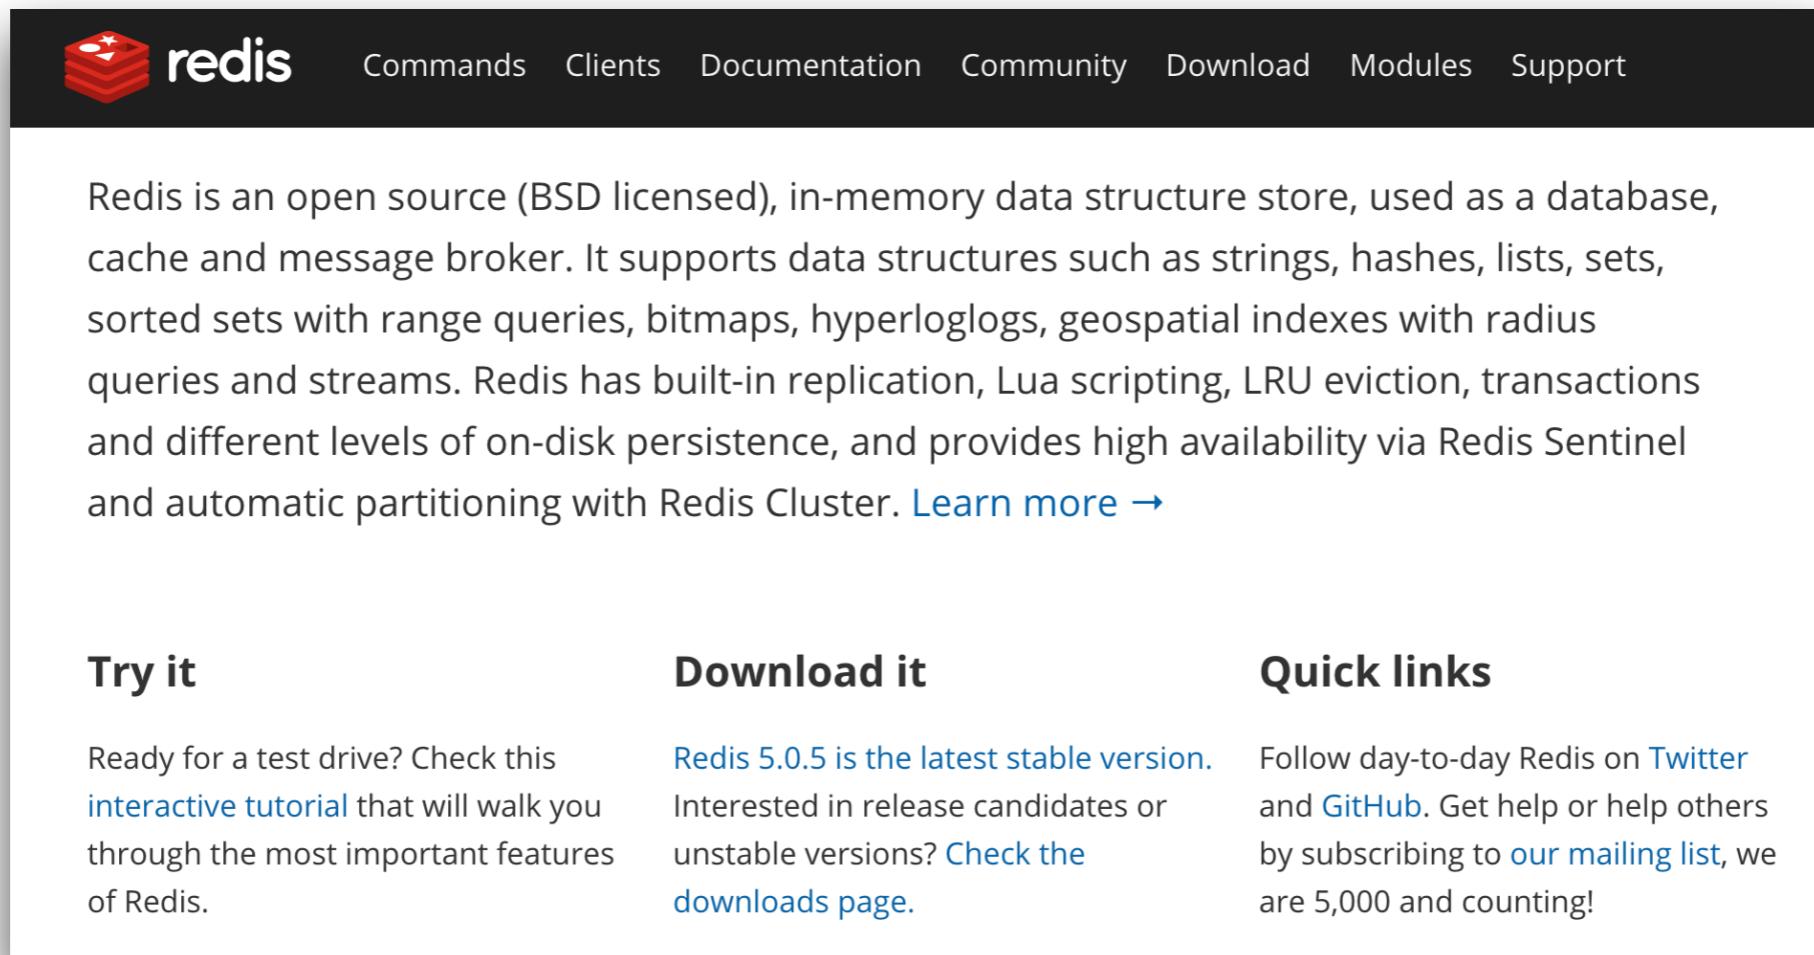
<https://redis.io/>



Redis

In-memory data structure store

Using for database cache and message broker



The screenshot shows the Redis homepage with a dark header bar containing the Redis logo and navigation links: Commands, Clients, Documentation, Community, Download, Modules, and Support.

The main content area contains a brief introduction to Redis, followed by three columns: Try it, Download it, and Quick links.

- Try it:** Ready for a test drive? Check this [interactive tutorial](#) that will walk you through the most important features of Redis.
- Download it:** Redis 5.0.5 is the latest stable version. Interested in release candidates or unstable versions? [Check the downloads page.](#)
- Quick links:** Follow day-to-day Redis on [Twitter](#) and [GitHub](#). Get help or help others by subscribing to [our mailing list](#), we are 5,000 and counting!

<https://redis.io/>

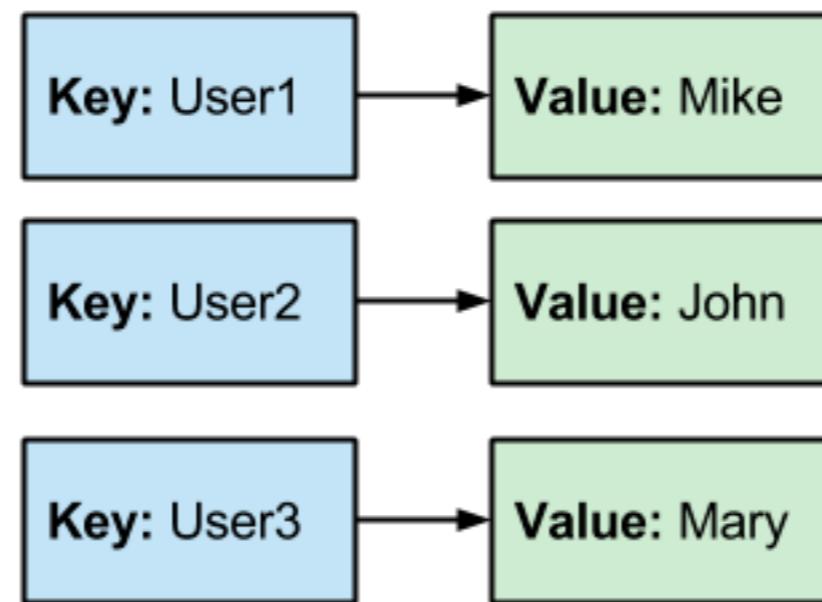


Redis

In-memory key-value store

Focus on performance and simplicity

All data keep in memory



Redis

REmote **D**Ictionary **S**erver

Data structure server

String, list, set, hash, sorted set

No index

No query language



Redis networking

Local or distributed access
Replication
Clustering



Redis networking

Simple Database



Master

Clustered Database



M1

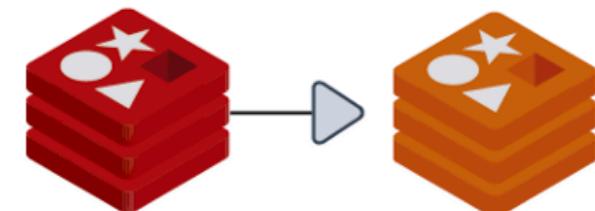


M2

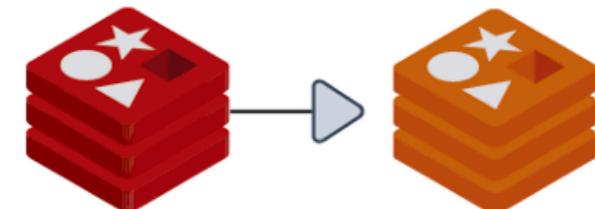


Mn

HA Clustered Database



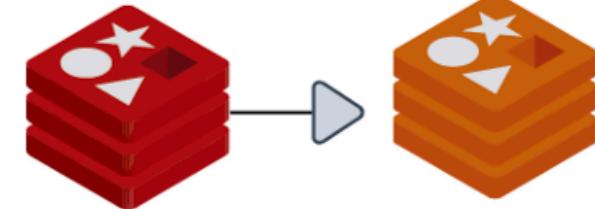
M1



S1



M2



S2

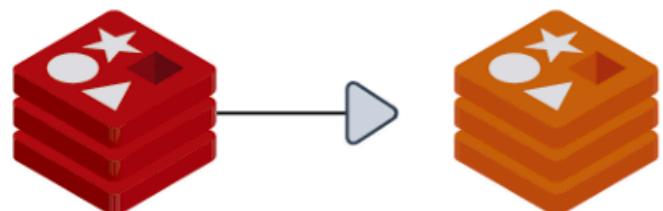


Mn



Sn

HA Database

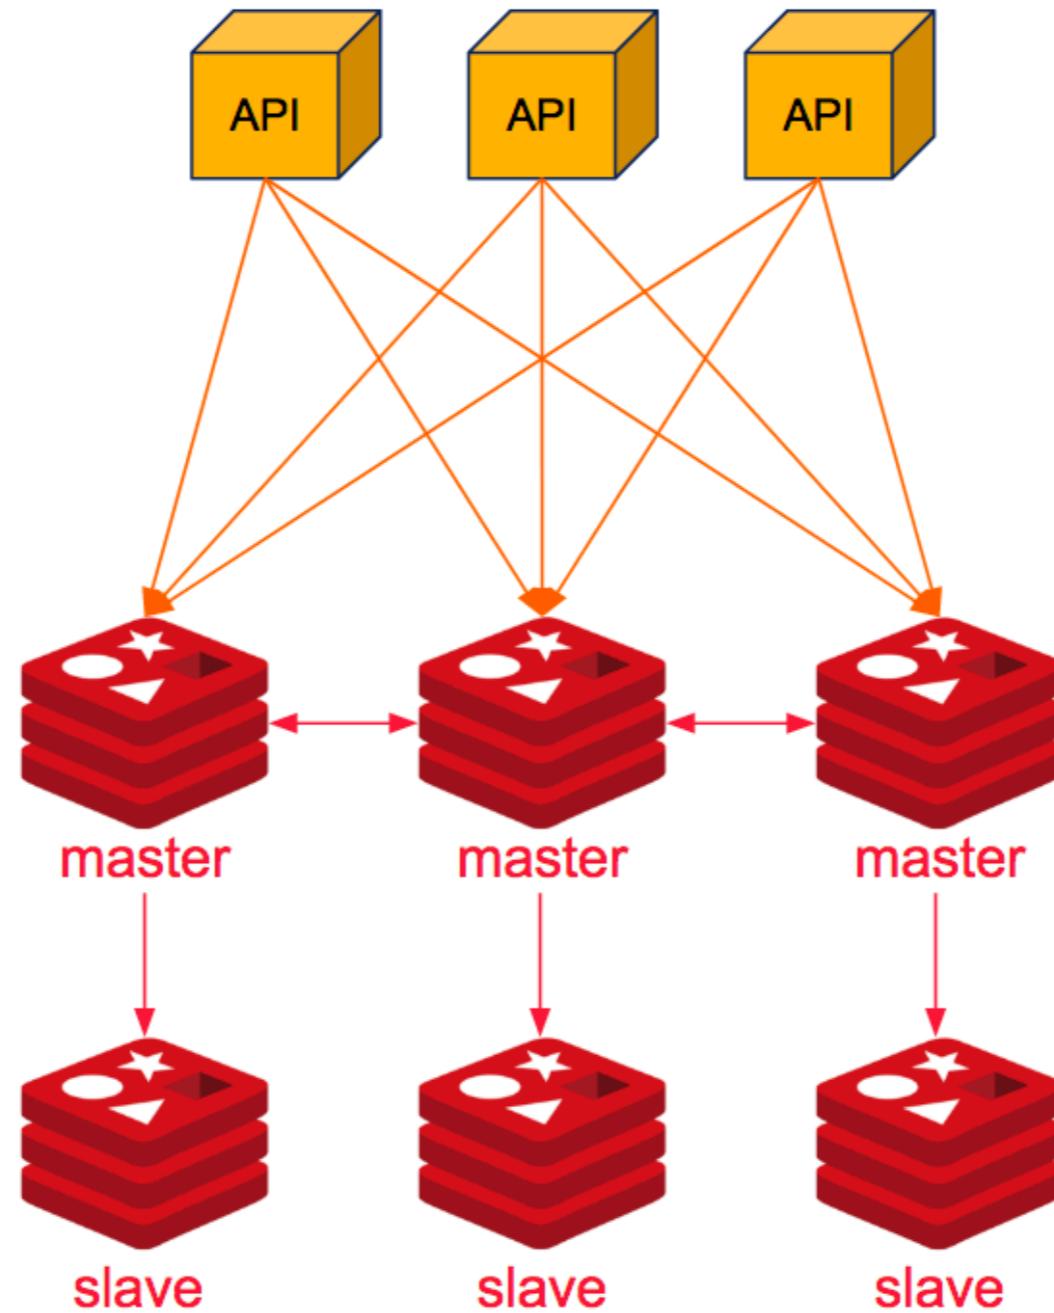


Master

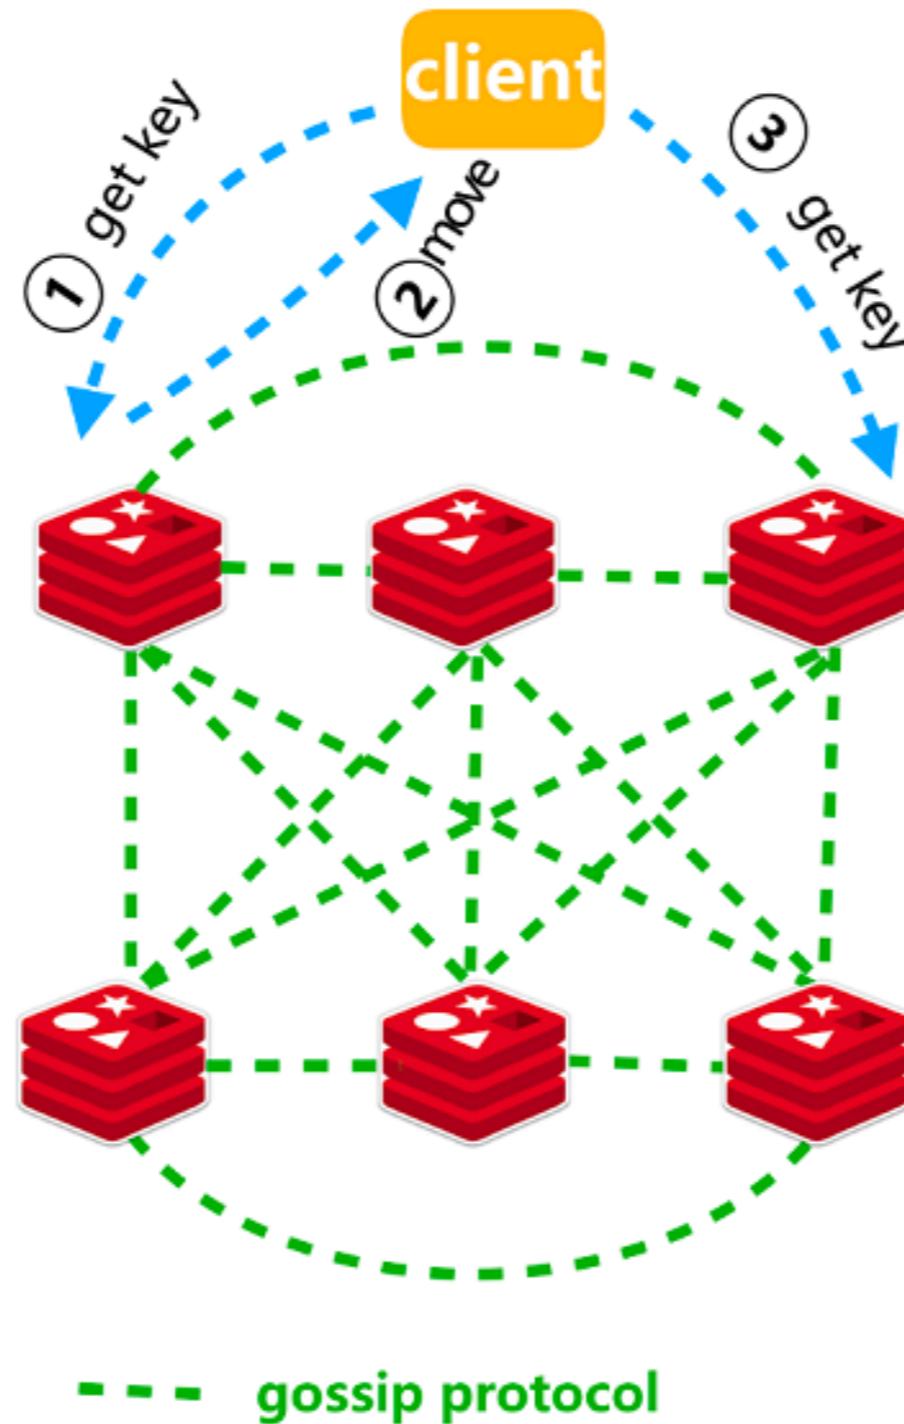
Slave



Redis replication



Redis cluster (Sentinel)



Durability options

Memory only

Append-only file (AOF)

Snapshot (RDB)



Memory vs AOF

Data-Persistence - The Wrong Way



Failed Instance

Data loss



New Empty Instance

Data-Persistence - The Right Way



Failed Instance

No Data loss



New Populated Instance

<https://redislabs.com/redis-enterprise/technology/durable-redis-2/>



Append-Only-File

Append new data to file system
Every second (fast but less safe)
Every write (safer but slower)



Snapshot

The entire point-in-time view of the dataset is written to persistent storage, across all shards of the database

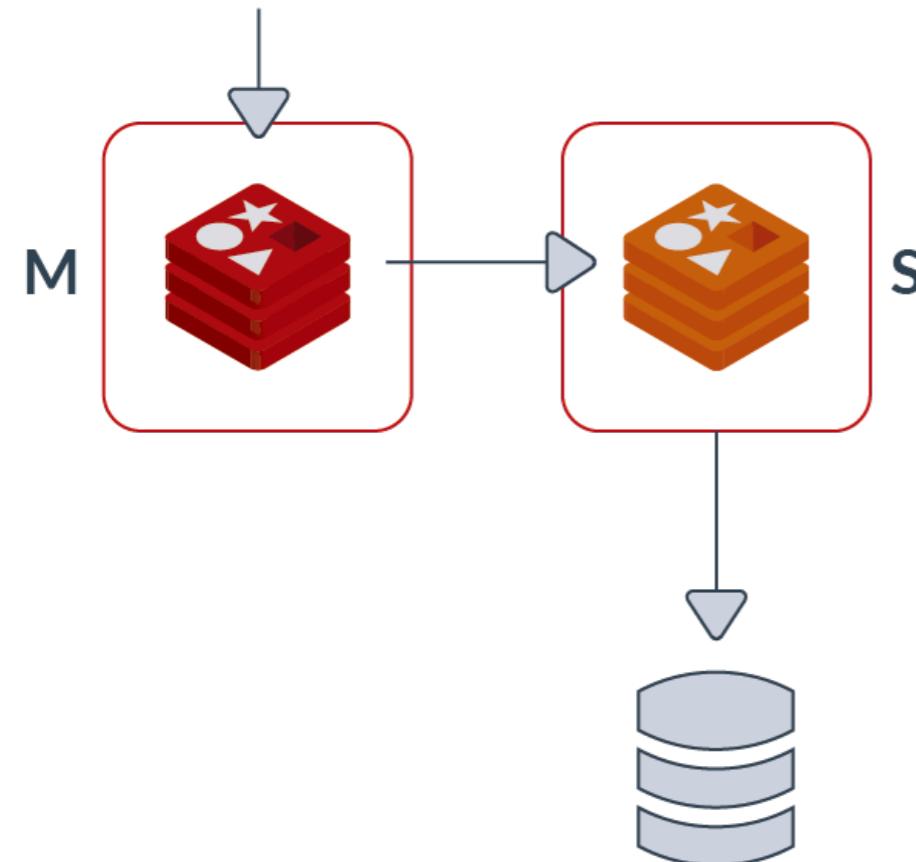
The snapshot time is **configurable**



Tuning for Master-Slave

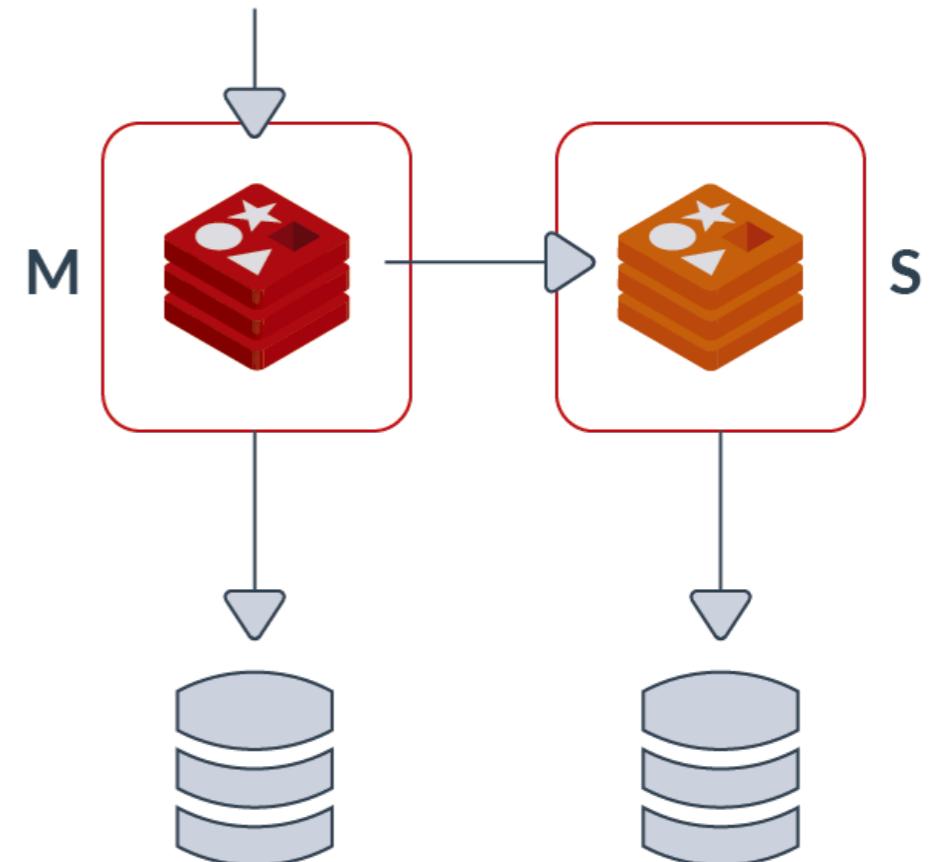
Tuned for Speed

Data-Persistence at the slave



Tuned for Reliability

Data-Persistence at the master & slave



Snapshot vs Backup



Snapshot vs Backup

Snapshot

supports data durability

(i.e. to automatically recover data when there is no copy of the dataset in memory)

Backup

supports disaster recovery

(i.e. when the entire cluster needs to be rebuilt from scratch)



Redis workshop

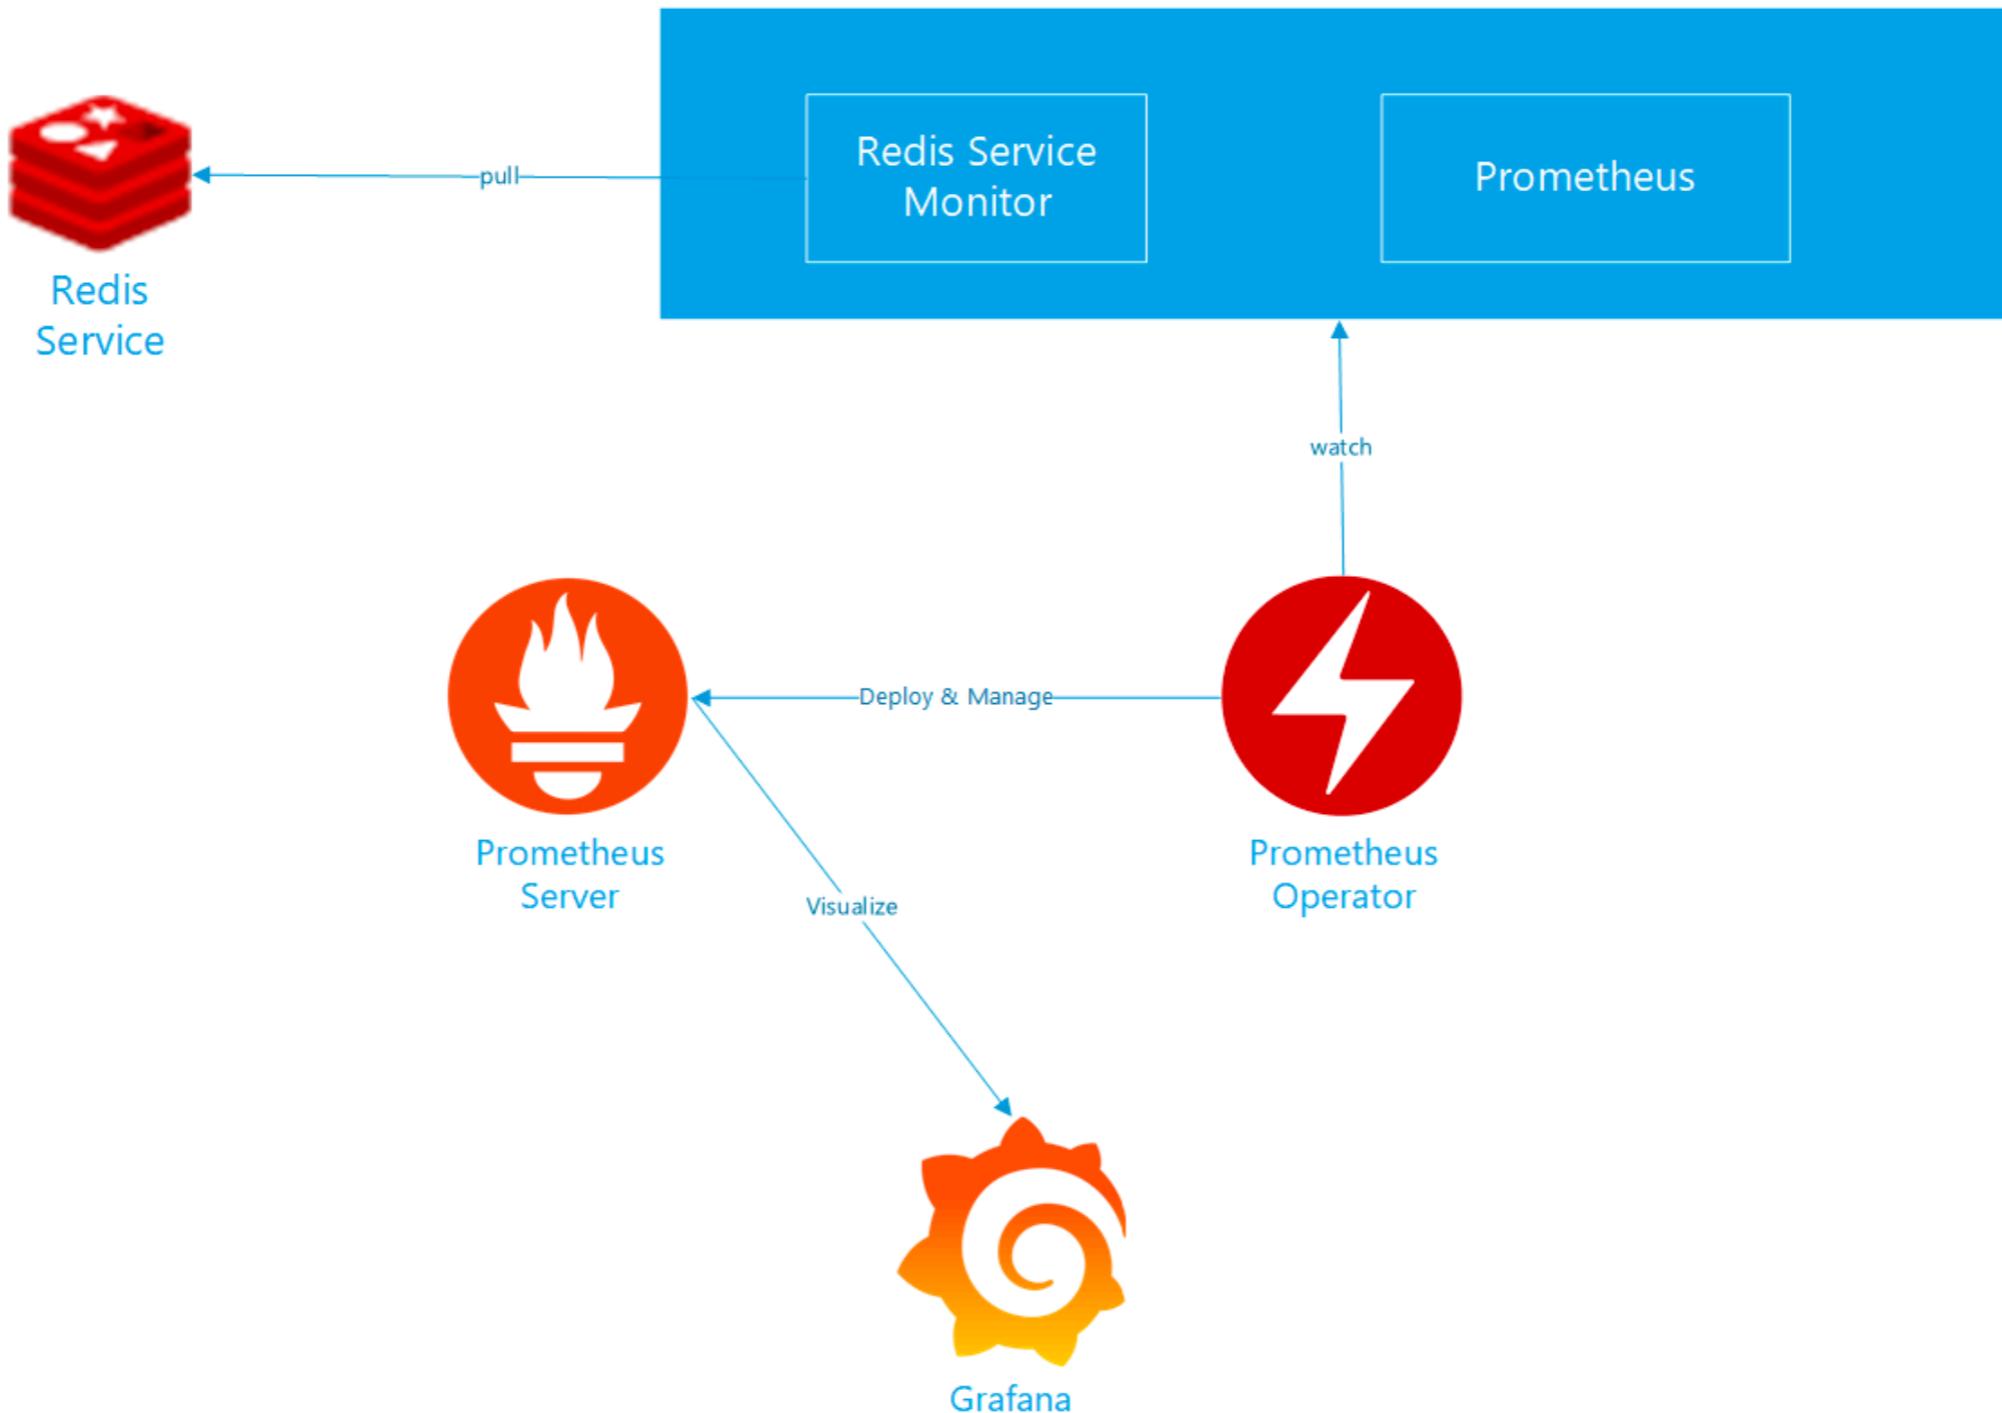
Basic of Redis operation

Clustering

Monitoring



Redis monitoring



PostgreSQL

<https://redis.io/>



PostgreSQL

RDBMS
Opensource
Active community

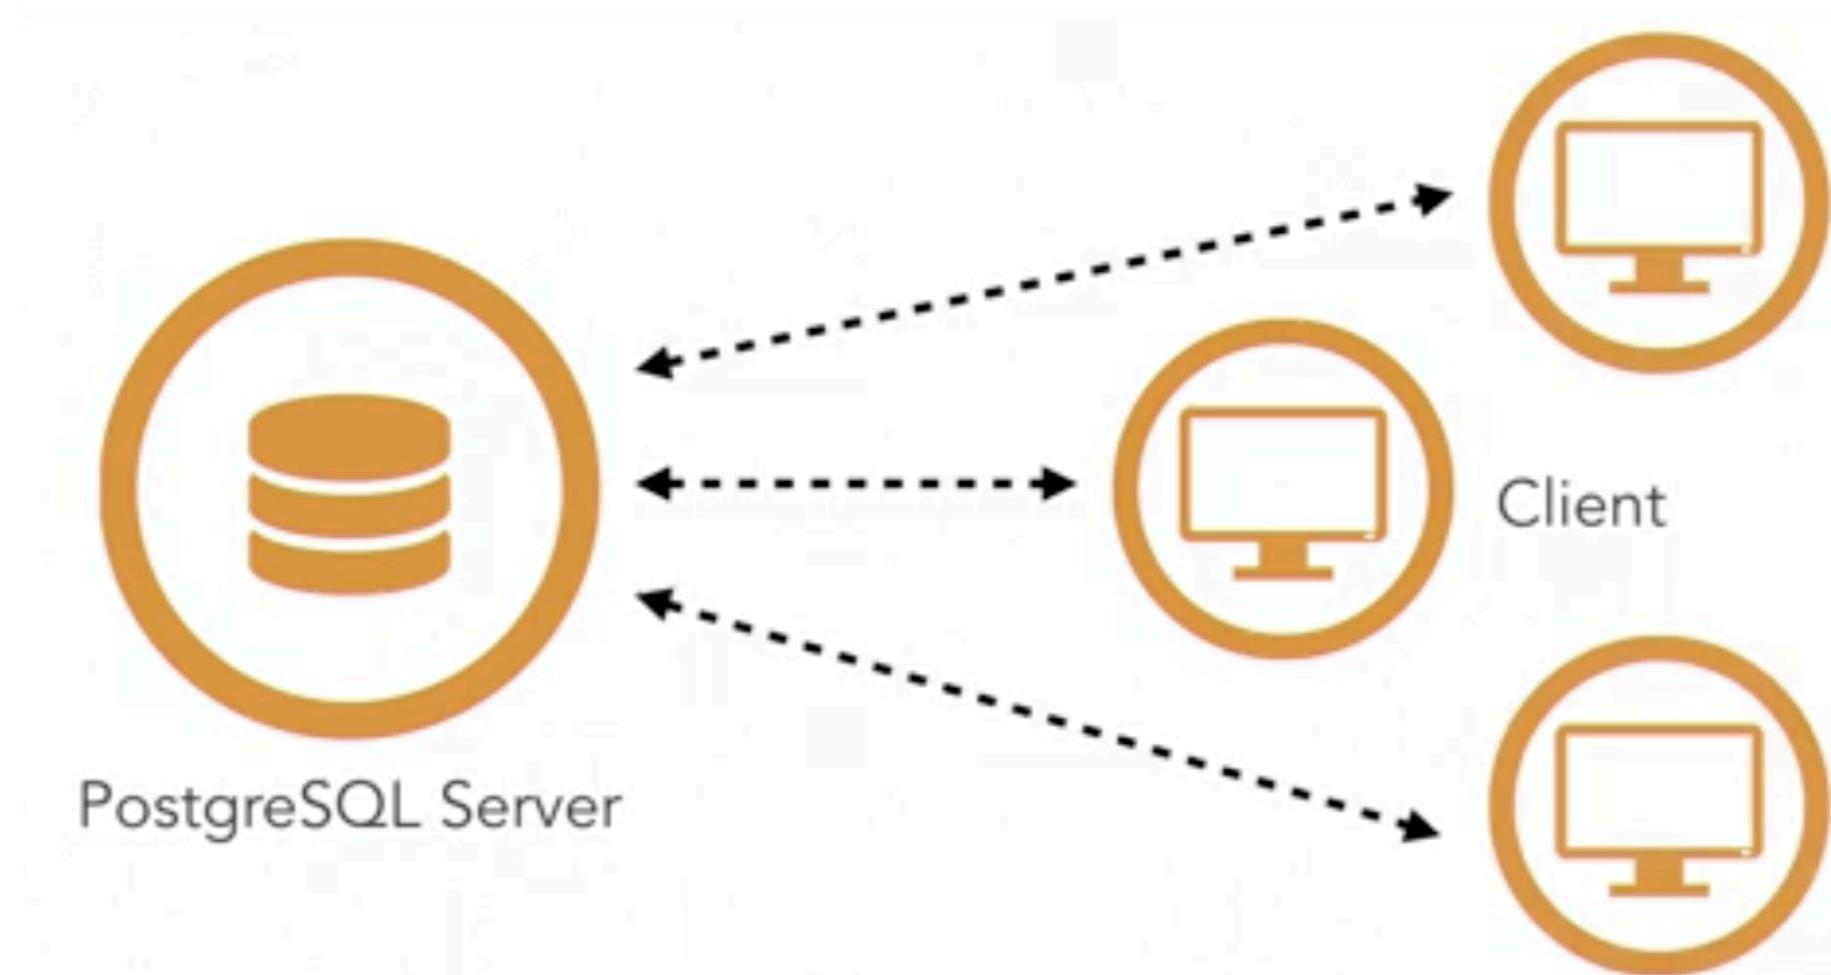


PostgreSQL

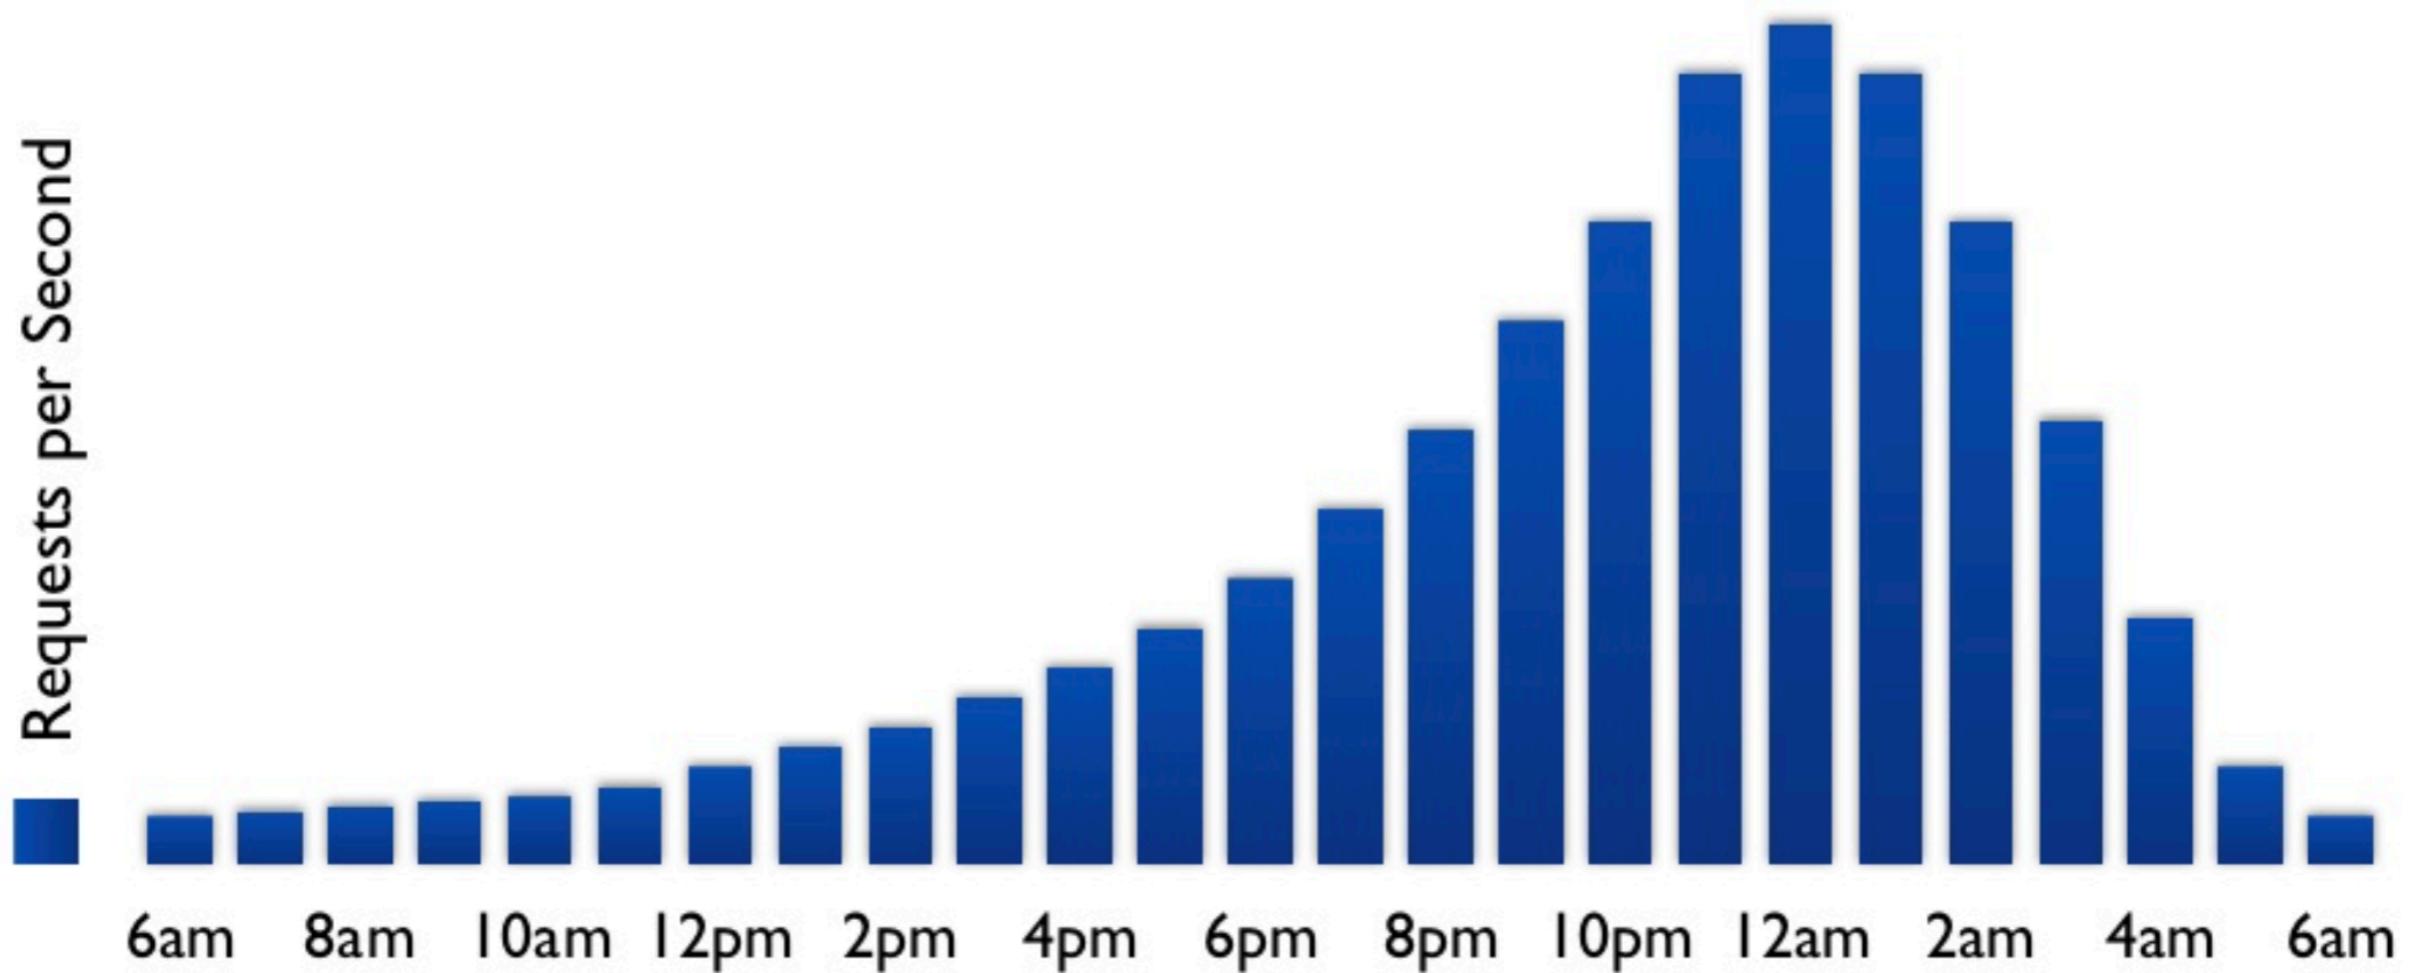
Rank			DBMS
Dec 2020	Nov 2020	Dec 2019	
1.	1.	1.	Oracle 
2.	2.	2.	MySQL 
3.	3.	3.	Microsoft SQL Server 
4.	4.	4.	PostgreSQL 
5.	5.	5.	MongoDB 



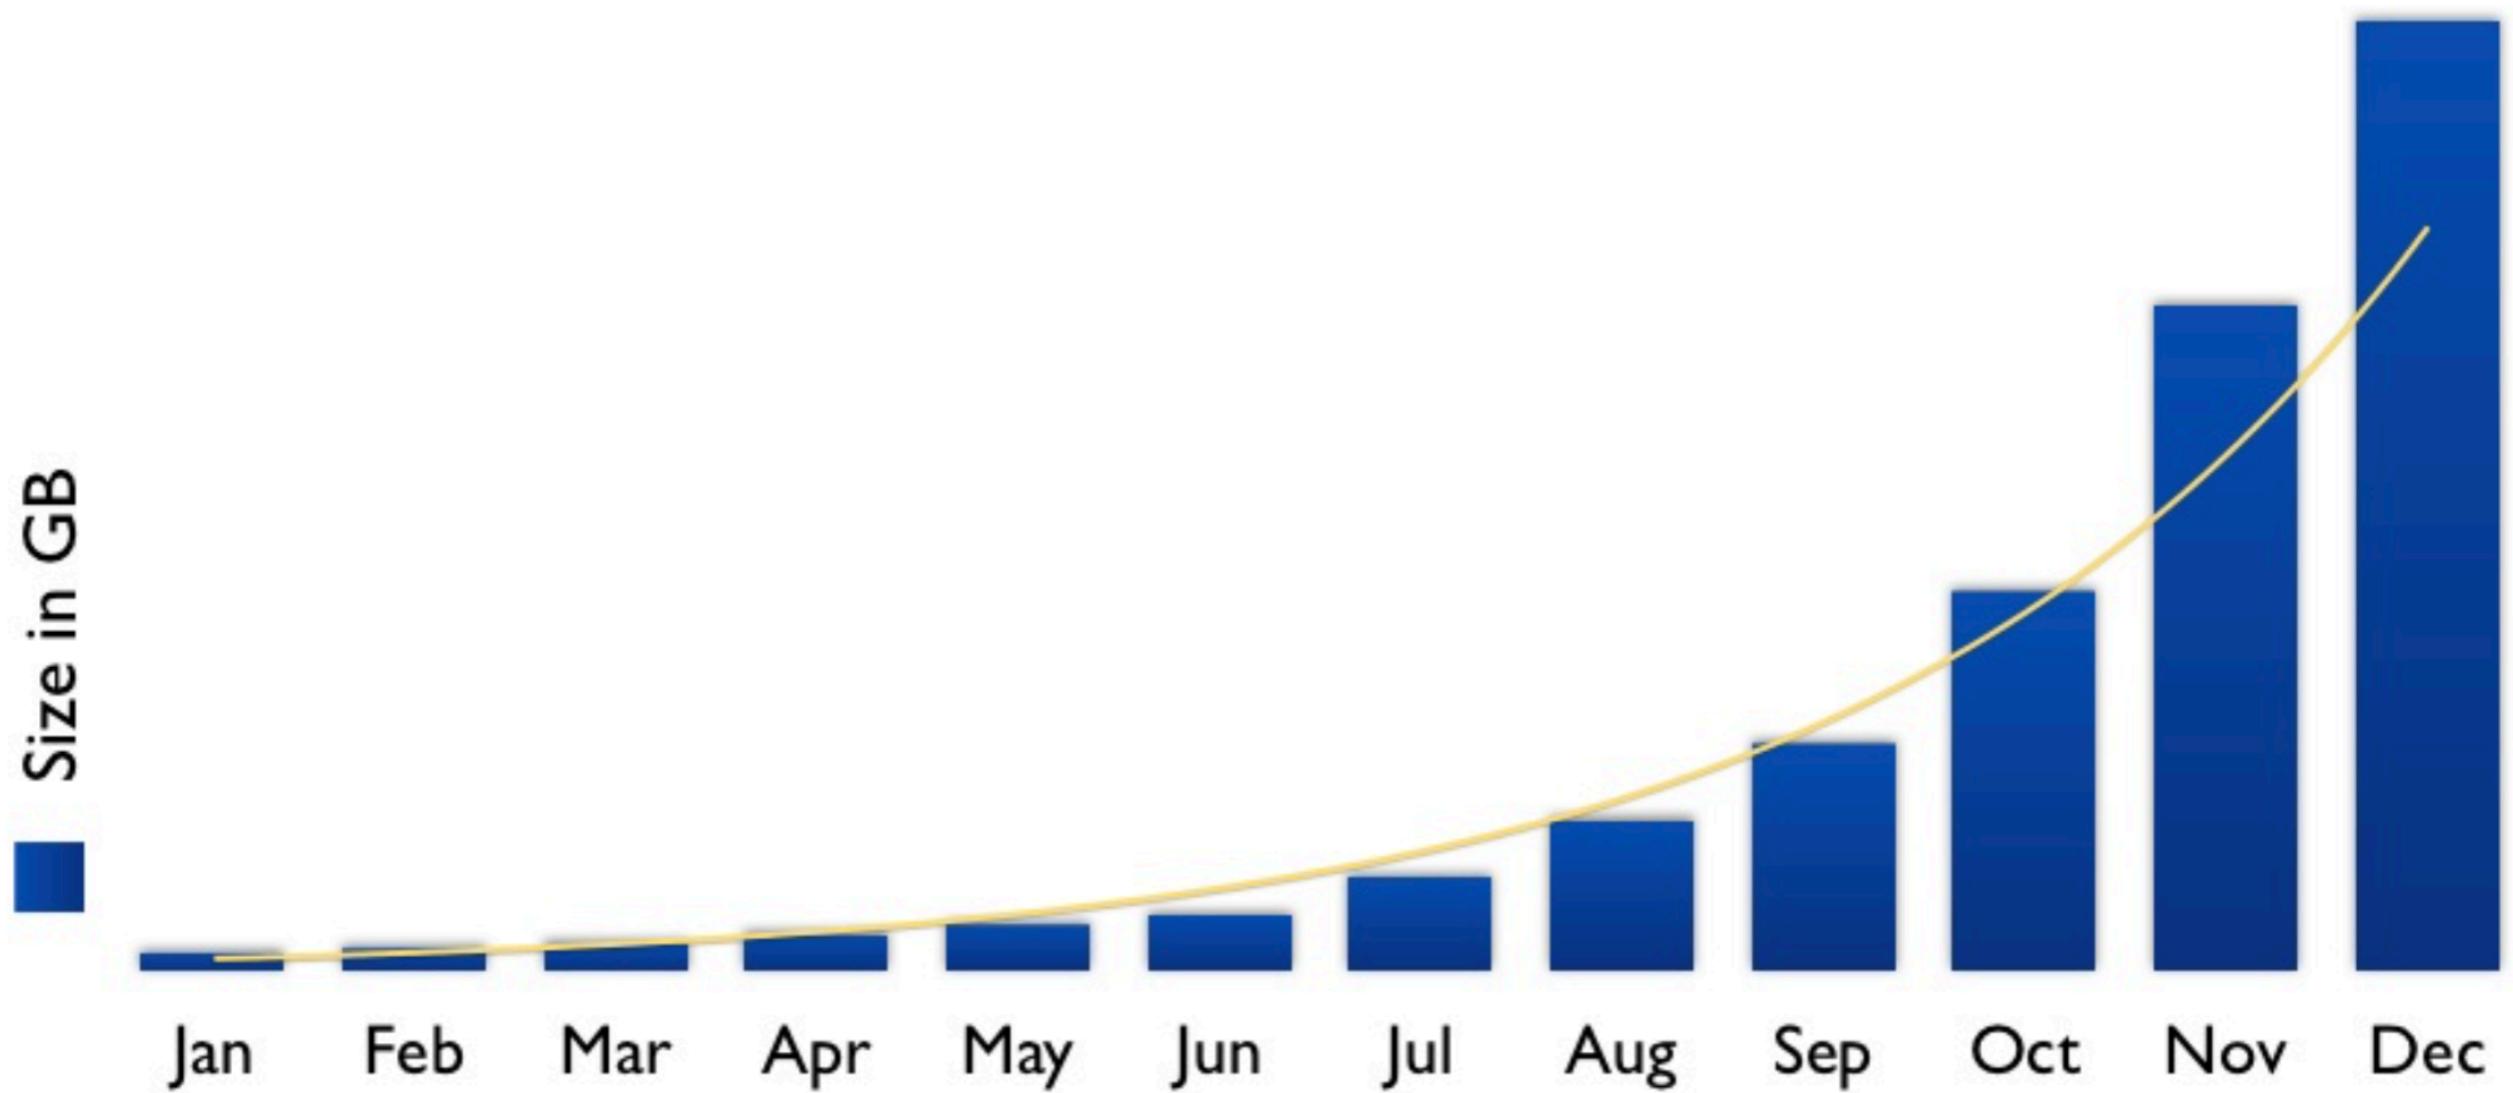
Client/Server model



Concurrency

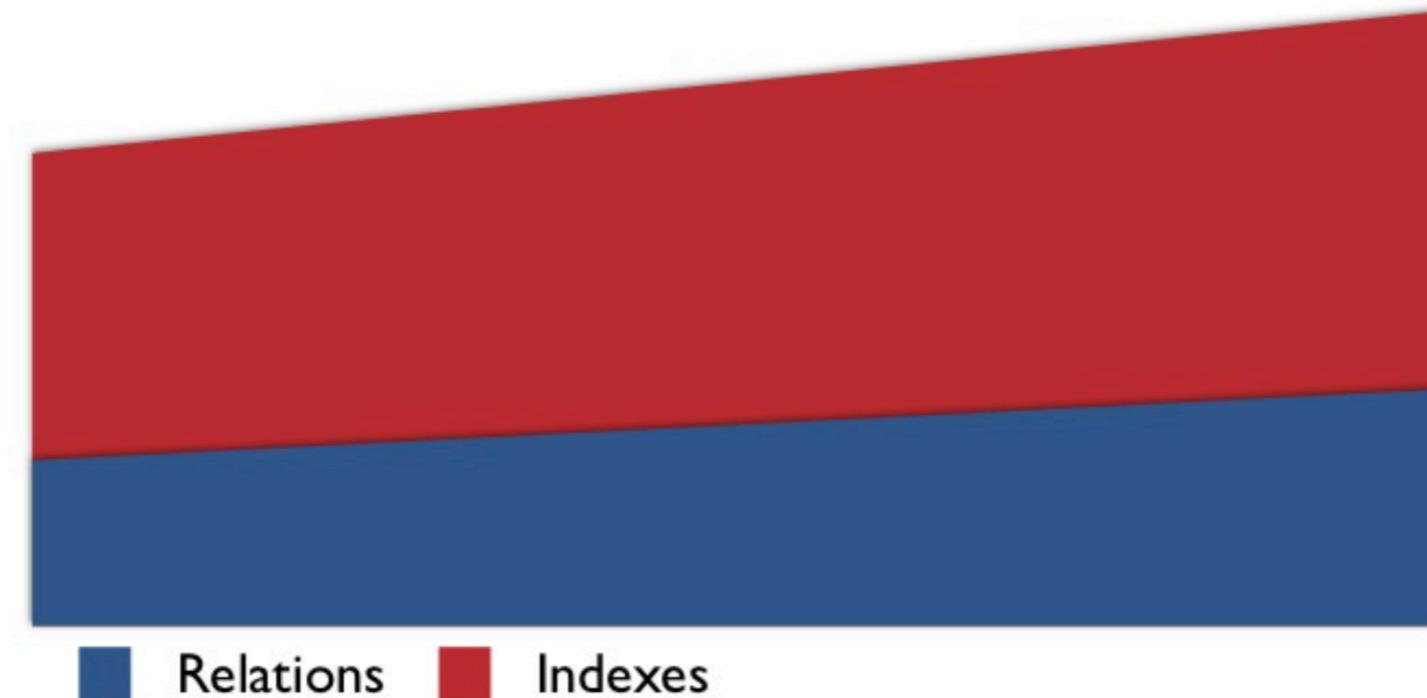


Data on disk



Data on disk

Tables sizes and indexes



Constraints

Available memory

Disk speed

I/O speed



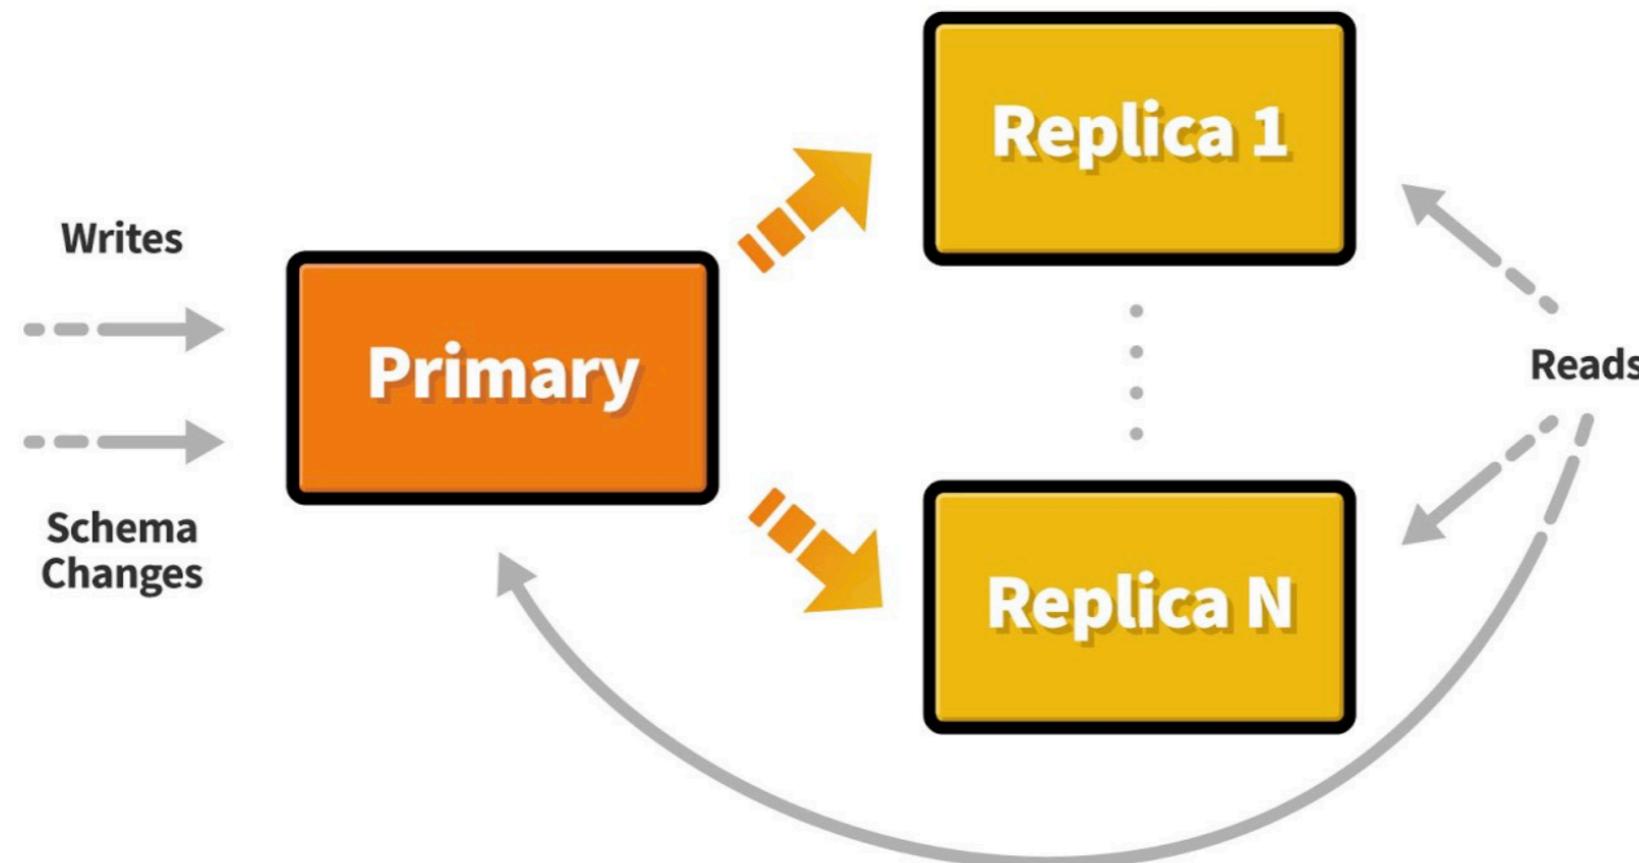
Scalability

Higher insert performance
Higher read performance
High availability

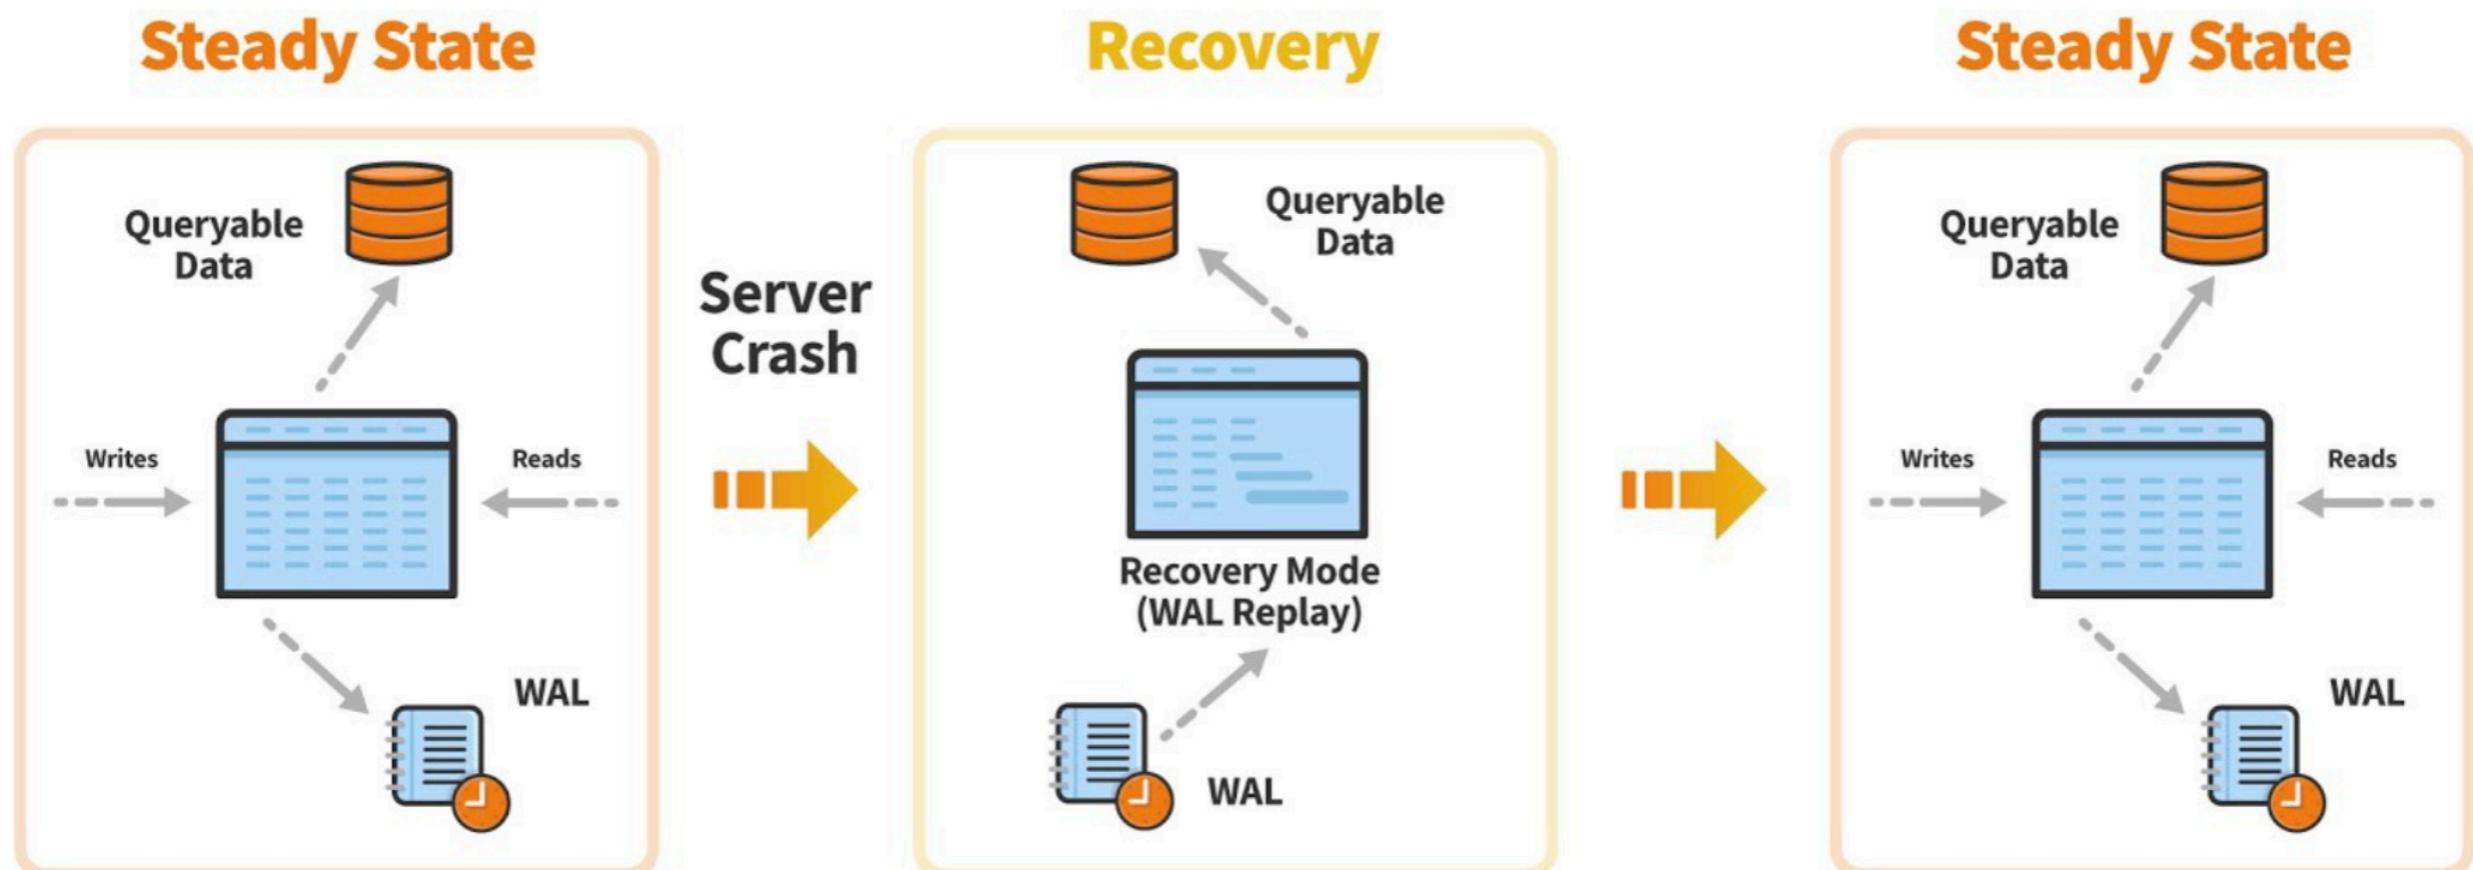


How an (N+1) node PostgreSQL cluster works ?

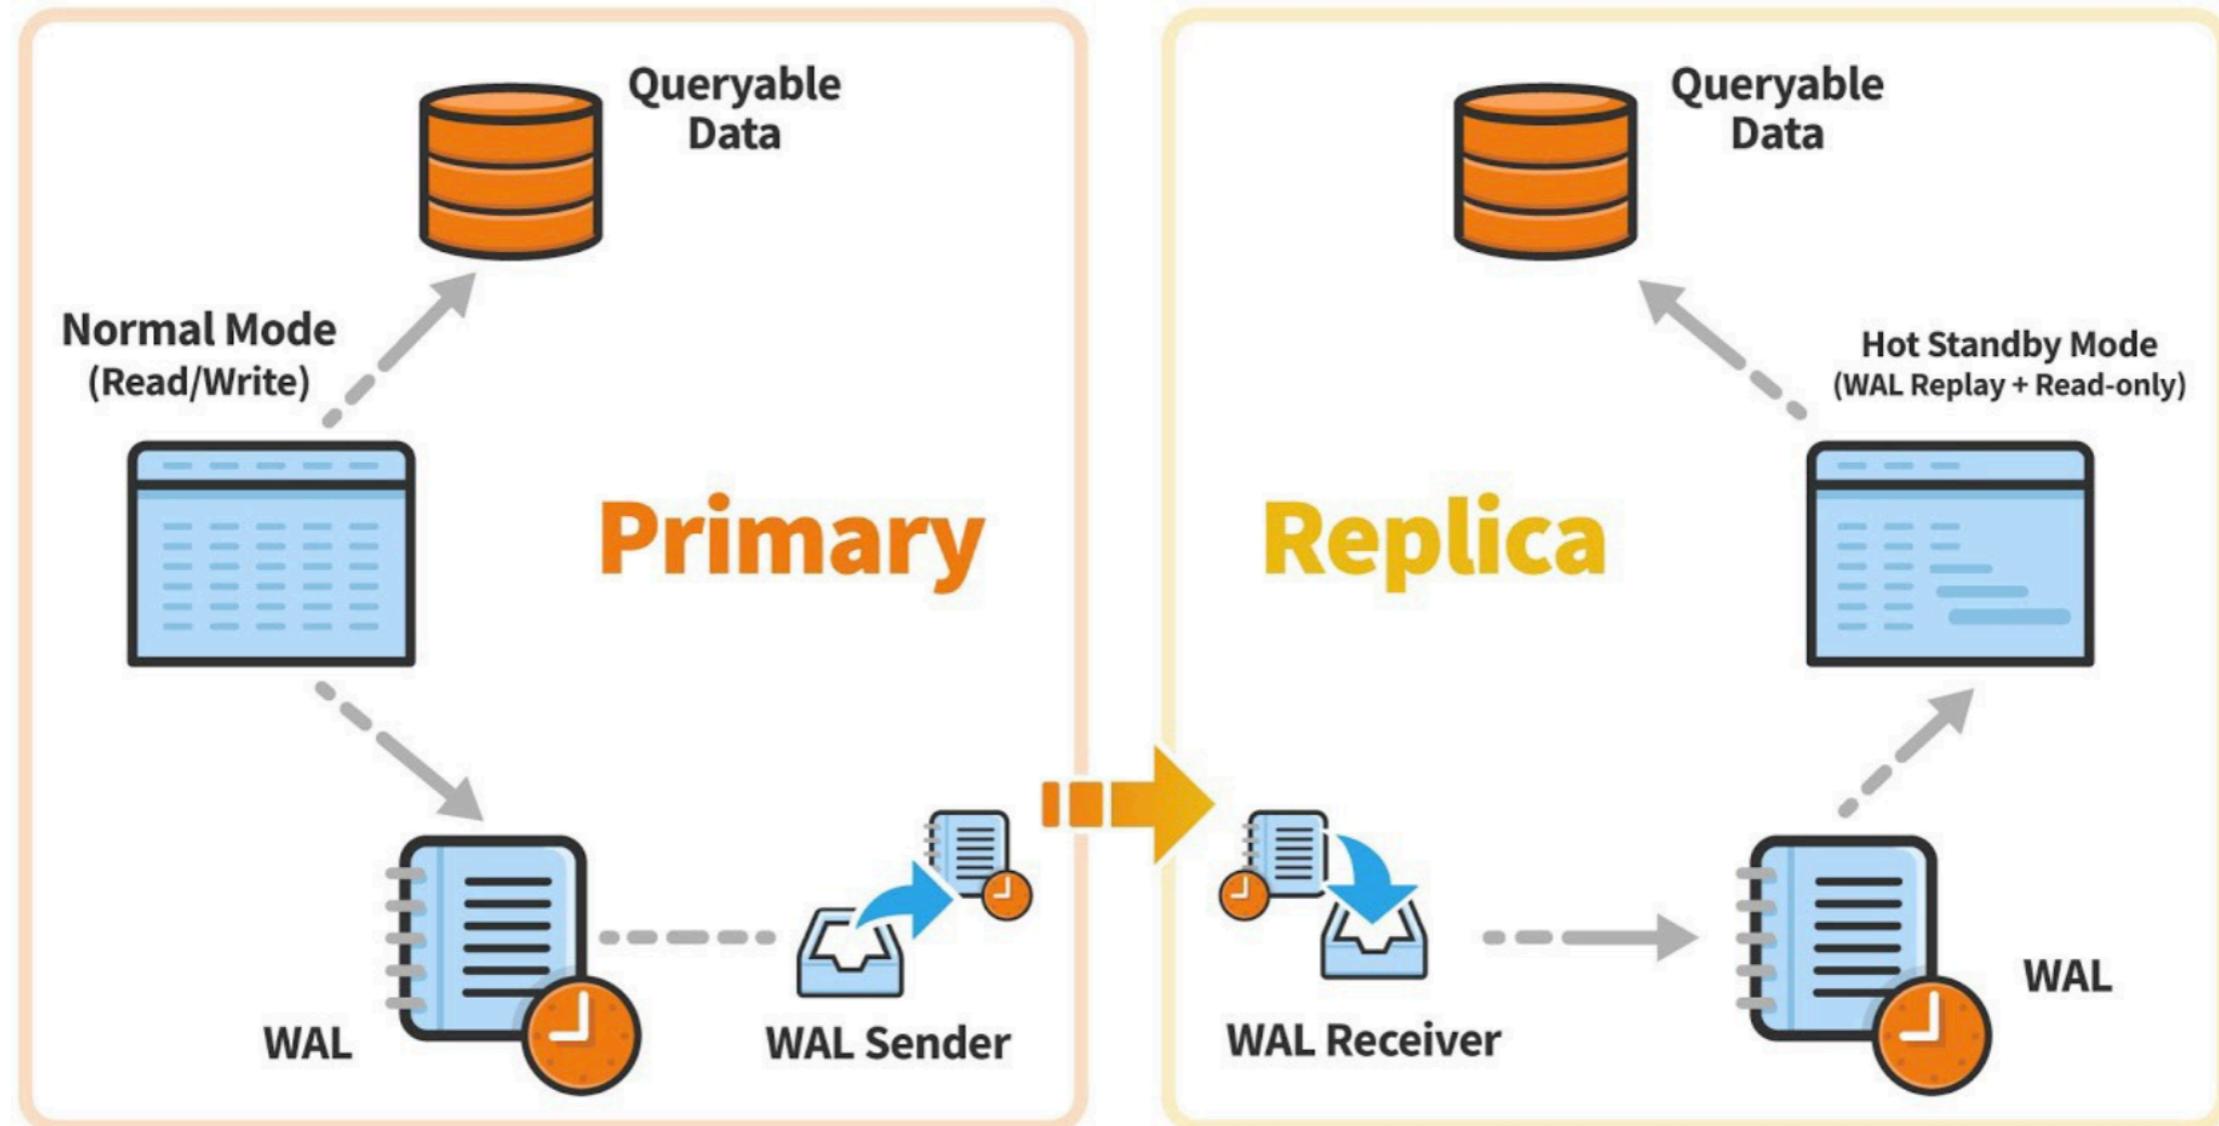
Using streaming replication



WAL (Write Ahead Log)



WAL and Replication



Replication modes

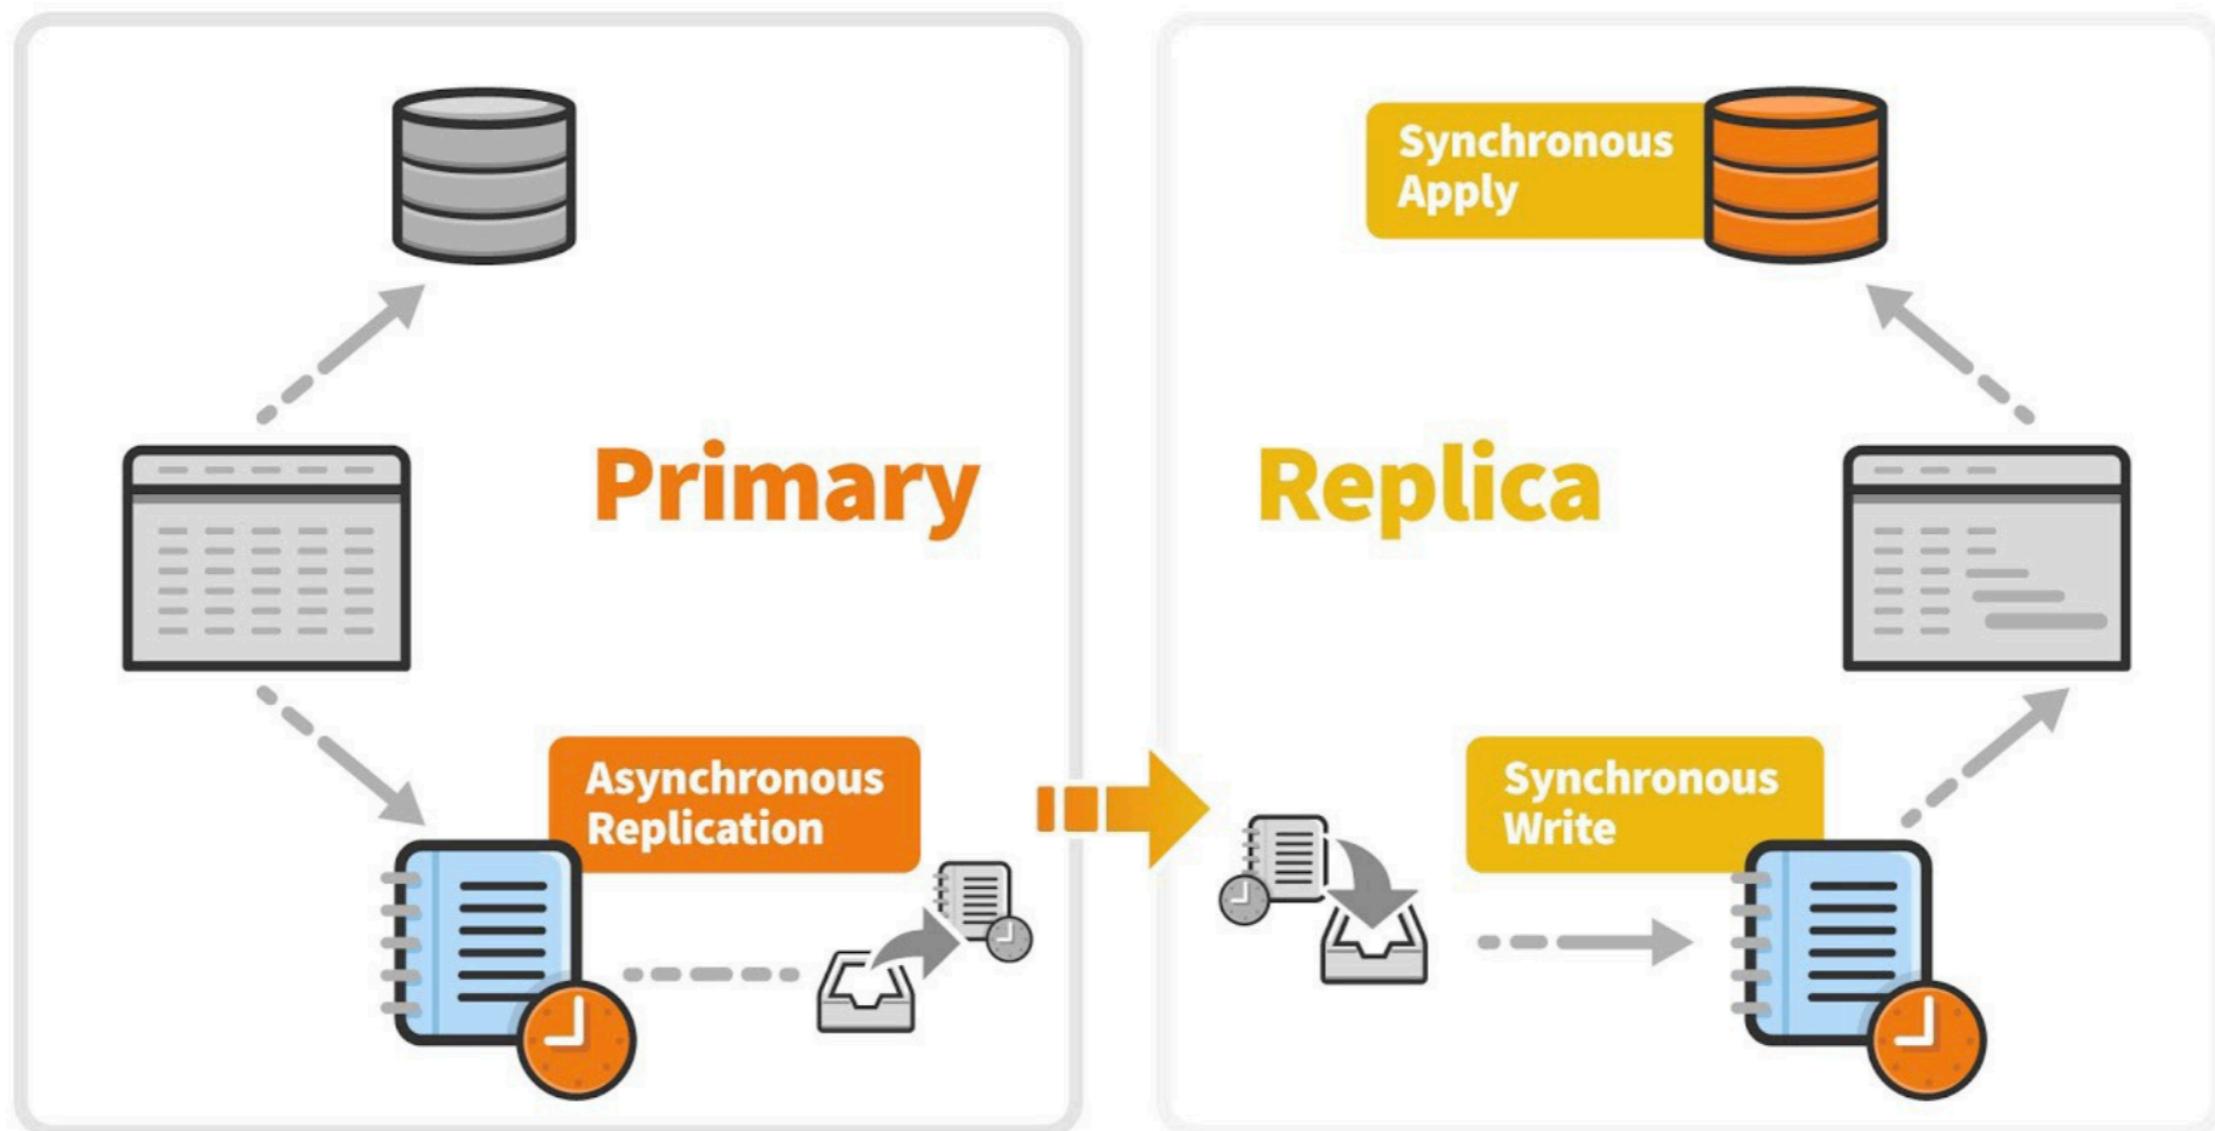
Asynchronous replication

Synchronous write replication

Synchronous apply replication



Replication modes



Performance

Replication mode	Write performance	Read consistency	Data loss
Async replication	Highest	Weakest, eventually consistent	Highest risk
Sync write replication	Medium	eventually consistent BUT less lag	Low risk
Sync apply replication	Lowest	Strong consistency	Lowest risk

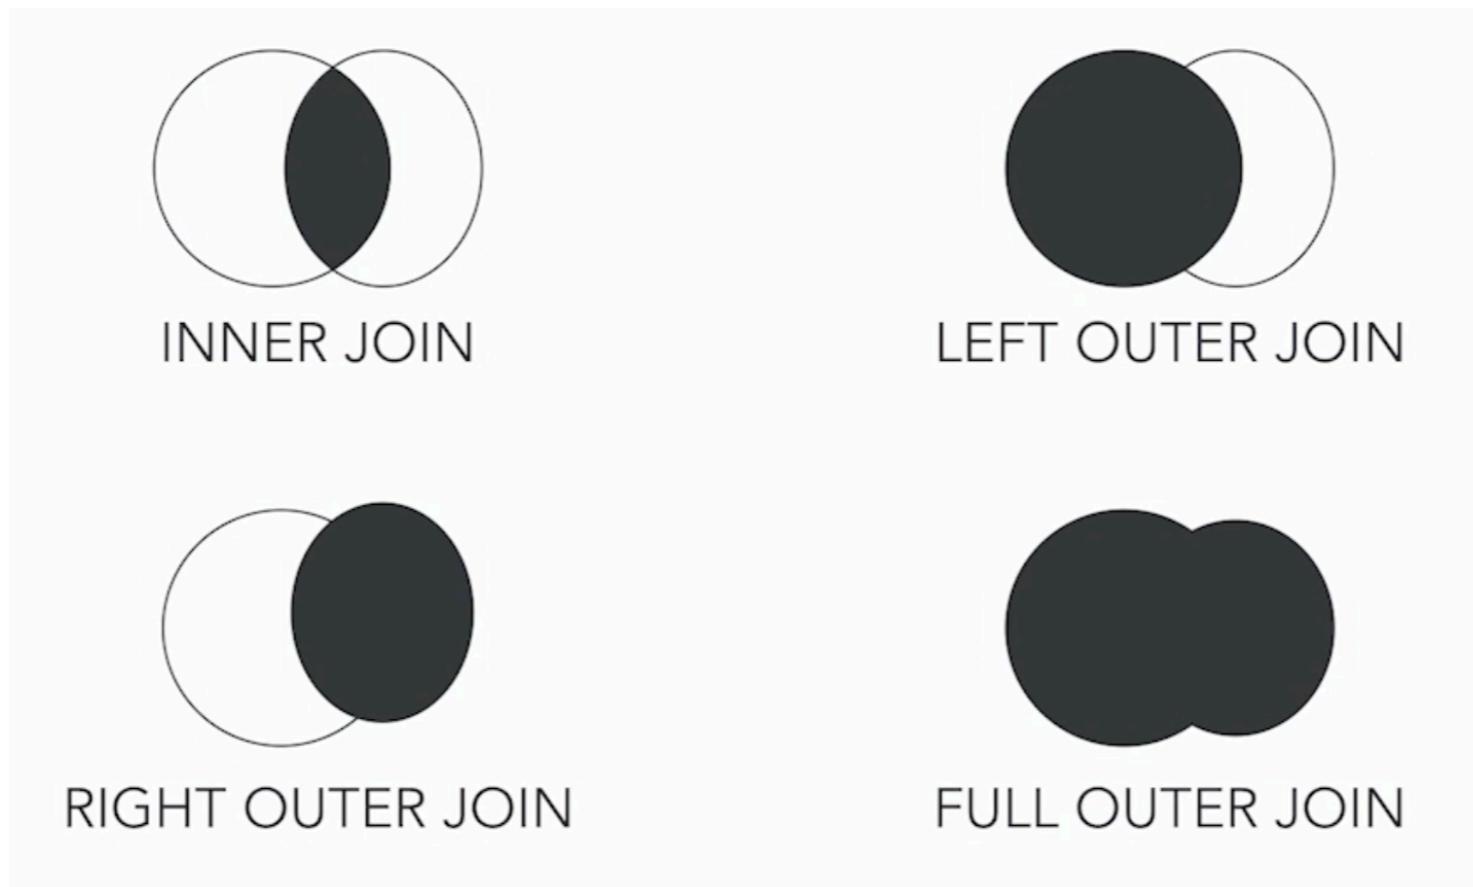


SQL Tuning with PostgreSQL



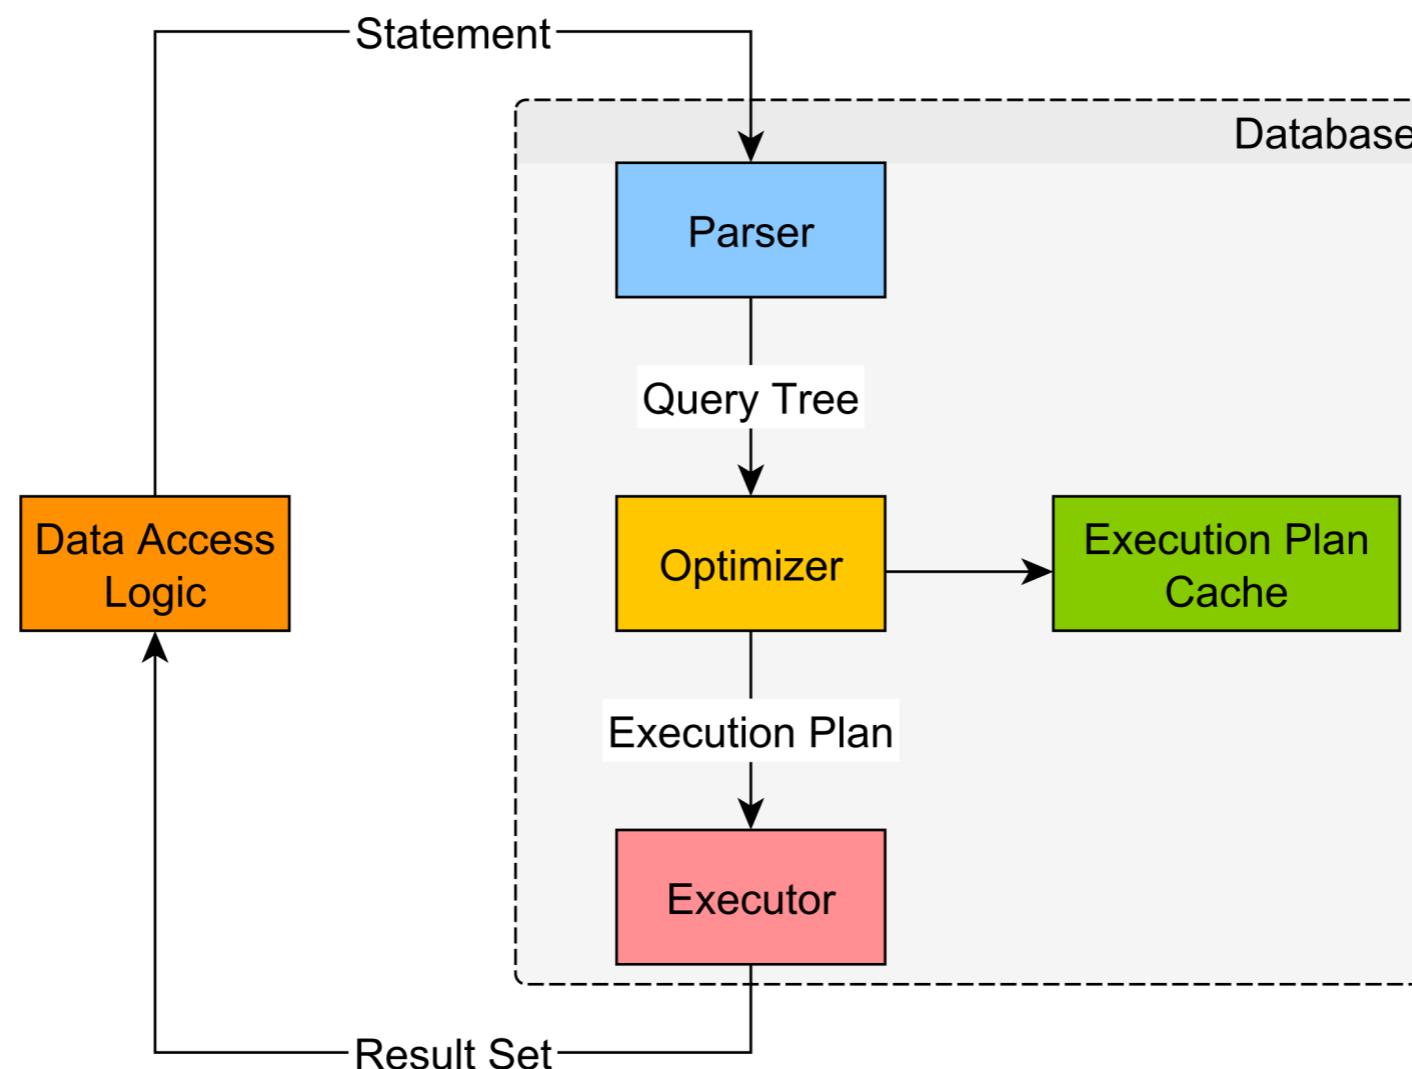
Tuning ?

Reduce query response time with query tuning
Explain query ?



SQL use execute plan

Set of step that scan, filter and join data



Efficient execution plan ?

Good or Bad ?

Good

Using index to find data

Bad

Scan 1M rows to find 100 rows



Step to learn

Understand query plan steps

Understand trade-offs in way to implement
query plan

Learn techniques that lead to efficient plans



Basic

Scanning tables and indexes

Joining tables

Partition data



Scanning tables



customer_id	fname	lname	last_purchase_date	status_level
1234	Alice	Smith	23-Jan-2016	silver
2345	Bob	Washington	14-July-2018	bronze
2352	Charlie	Johnson	20-Dec-2018	gold
2678	Chendong	Mao	19-Dec-2017	gold
3422	Javier	Valdez	23-Dec-2018	gold
3577	Avery	Winston	5-Jan-2019	silver
4240	Charles	Davis	15-Jan-2017	bronze
6235	Vihaan	Patel	12-Nov-2018	gold
2352	Katherine	Coltrane	04-Apr-2016	silver



Scanning tables

Scanning looks at each row

Fetch data block containing row

Apply filters or conditions

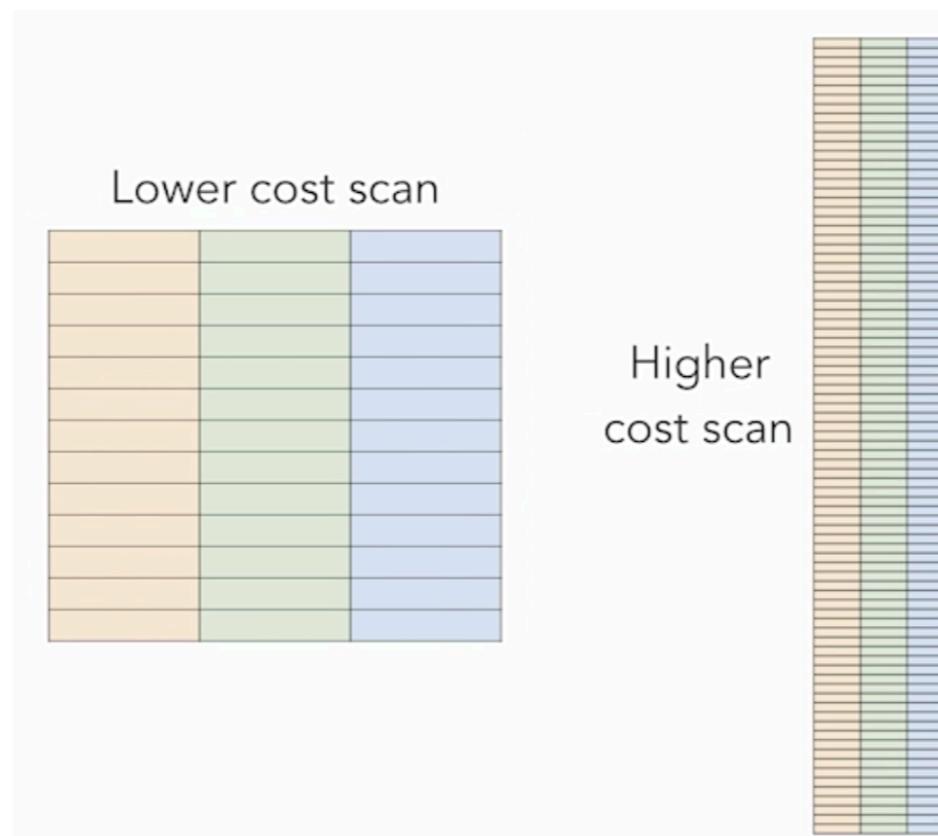
**Cost of query based on number of rows in
the table**



Cost based on number of rows

Scanning on small tables is efficient

Scanning on large tables can be efficient if
have few queries



Indexing reduce full table scan

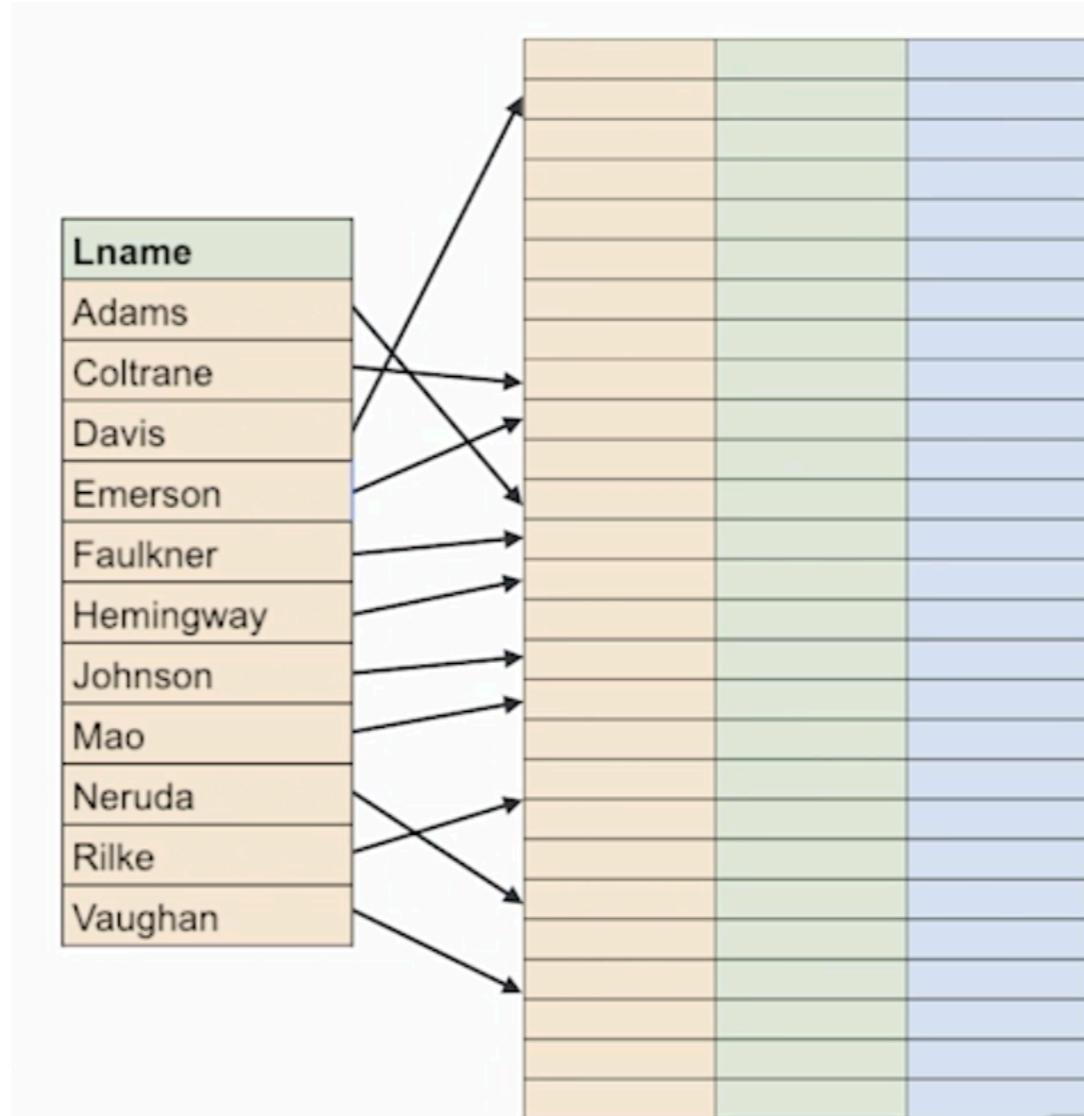
Indexes are ordered

Faster to search index for an attribute value

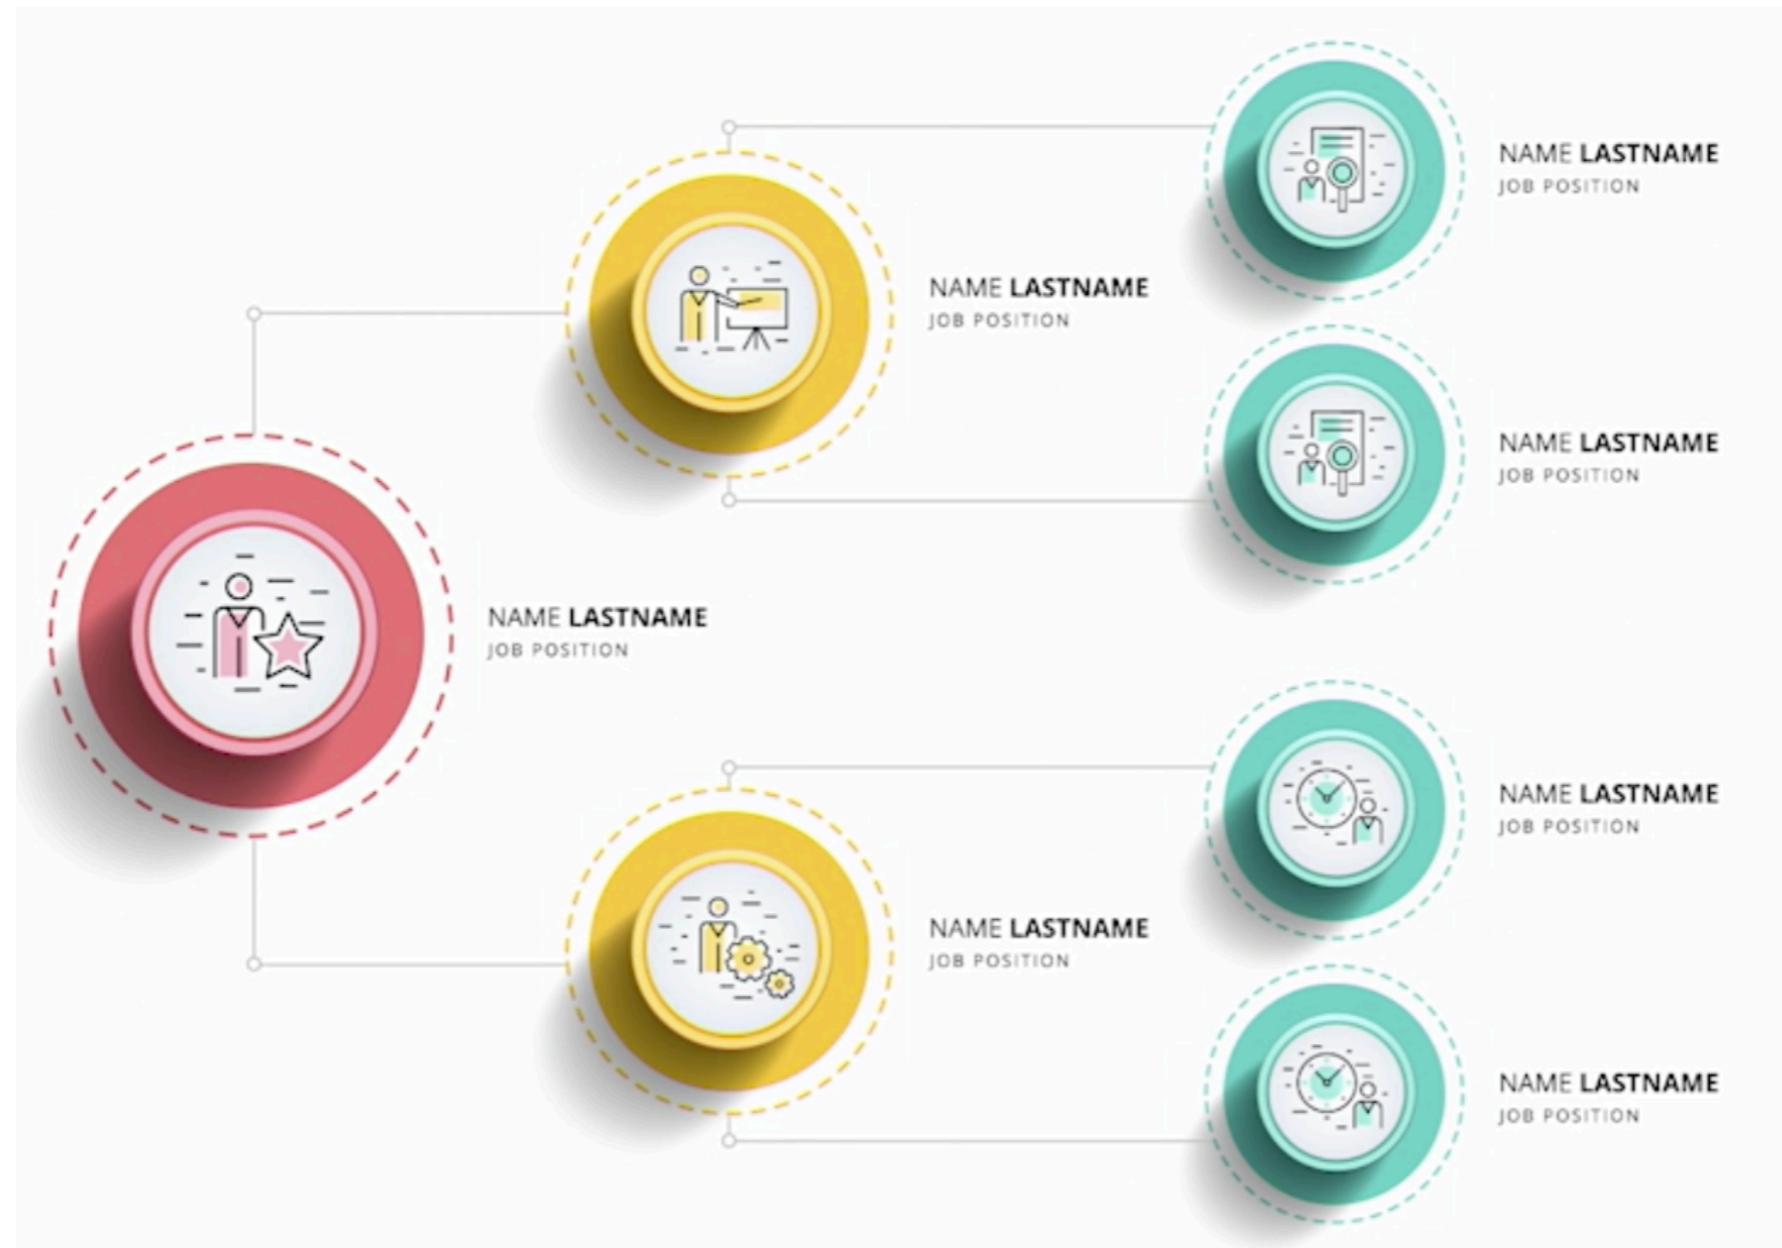
Point to location of row



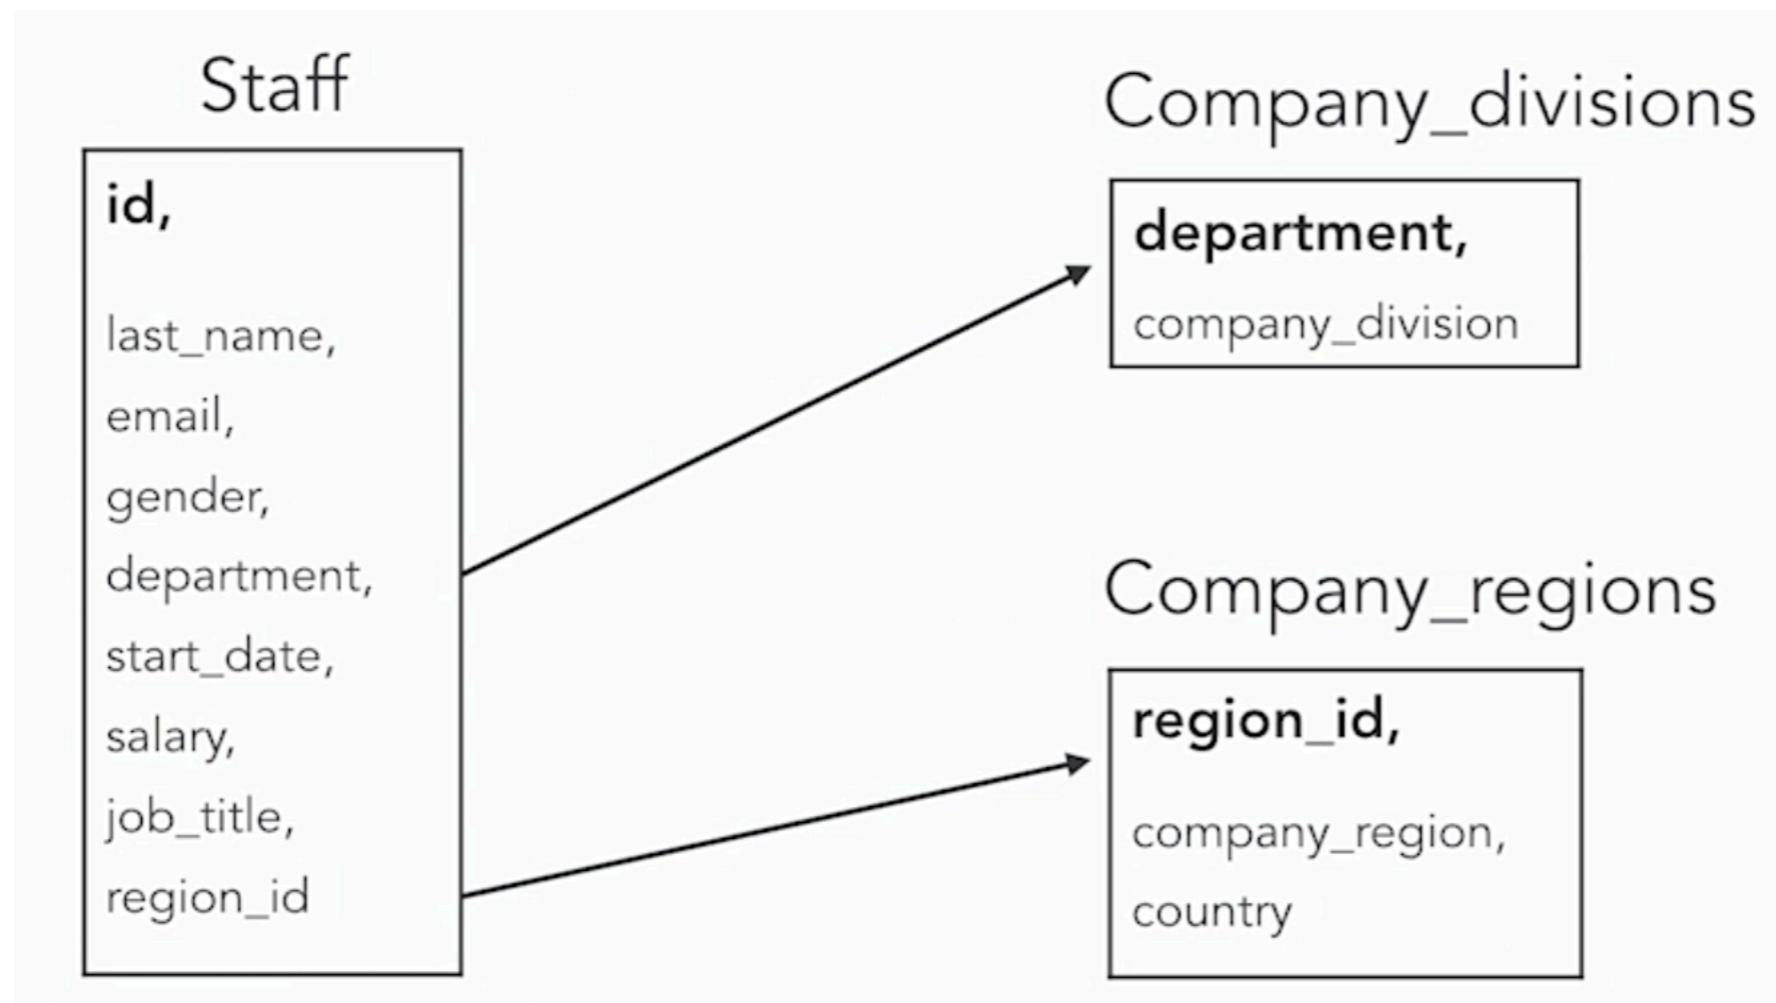
Indexing reduce full table scan



Indexes are optimised for search



Indexes are optimised for search



Purposes of indexes

Speed up access to data

Help enforce constraints

Indexes are ordered

Typically smaller than tables



Speeds

1

Main Memory

100 ns

2

**Read 1 MB
SSD**

1,000,000 ns

(1 ms)

3

**Read 1 MB
HDD**

20,000,000 ns

(20 ms)

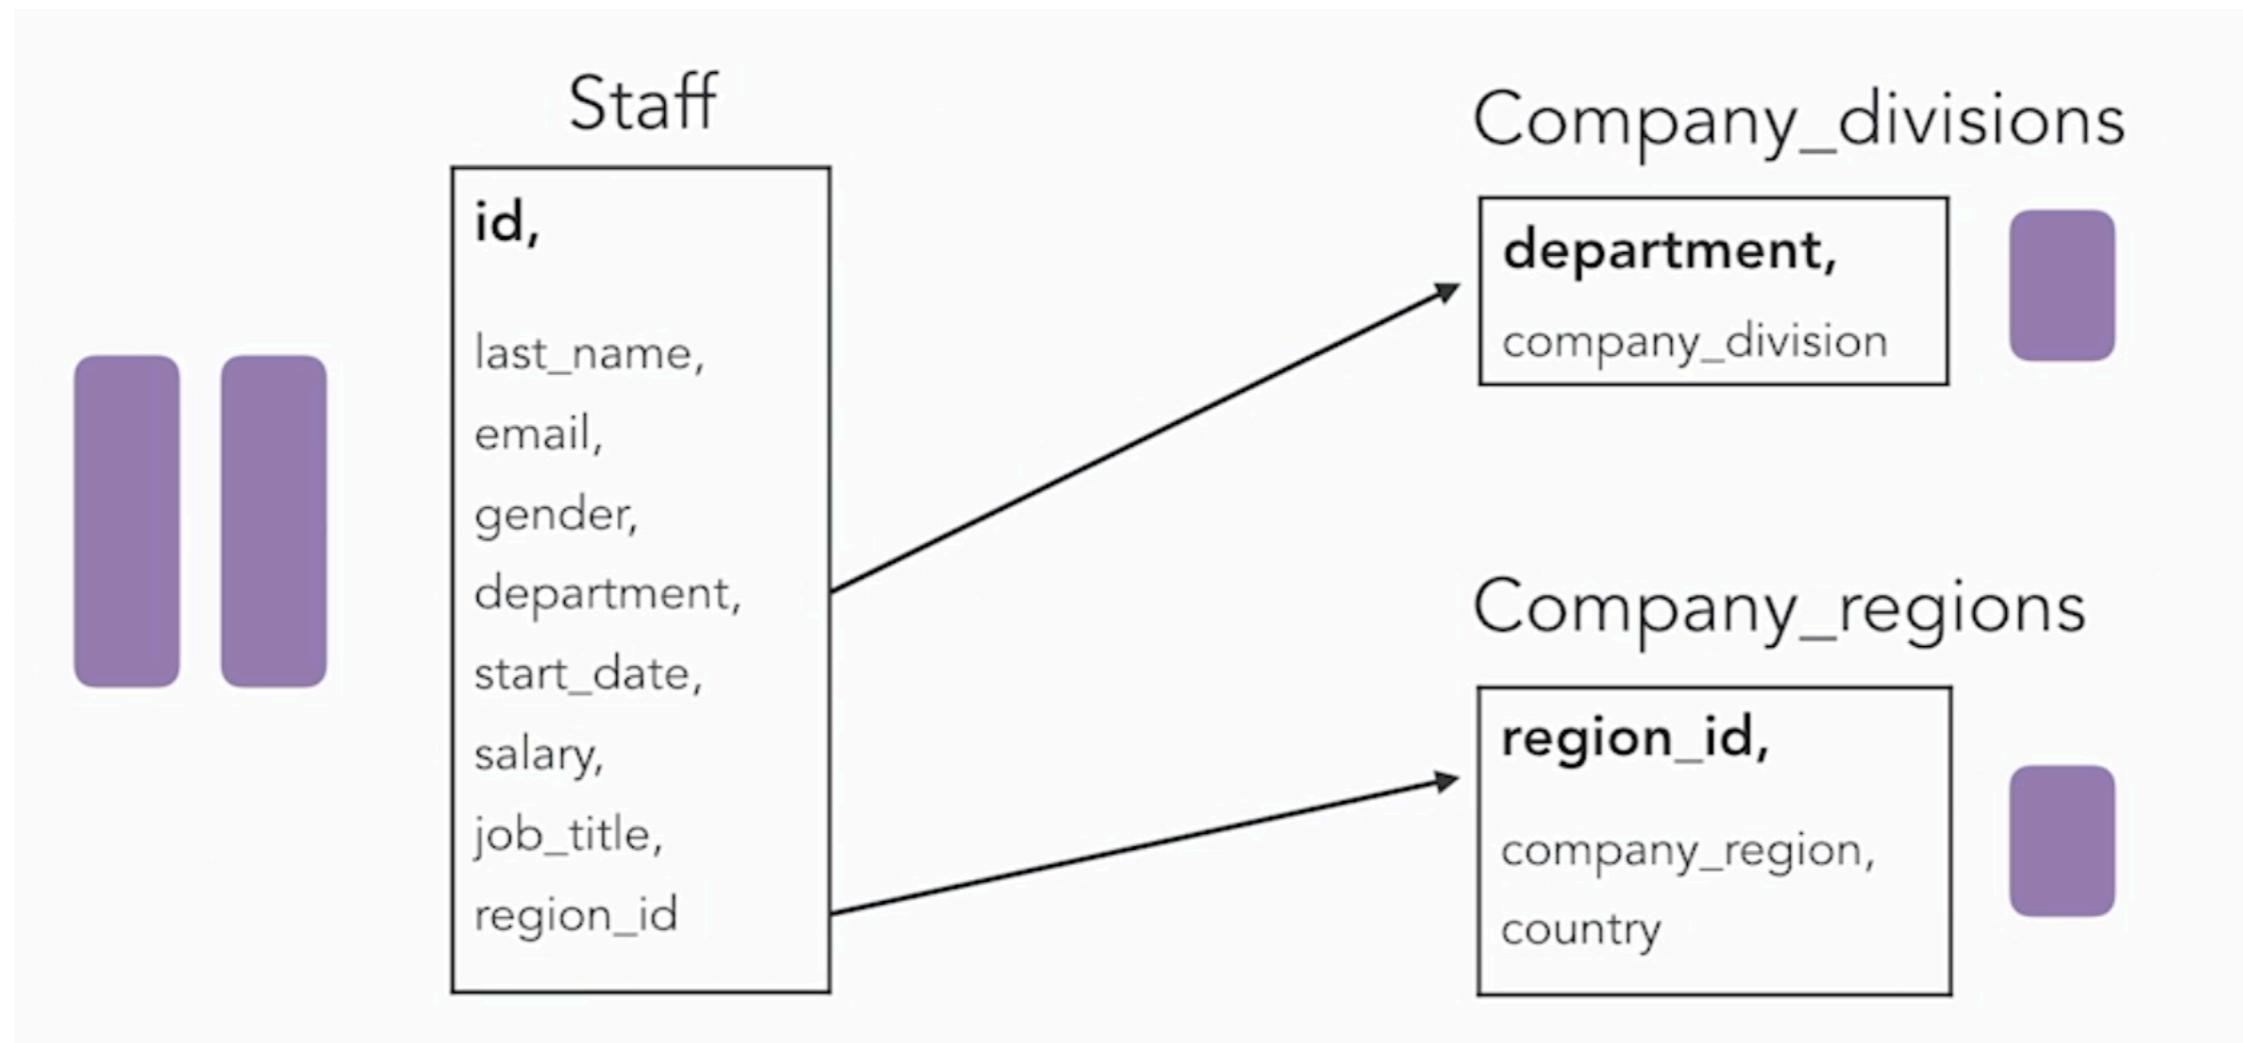


Implementation of indexes

Data structure separate from table
Sometime duplicate some data (e.g. key)
Organised differently than table data



Implementation of indexes



Types of indexes

B-tree for equality and range query

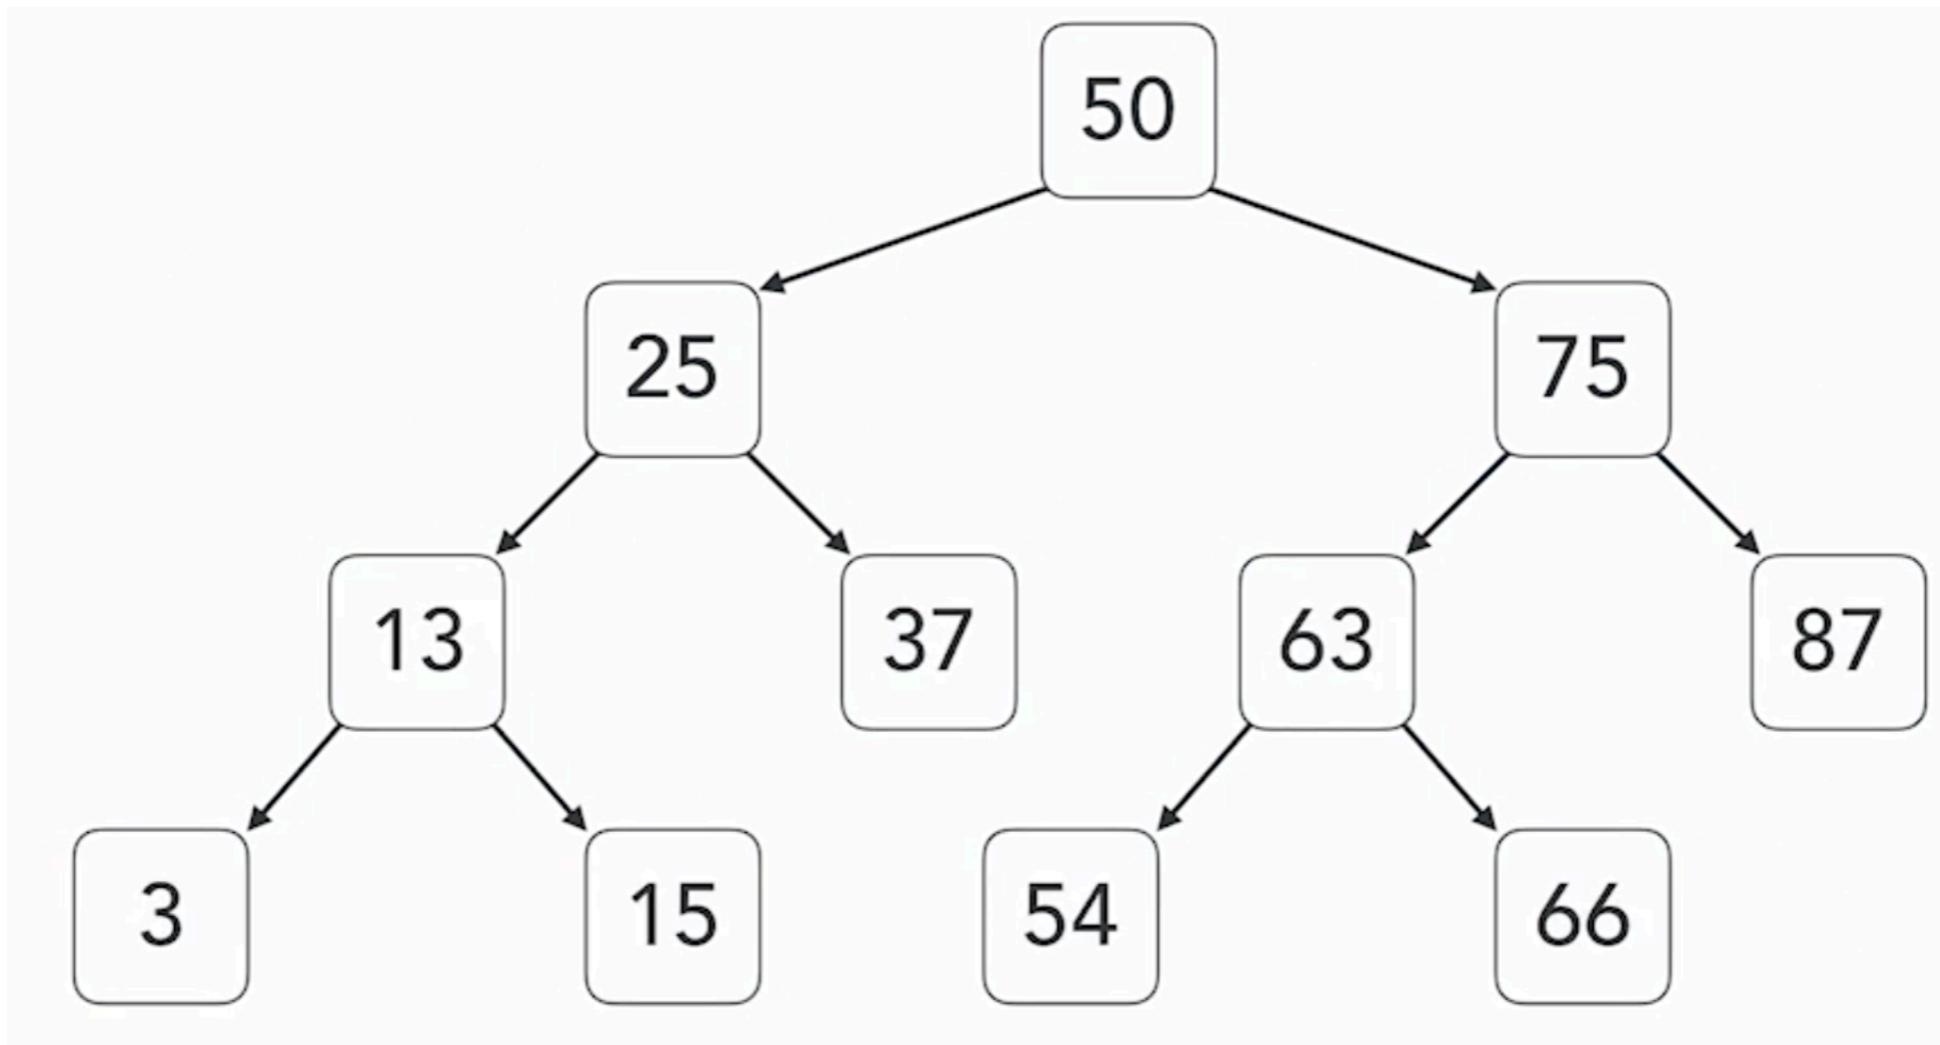
Hash for equality query

Bitmap for inclusion

Specialize for geo-spatial or user-defined



B-tree indexes



B-tree indexes

Most common type of index

Use when large number of possible values in a column

Rebalances as need

Time to access is based on depth of tree



Bitmap indexes

Use when small number of possible values in a column

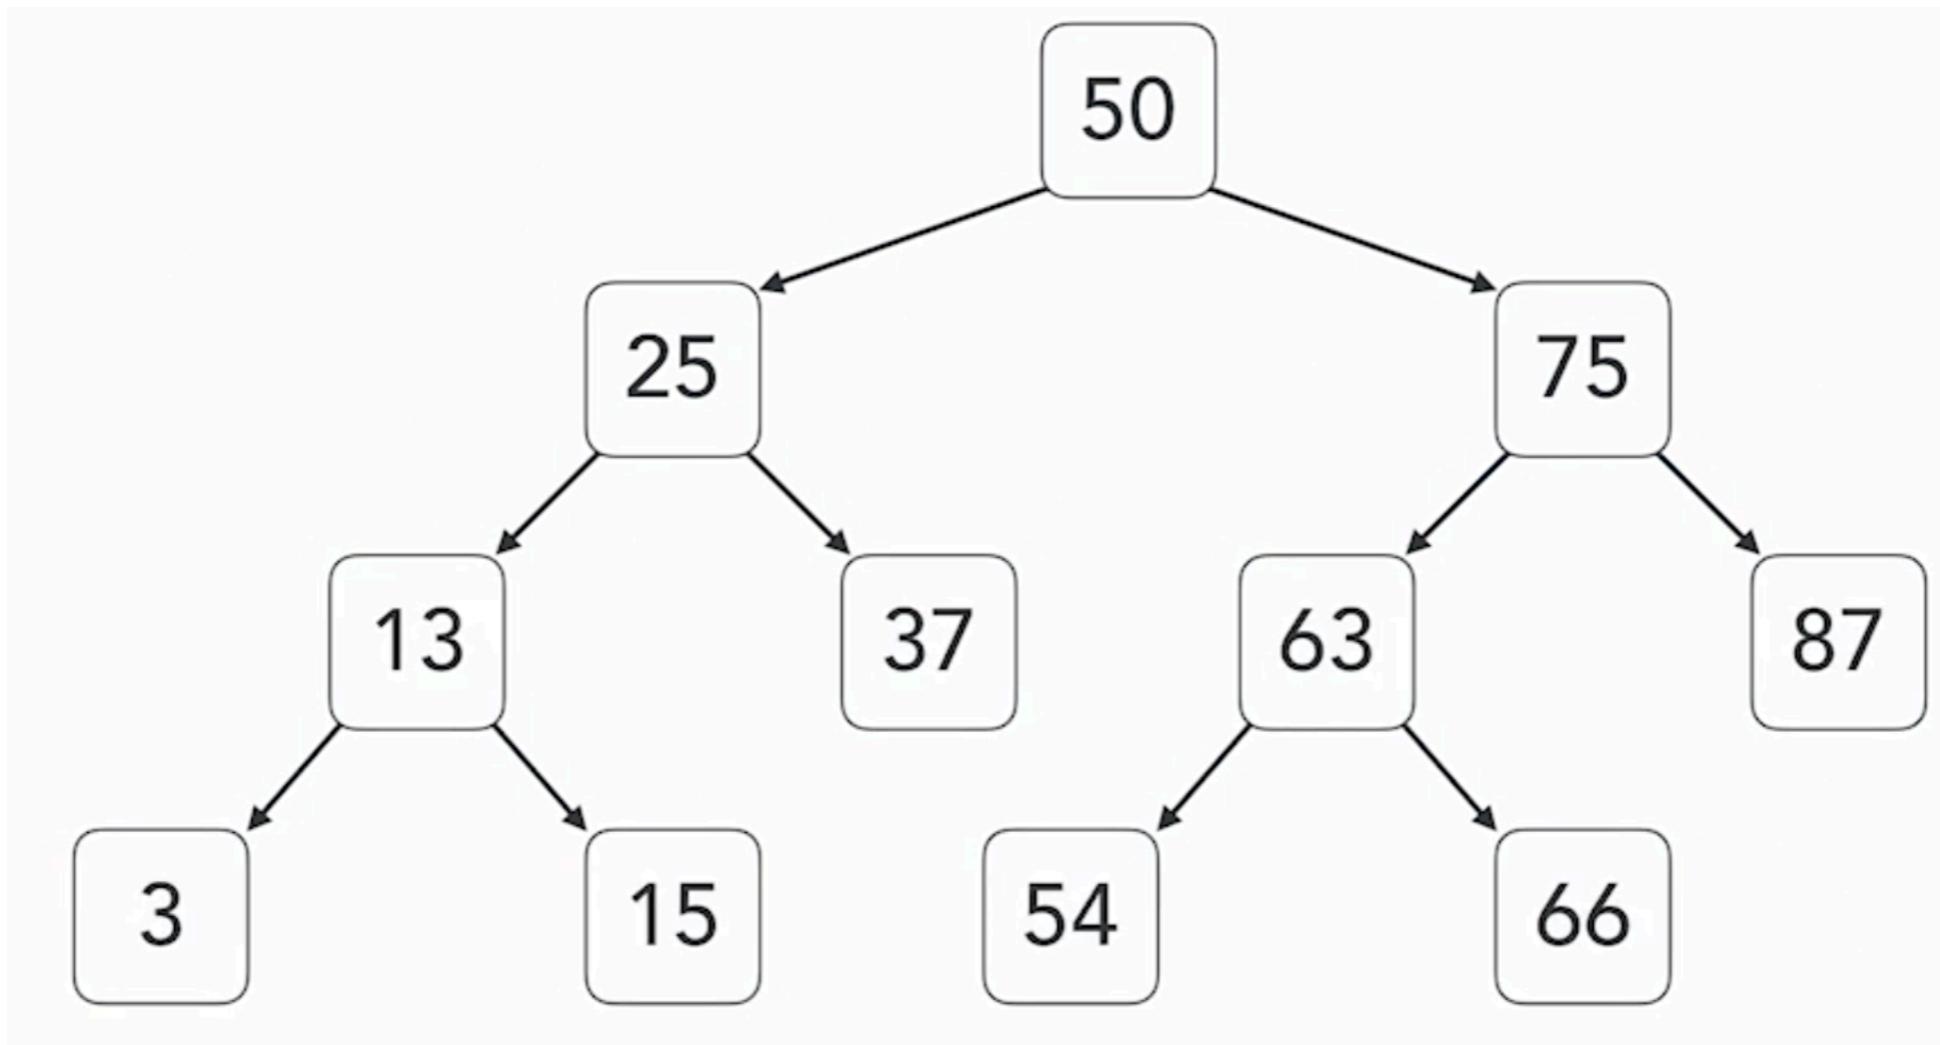
Filter by bitwise operations (and, or, not)

Time to access is based on time to perform bitwise operation

Use for read intensive, few write



B-tree indexes



Hash indexes

1

Equality Only

Hash indexes used when “=” is used, but not for ranges of values.

2

Smaller Size Than B-Tree

Latest versions of PostgreSQL (10+) have improved hash indexes.

3

As Fast as B-Tree

Builds and lookups are comparable; advantage is size; may fit in memory.



Workshop



Joining tables

customer_id	fname	lname	last_purchase_date	status_level
1234	Alice	Smith	23-Jan-2016	silver
2345	Bob	Washington	14-July-2018	bronze
2352	Charlie	Johnson	20-Dec-2018	gold
2678	Chendong	Mao	19-Dec-2017	gold
3422	Javier	Valdez	23-Dec-2018	gold
3577	Avery	Winston	5-Jan-2019	silver
4240	Charles	Davis	15-Jan-2017	bronze
6235	Vihaan	Patel	12-Nov-2018	gold
2352	Katherine	Coltrane	04-Apr-2016	silver

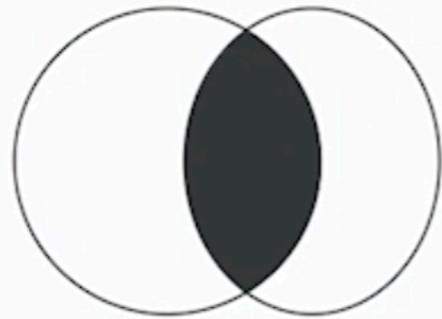


status_level	min_purchase_amt	discount_pct
gold	2500	7
silver	1500	5
bronze	500	3
platnimum	5000	10

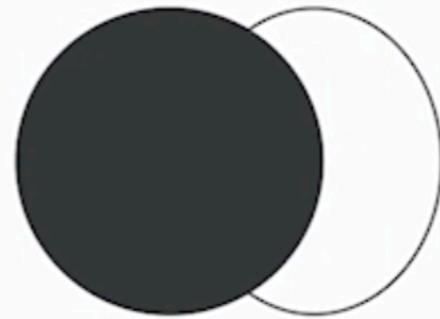
customer_id	fname	lname	last_purchase_date	status_level	status_level	min_purchase_amt	discount_pct
2345	Bob	Washington	14-July-2018	bronze	bronze	500	3



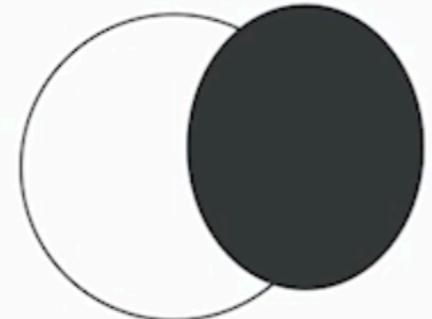
Types of joining



INNER JOIN



LEFT OUTER JOIN



RIGHT OUTER JOIN



FULL OUTER JOIN



Inner join

Most common join

Return rows from both tables that have corresponding row in other table

```
SELECT
  *
FROM
  company_region cr,
INNER JOIN
  staff s
ON
  cr.region_id = s.region_id
```



Left outer join

Return all rows from left table

Return rows from right table that have matching key

```
SELECT
  *
FROM
  company_region cr,
LEFT OUTER JOIN
  staff s
ON
  cr.region_id = s.region_id
```



Right outer join

Return all rows from right table

Return rows from left table that have matching key

```
SELECT
  *
FROM
  company_region cr,
RIGHT OUTER JOIN
  staff s
ON
  cr.region_id = s.region_id
```



Full Outer join

Return all rows from both tables

```
SELECT
  *
FROM
  company_region cr,
  FULLER OUTER JOIN
    staff s
  ON
    cr.region_id = s.region_id
```



Joining

1

Prefer INNER JOINS

Often fastest

2

OUTER and FULL JOINS

Require additional steps in addition to INNER JOIN

3

When need to NULLs

LEFT, RIGHT, and FULL OUTER JOINS used when need to NULLs



How to match rows ?

Foreign keys in one table
Primary key in another table
How to find matching keys ?

customer_id	fname	lname	last_purchase_date	status_level
1234	Alice	Smith	23-Jan-2016	silver
2345	Bob	Washington	14-July-2018	bronze
2352	Charlie	Johnson	20-Dec-2018	gold
2678	Chendong	Mao	19-Dec-2017	gold
3422	Javier	Valdez	23-Dec-2018	gold
3577	Avery	Winston	5-Jan-2019	silver
4240	Charles	Davis	15-Jan-2017	bronze
6235	Vihaan	Patel	12-Nov-2018	gold
2352	Katherine	Coltrane	04-Apr-2016	silver

status_level	min_purchase_amt	discount_pct
gold	2500	7
silver	1500	5
bronze	500	3
platnimum	5000	10



Joining tables

1

Nested Loop Join

Compare all rows in both tables to each other.

2

Hash Join

Calculate hash value of key and join based on matching hash values.

3

Sort Merge Join

Sort both tables and then join rows while taking advantage of order.



Workshop



Sub-queries vs Join



Sub-queries

Returns a value from a related tables

```
SELECT
    s.id, s.last_name, s.department,
    (SELECT
        company_regions
    FROM
        company_regions cr
    WHERE
        cr.region_id = s.region_id) region_name
FROM
    staff s
```



Sub-queries vs Join

Same logical outcome

SQL queries specify “what”

More than one way to express the same thing



Sub-queries vs Join

1

Conventional Wisdom

Always use joins;
they are more
efficient.

2

Improved Query Plans

Query plan builders
are more effective
at optimizing
subqueries.

3

Maximize Clarity

Both will work well
in many cases. Opt
for what makes
intention clear.



Partition data ?

Storing table data in multiple sub-tables
Called “Partition”



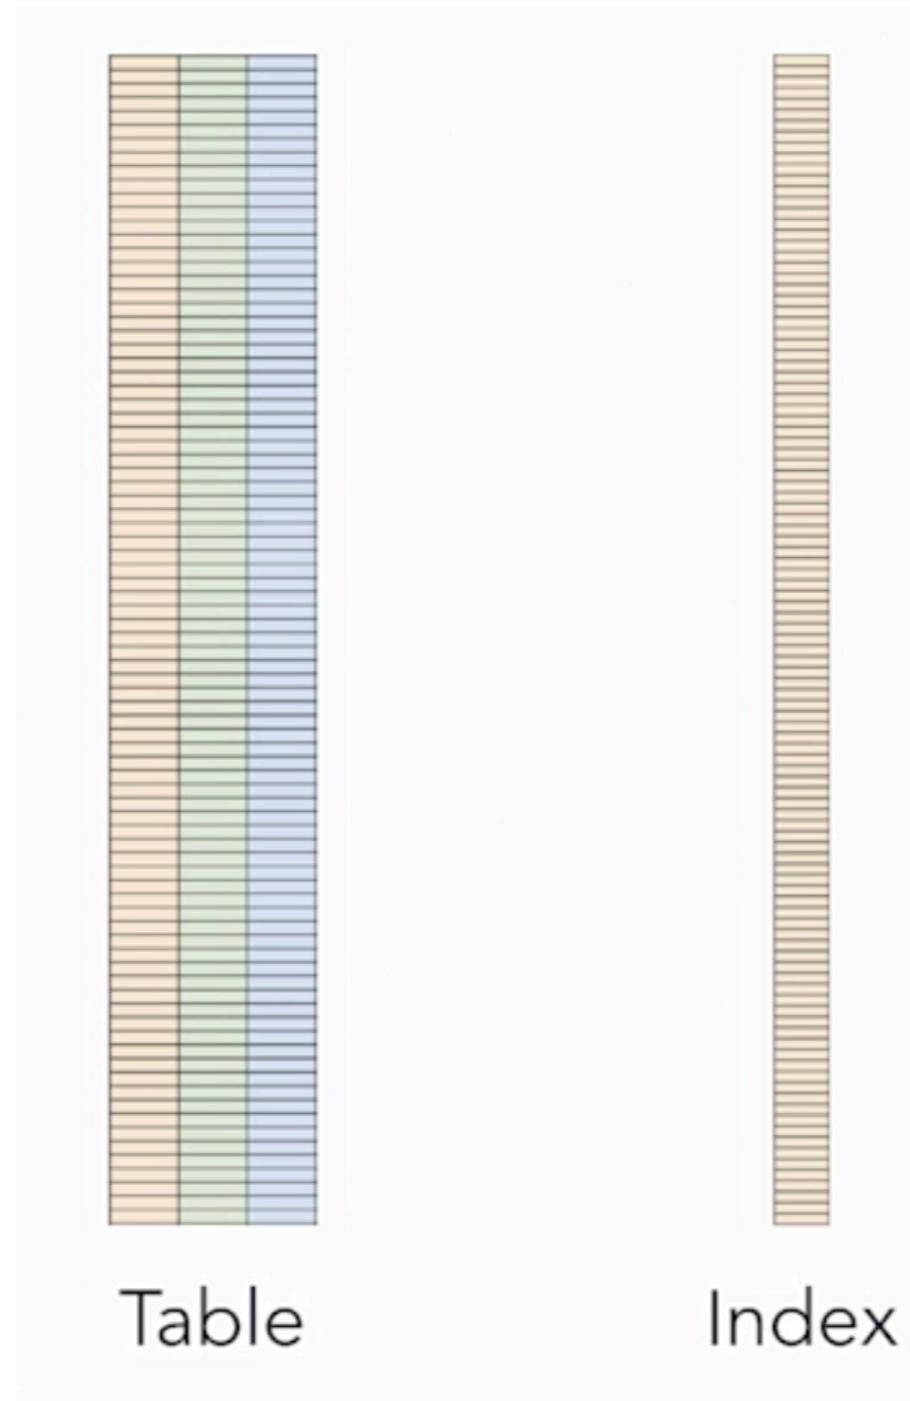
Partition data ?

Used to improve query, load and delete operation

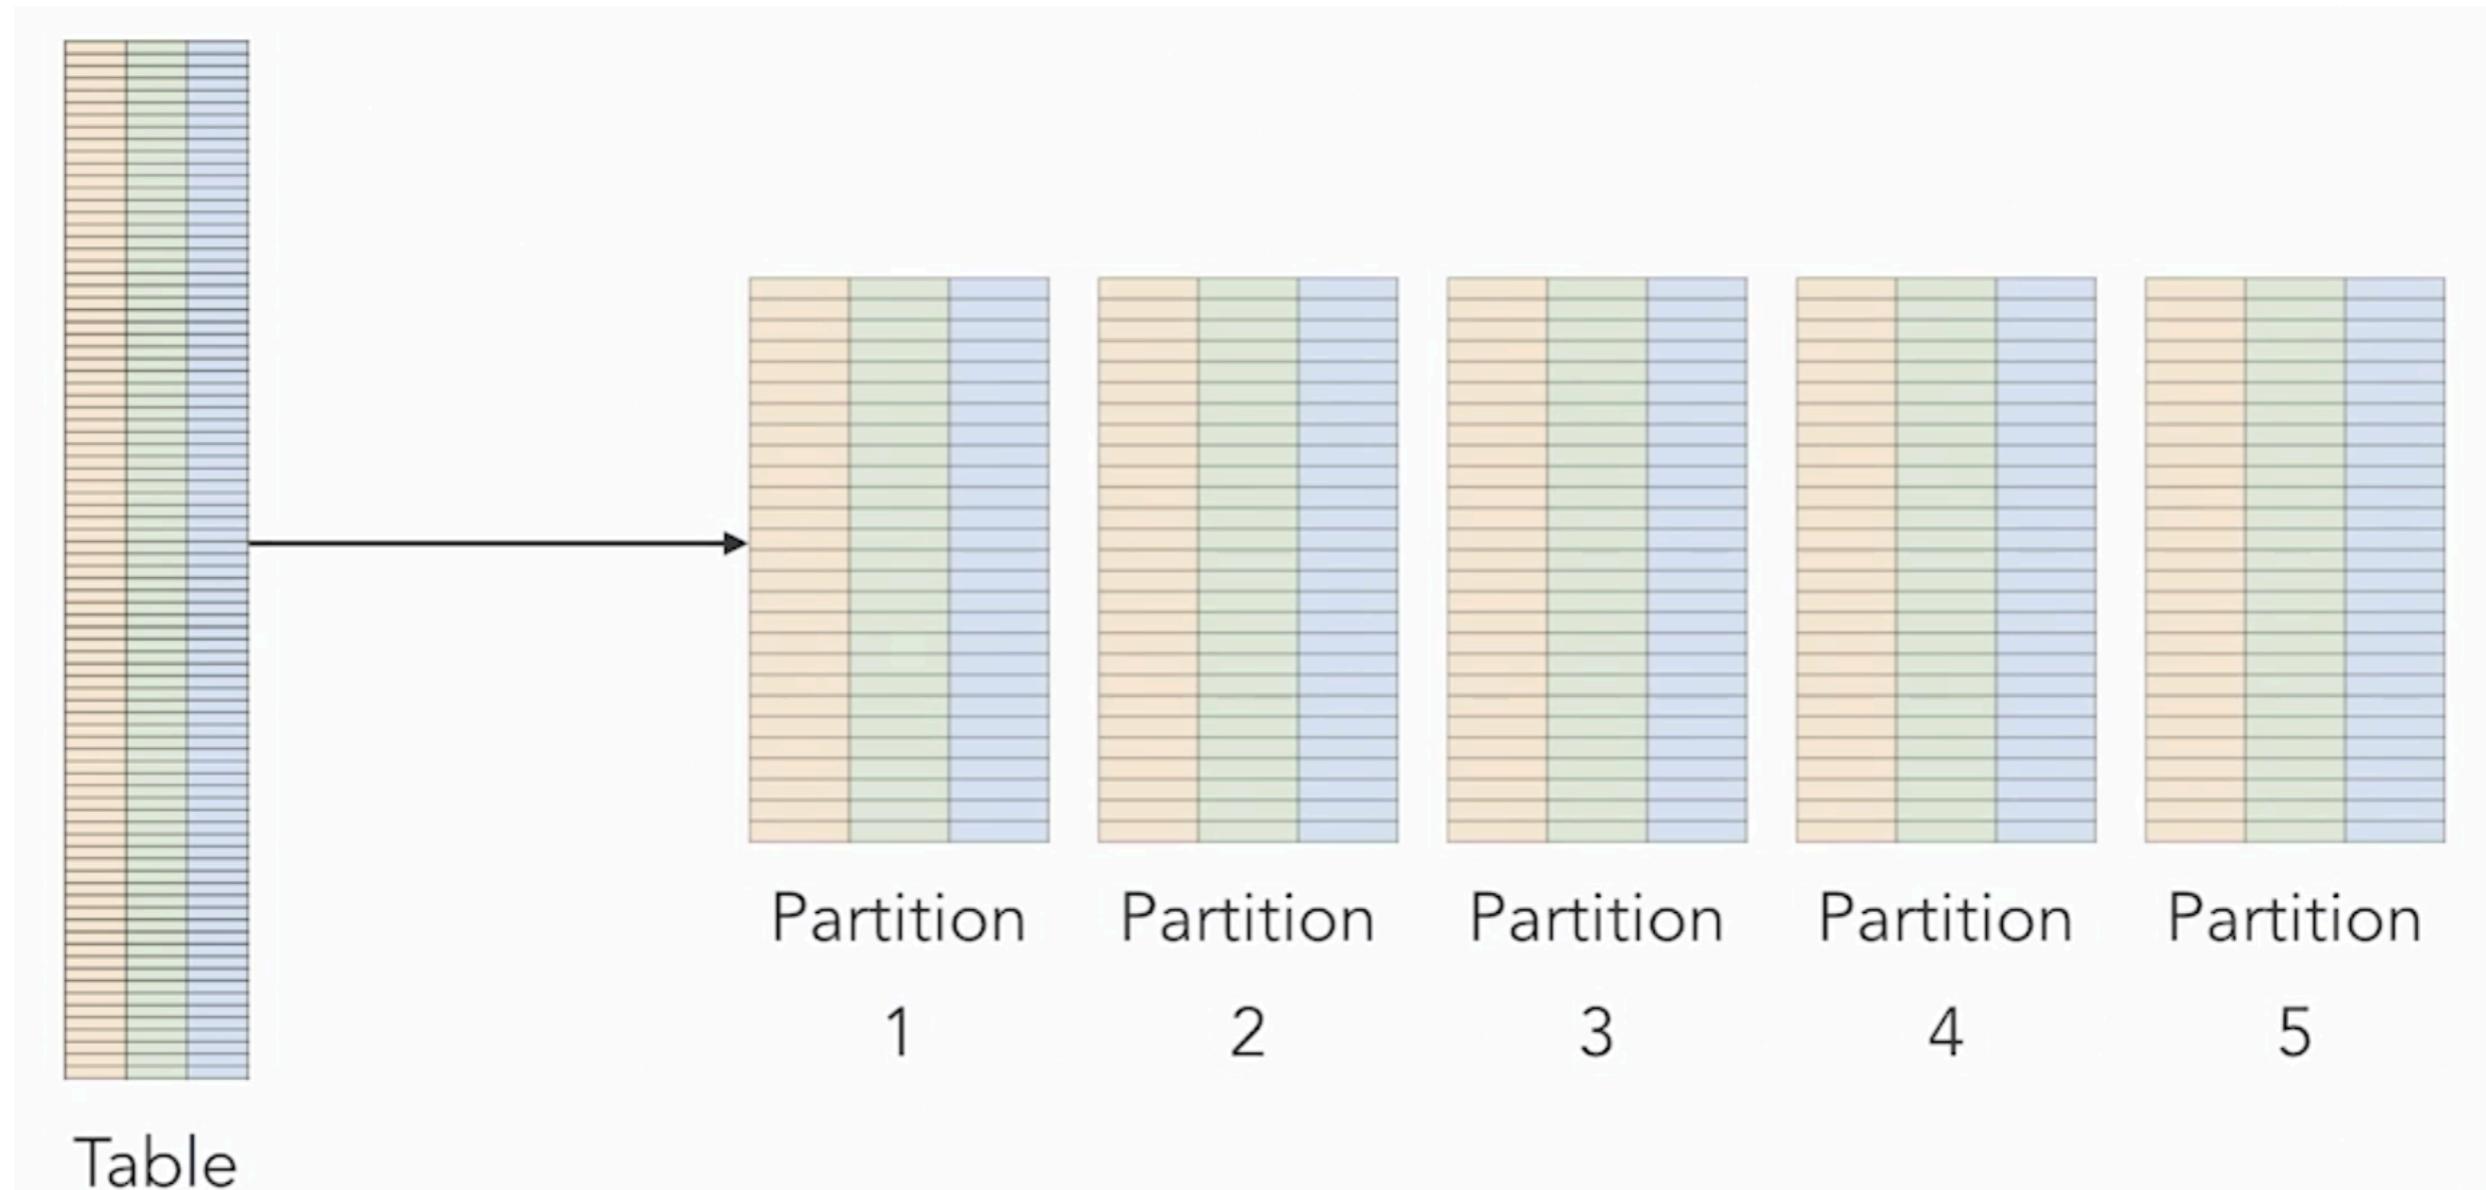
Used for large tables
Expensive



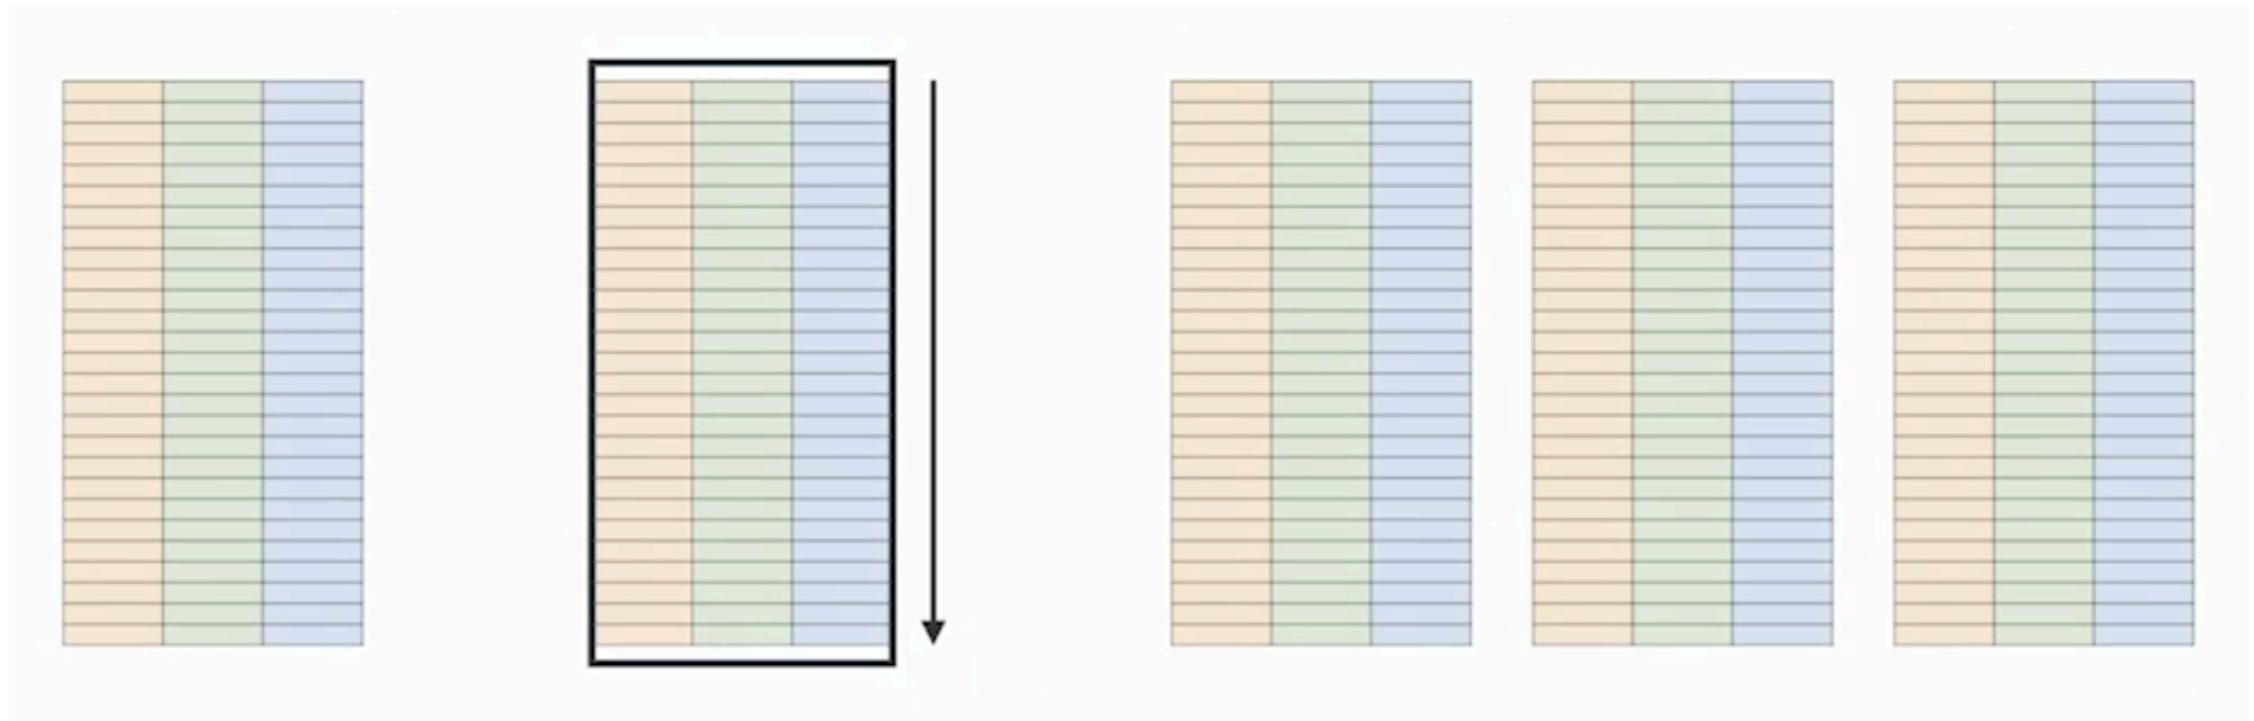
Large table = Large index



Partition



Faster Scan

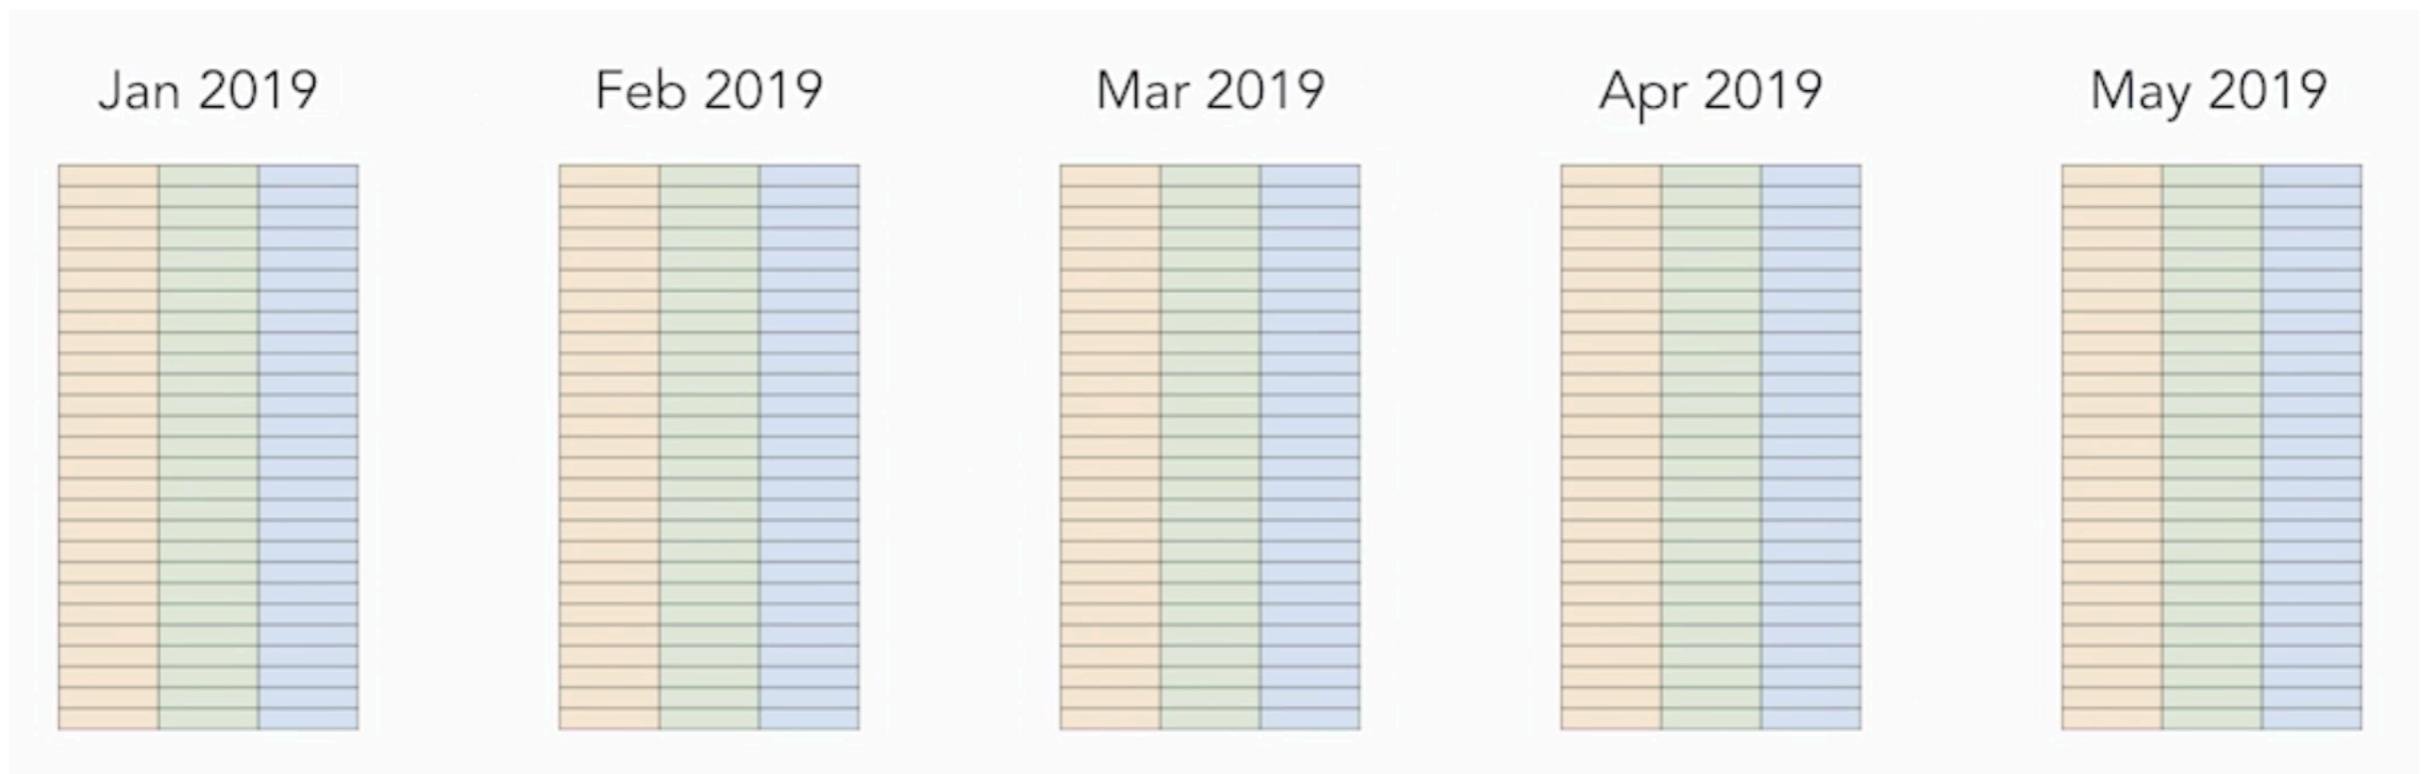


Partition keys

By range

By list

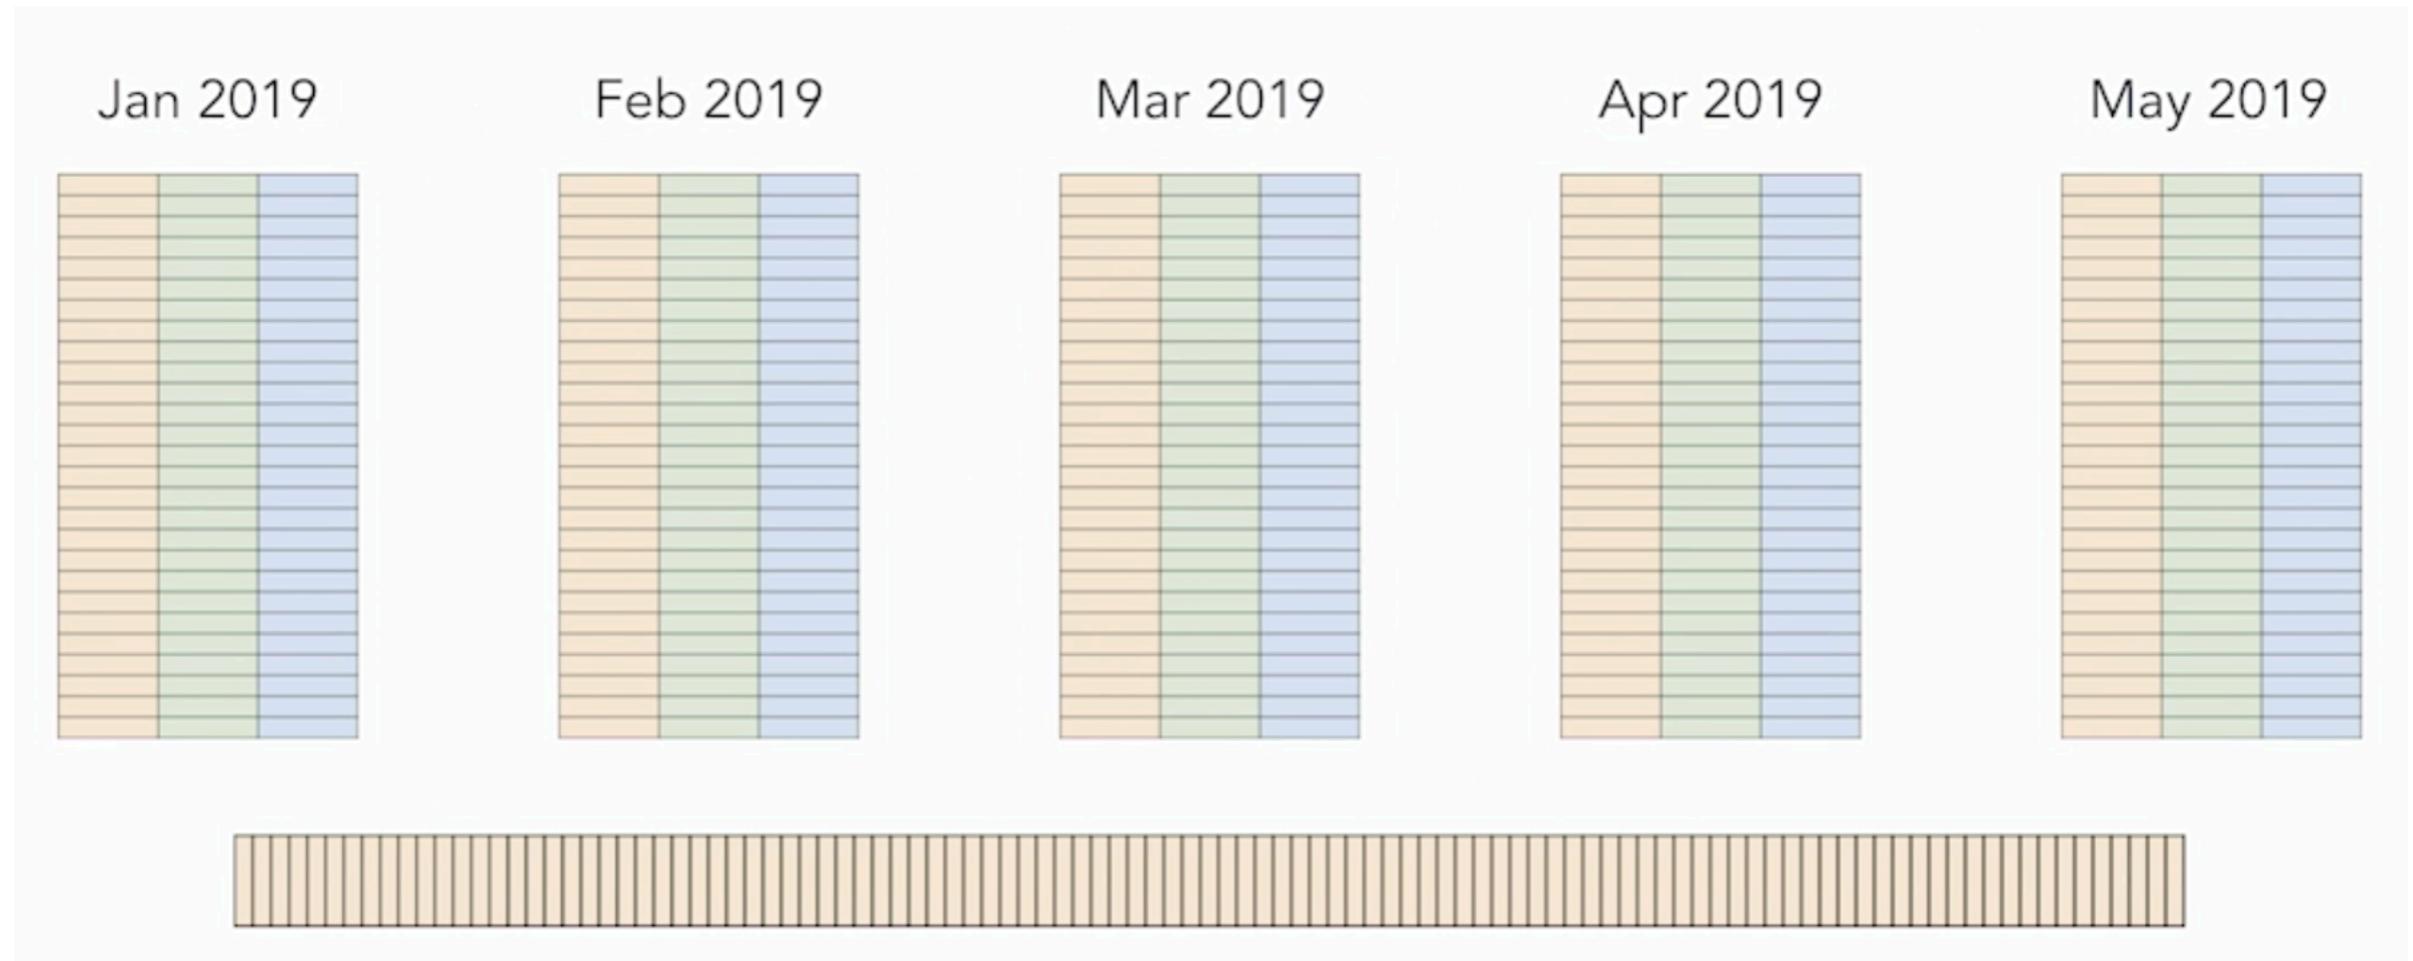
By hash



Local index in each partition



Global index for all partitions

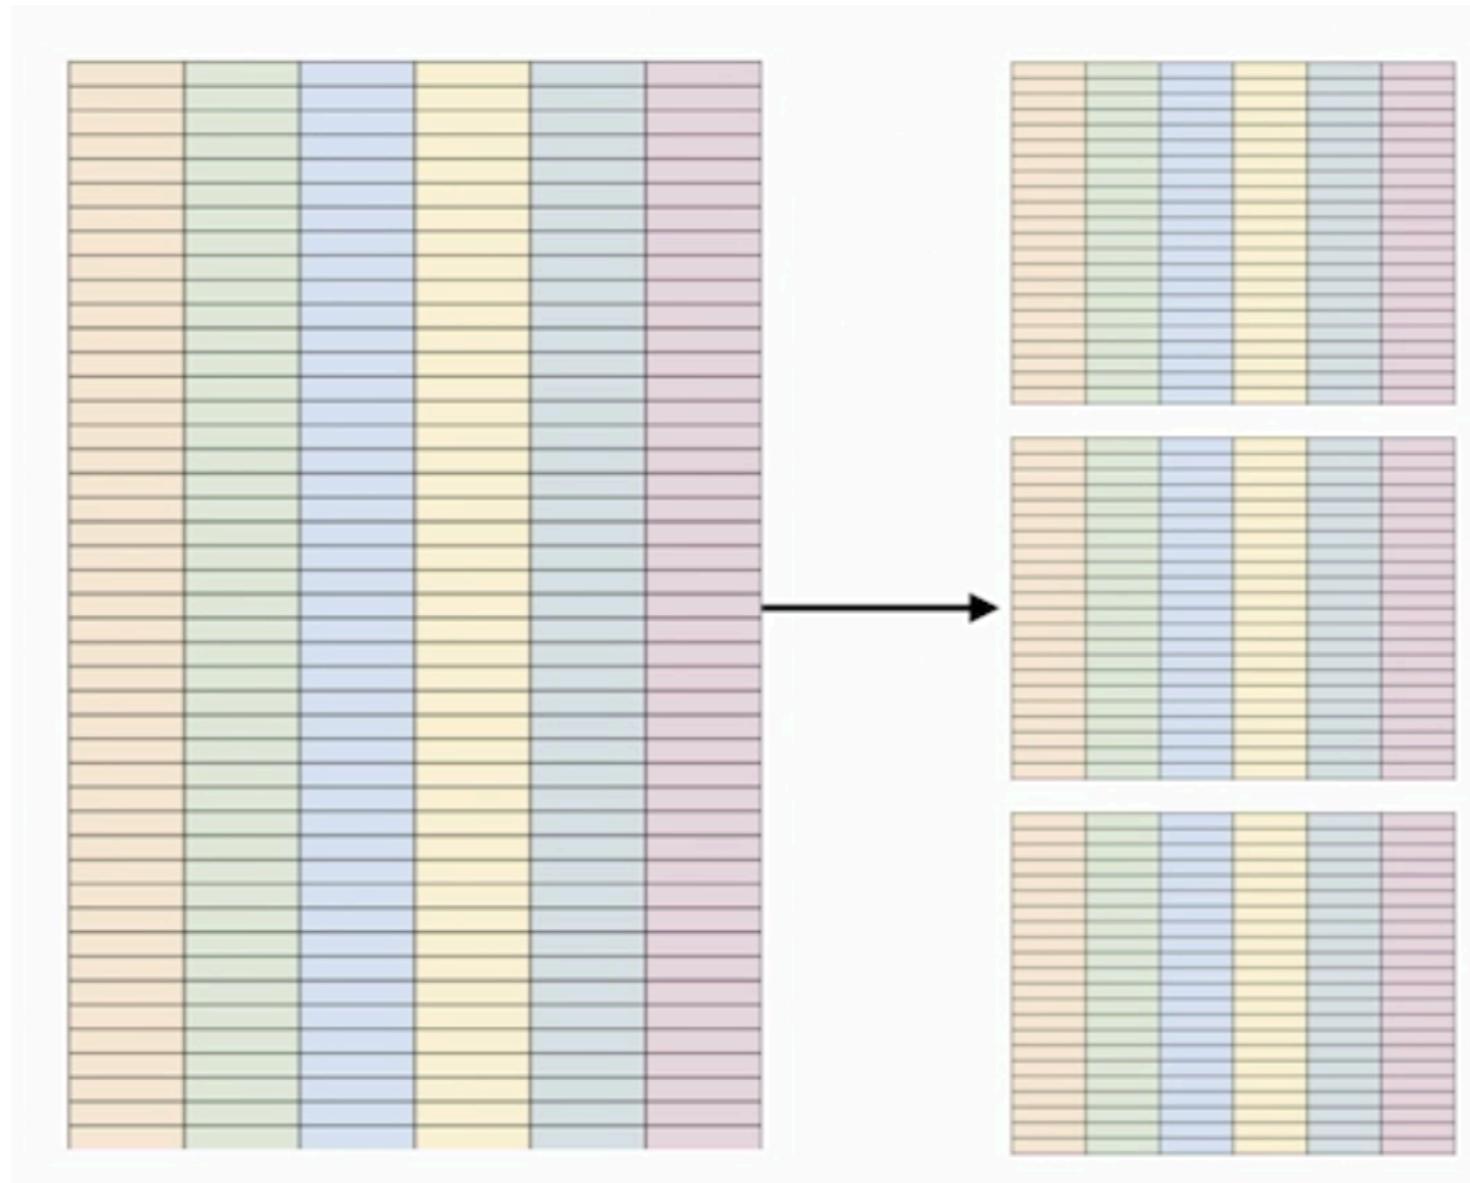


Partition types

Horizontal
Vertical



Horizontal partition



Horizontal partition

Large tables can be difficult to query efficient

Split tables by rows in to partitions

Treat each partition like a table



Types of horizontal partition

Partition by range

Partition by list

Partition by hash



Partition by range

Partition by date

Partition by numeric

Partition by alphabet

By non-overlapping keys



Partition by range

1

Partition Key

Determines which partition is used for data

2

Partition Bounds

Minimum and maximum values allowed in the partition

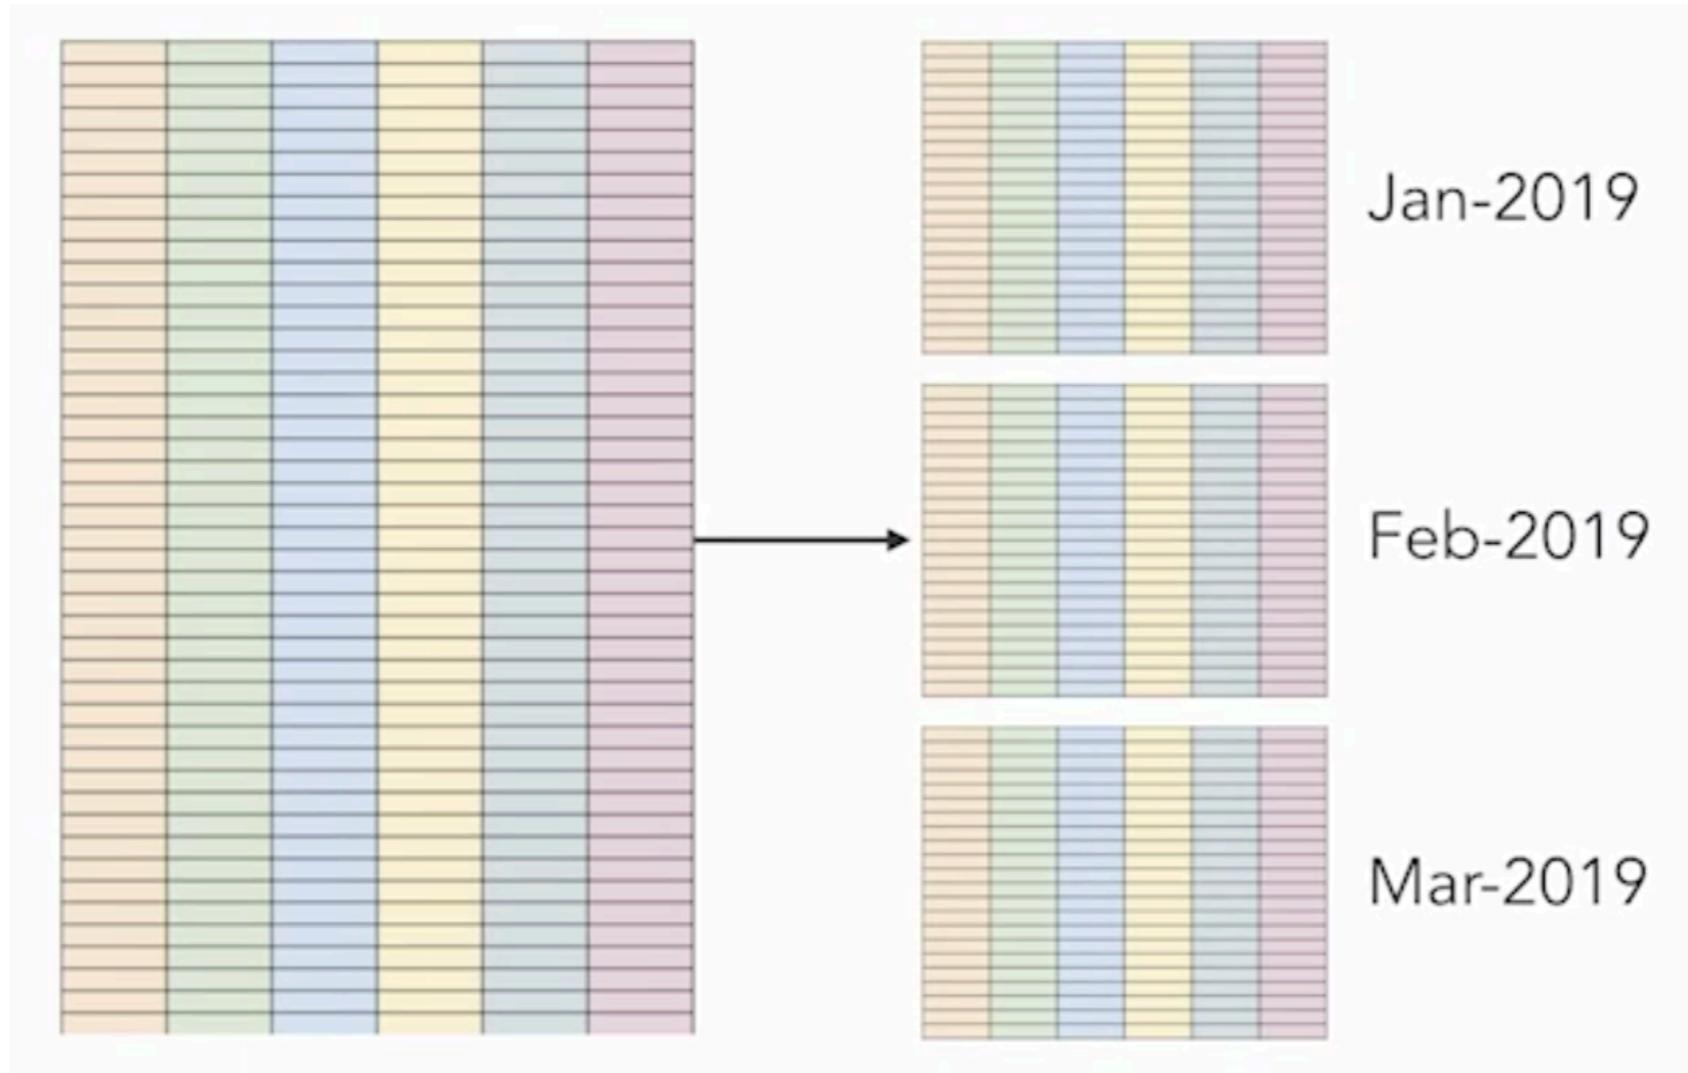
3

Constraints

Each partition can have its own indexes, constraints, and defaults



Example partition by date



IoT data

Partition by date

```
CREATE TABLE iot_measurement
  ( location_id int not null,
    measure_date date not null,
    temp_celcius int,
    rel_humidity_pct int)
    PARTITION BY RANGE (measure_date);
```



IoT data

Partition by date

```
CREATE TABLE iot_measurement
( location_id int not null,
  measure_date date not null,
  temp_celcius int,
  rel_humidity_pct int)
PARTITION BY RANGE (measure_date);
```

```
CREATE TABLE iot_measurement_wk1_2019 PARTITION OF iot_measurement
  FOR VALUES FROM ('2019-01-01') TO ('2019-01-08');

CREATE TABLE iot_measurement_wk2_2019 PARTITION OF iot_measurement
  FOR VALUES FROM ('2019-01-08') TO ('2019-01-15');

CREATE TABLE iot_measurement_wk3_2019 PARTITION OF iot_measurement
  FOR VALUES FROM ('2019-01-15') TO ('2019-01-22');
```



When to use ?

Query latest data

Comparative queries

Report within range

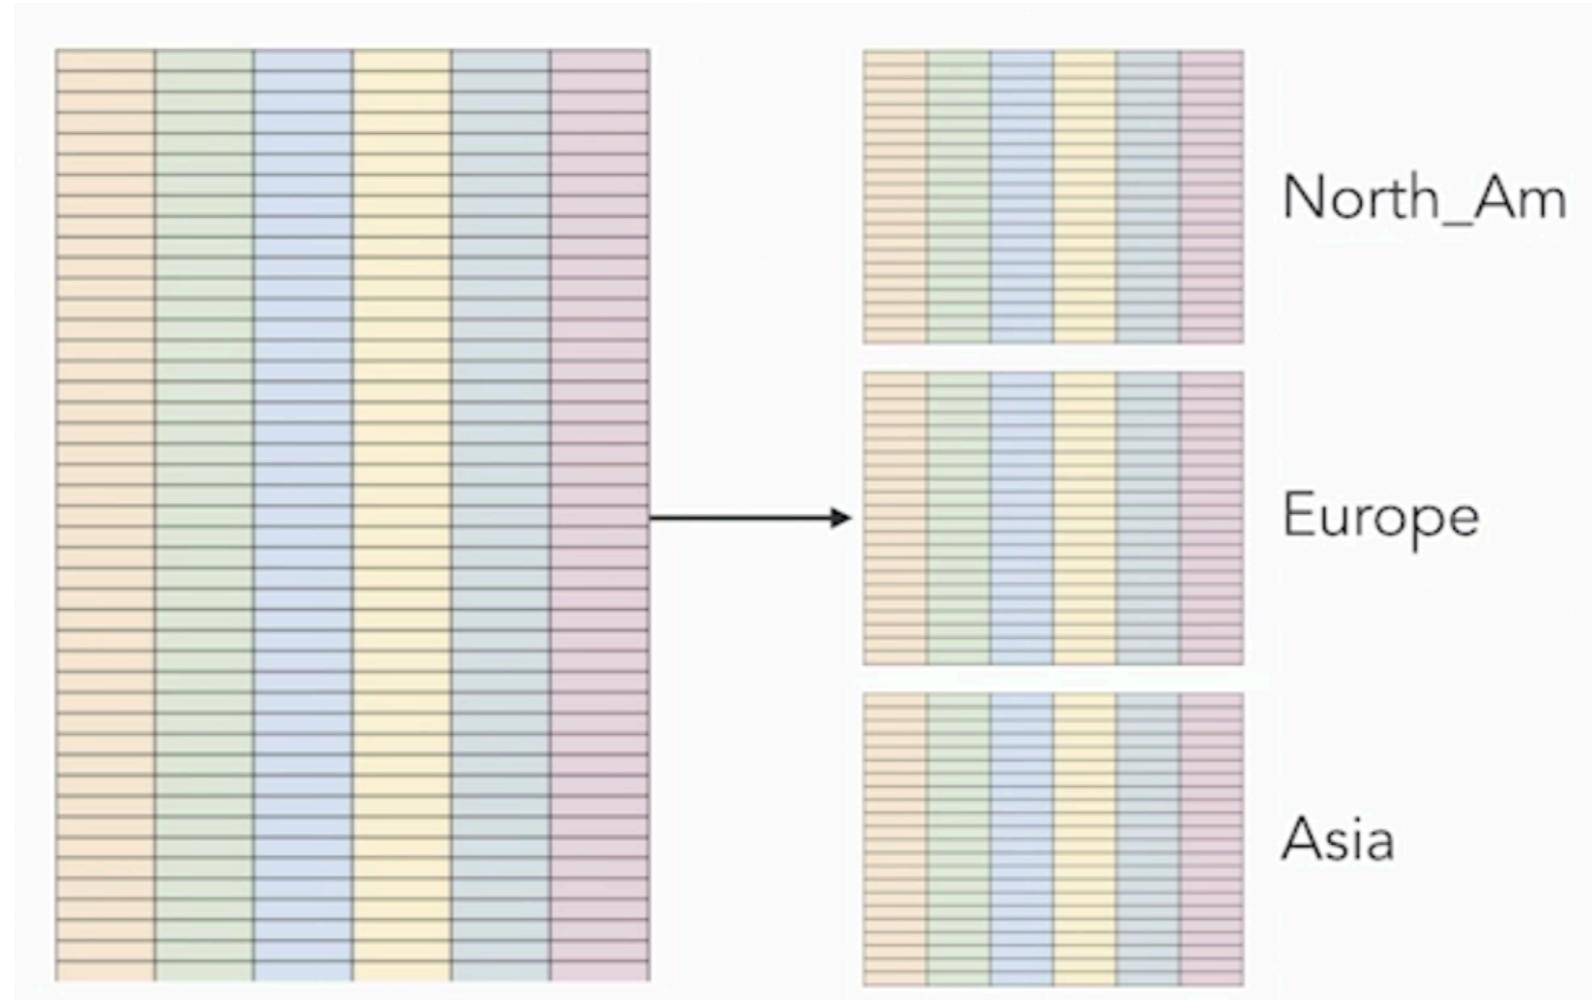
Drop data after a period of time



Workshop



Partition by list



Partition by list

Partition on value of list of values

By non-overlapping keys



Partition by list

1

Partition Key

Determines which partition is used for data

2

Partition Bounds

List of values for a partition

3

Constraints

Each partition can have its own indexes, constraints, and defaults



Example :: Product catalog

Partition by list of product catalog

```
CREATE TABLE products
  (prod_id    int not null,
   prod_name  text not null,
   prod_short_descr text not null,
   prod_long_descr text not null,
   prod_category varchar)
PARTITION BY LIST (prod_category);
```



Example :: Product catalog

Partition by list of product catalog

```
CREATE TABLE products
(prod_id  int not null,
 prod_name text not null,
 prod_short_descr text not null,
 prod_long_descr text not null,
 prod_category varchar)
PARTITION BY LIST (prod_category);
```

```
CREATE TABLE product_clothing PARTITION OF products
FOR VALUES IN ('casual_clothing', 'business_attire', 'formal_clothing');

CREATE TABLE product_electronics PARTITION OF products
FOR VALUES IN ('mobile_phones', 'tablets', 'laptop_computers');

CREATE TABLE product_kitchen PARTITION OF products
FOR VALUES IN ('food_processor', 'cutlery', 'blenders');
```



When to use ?

Data logically groups into subgroups

Often query within subgroups

Data not time oriented

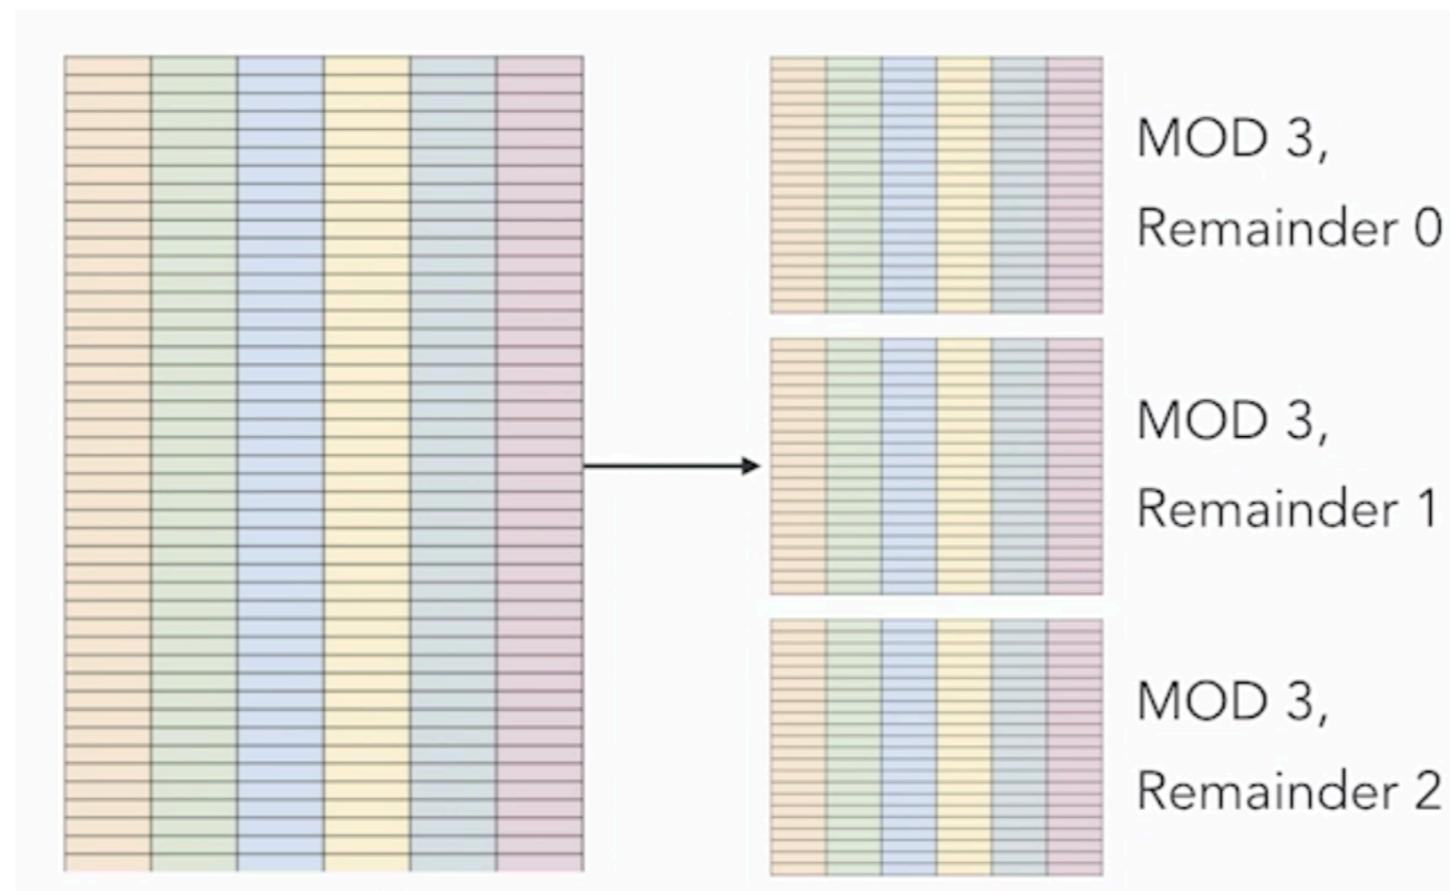


Workshop



Partition by hash

Partition on modulus of hash of partition keys



Partition by hash

1

Partition Key

Determines which partition is used for data

2

Modulus

Number of partitions

3

Availability

PostgreSQL 11+, Oracle, MySQL



Example :: Web Path Analytic

Hash function on ci_id

```
CREATE TABLE customer_interaction
  (ci_id int not null,
   ci_url text not null,
   time_at_url int not null,
   click_sequence int not null)
PARTITION BY HASH (ci_id);
```



Example :: Web Path Analytic

Hash function on ci_id

```
CREATE TABLE customer_interaction
(ci_id int not null,
 ci_url text not null,
 time_at_url int not null,
 click_sequence int not null)
PARTITION BY HASH (ci_id);
```

```
CREATE TABLE customer_interaction_1 PARTITION OF customer_interaction
FOR VALUES WITH (MODULUS 5, REMAINDER 0);
CREATE TABLE customer_interaction_2 PARTITION OF customer_interaction
FOR VALUES WITH (MODULUS 5, REMAINDER 1);
CREATE TABLE customer_interaction_3 PARTITION OF customer_interaction
FOR VALUES WITH (MODULUS 5, REMAINDER 2);
CREATE TABLE customer_interaction_4 PARTITION OF customer_interaction
FOR VALUES WITH (MODULUS 5, REMAINDER 3);
CREATE TABLE customer_interaction_5 PARTITION OF customer_interaction
FOR VALUES WITH (MODULUS 5, REMAINDER 4);
```



When to use ?

Data doesn't logically group into subgroups

Distribute data across partitions

No need for subgroup-operations



Workshop



Benefits

- Limit scans to subset of partitions
- Local indexes for each partition
- Efficient for add and delete data



Use cases

1

Data Warehouses

- Partition on time
- Query on time
- Delete by time

2

Timeseries

- Most likely query latest data
- Summarize data in older partitions

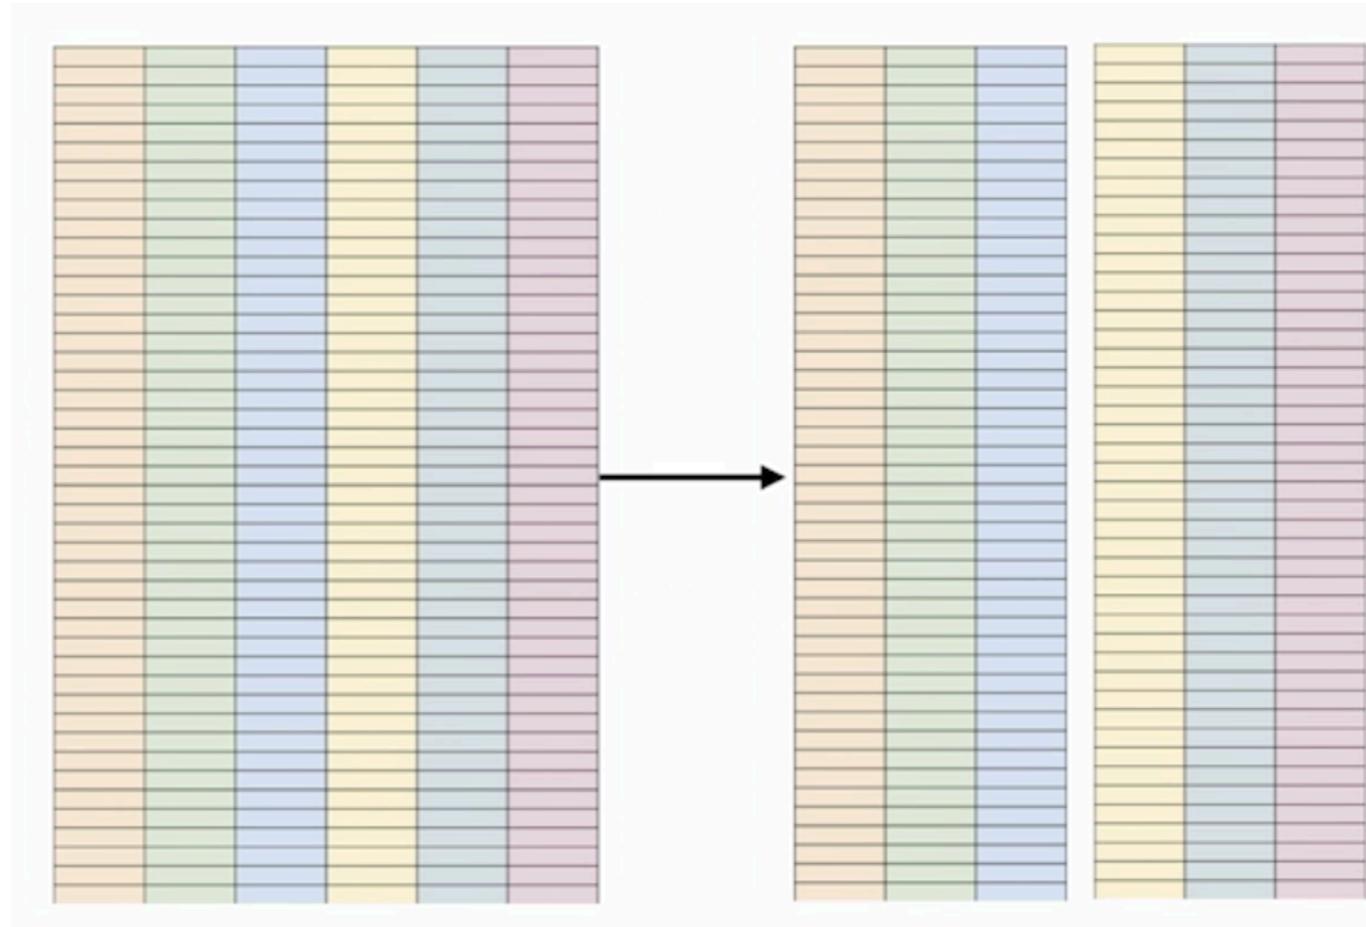
3

Naturally Partition Data

- Retailer, by geography
- Data science, by product category



Vertical partition



Vertical partition

Implement as separate tables

No partitioning-specific definitions are required



Vertical partition

Separate columns into multiple tables

Keep frequently queried columns together

Used same primary key in all tables



Benefits

Increase number of rows in data block
Global indexes for each partition
Reduce I/O



Use cases

1

Data Warehouses

Partition on groups
of attributes

2

Many Attributes

Wide variety of
products, each with
different attributes

3

Data Analytics

Statistics on subset
of attributes; after
factor analysis



Workshop

